

Chapter 8: Gashlycrumb: Looking items up in a dictionary

Every time you log into a website, the code behind it has to look up your username and password to compare to the values you put into the login form. Whenever you give your phone number at the hardware store or scan your library card to checkout a book, a computer program uses one piece of information to find other things like how often you buy compost or if you have any overdue books. Probably all these examples would be using a database to find that information. We're going to use a dictionary that we will fill with information from an input file.




Figure 1. *N is for Neville who died of ennui.*

In this exercise, we're going to look up a line of text from an input file that starts with the letter provided by the user. The text will come from an input file which will default to Edward Gorey's "The Gashlycrumb Tinies," an abecedarian book that describes various and ghastly ways in which children die. Instead of "A is for artichoke, B is for blackberry," we get:

“*A is for Amy who fell down the stairs. B is for Basil assaulted by bears.*

Our `gashlycrumb.py` program will take a letter of the alphabet as a positional argument and will look up the line of text from an *optional* input file that starts with that letter. The input file will have each letter on a separate line:

```
$ head -2 gashlycrumb.txt
A is for Amy who fell down the stairs.
B is for Basil assaulted by bears.
```

When our unfortunate user runs our program, here is what they would see. Note that we will consider the letter in a *case-insensitive* fashion:

```
$ ./gashlycrumb.py a
A is for Amy who fell down the stairs.
```

In this exercise, you will:


- Accept a single positional argument we'll call `letter`
- Accept an optional `--file` argument which must be a file. The default value will be `'gashlycrumb.txt'` (provided).
- Read the file, find the first letter of each line, and build a data structure that associates that letter to the line of text.
- Check if the given `letter` is a present in the given `file`.
- Print the line of text for the `letter` if present or an error if it isn't.
- Learn how to "pretty print" a data structure.

You can draw from several previous programs:

- From the "Word Count" program, you know how to take a file input and read it line-by-line.
- From the "Article" program, you know how to get the first letter of a bit of text.

- From the "Jump The Five" program, you know how to build a dictionary and lookup a value.

Now you'll put all those skills together to recite morbid poetry!



Writing gashlycrumb.py

Before you begin writing, I would encourage you to run the tests with `make test` or `pytest -xv test.py` in the `gashlycrumb` directory. The first test should fail:

```
test.py::test_exists FAILED
```

This is just a reminder that the first thing to do is to create the file called `gashlycrumb.py`. You can do this however you like, maybe by running `new.py gashlycrumb.py` in the `gashlycrumb` directory or by copying the `template/template.py` file, or by just starting a new file from scratch.

Run your tests again and you should pass the first and possibly the second tests if your program produces a usage statement. Next let's get the arguments straight. Modify your program's parameters in the `get_args` function so that it will produce the following usage when the program is run with *no arguments* or with the `-h` or `--help` flags:

```
$ ./gashlycrumb.py -h
usage: gashlycrumb.py [-h] [-f str] str

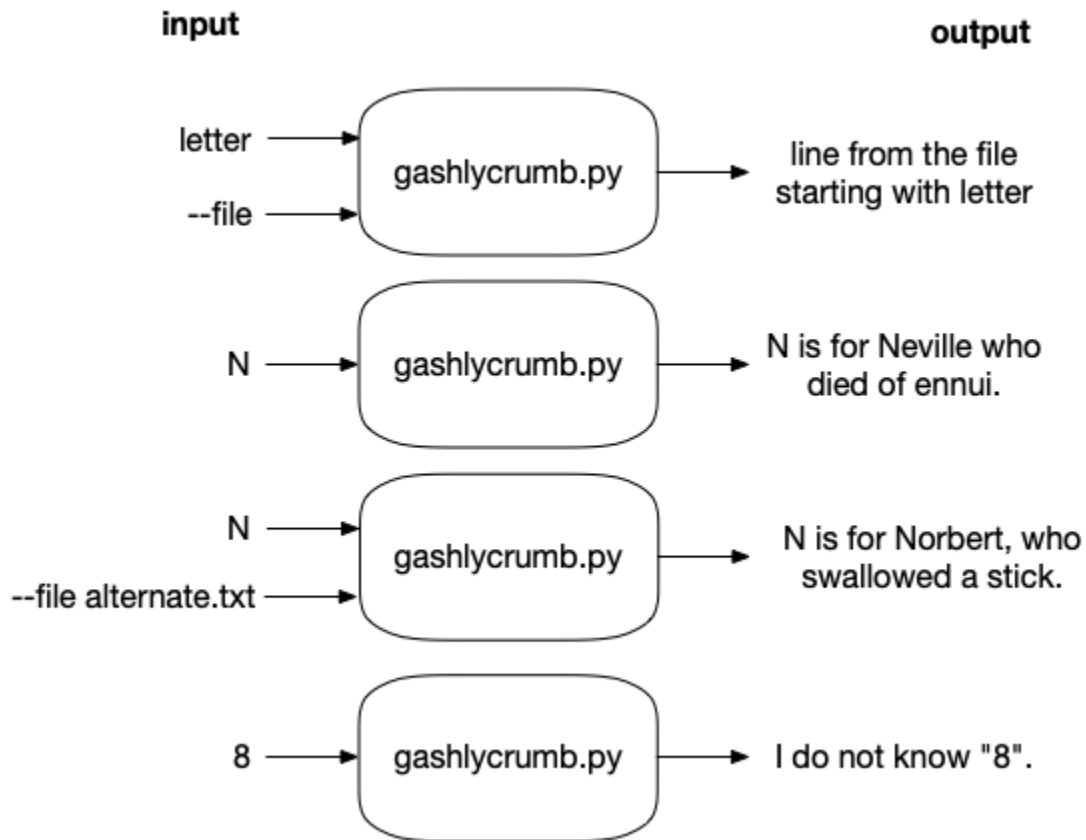
Gashlycrumb

positional arguments:
  str                Letter 1

optional arguments:
  -h, --help          show this help message and exit.      2
  -f str, --file str  Input file (default: gashlycrumb.txt)  3
```

- 1 The "letter" is a required positional argument.
- 2 The `-h` and `--help` arguments are created automatically by `argparse`.
- 3 The `-f` or `--file` argument is an option with a default value of "gashlycrumb.txt"

Here is a string diagram showing how the program will work:



Once you have the correct usage, start off by echoing the letter :

```
def main():
    args = get_args()
    print(args.letter)
```

Try running it to make sure it works:

```
$ ./gashlycrumb.py a
a
```

Next read the file line-by-line using a for loop. Note that I'm using `end=''` with the `print` so that it won't print the newline that's already attached to each line of the file :

```
def main():
    args = get_args()
    print(args.letter)

    for line in args.file:
        print(line, end='')

```

Try running it to ensure you can read the input file :

```
$ ./gashlycrumb.py a | head -3
a
A is for Amy who fell down the stairs.
B is for Basil assaulted by bears.
```

Use the "alternate.txt" file, too:

```
$ ./gashlycrumb.py a -f alternate.txt | head -3
a
A is for Alfred, poisoned to death.
B is for Bertrand, consumed by meth.
```

If provided a `--file` argument that does not exist, your program should exit with an error and message. Note that, if you use `type=argparse.FileType('r')`, this error should be produced automatically by `argparse`:

```
$ ./gashlycrumb.py -f blargh b
usage: gashlycrumb.py [-h] [-f str] str
gashlycrumb.py: error: argument -f/--file: can't open 'blargh': \
[Errno 2] No such file or directory: 'blargh'
```

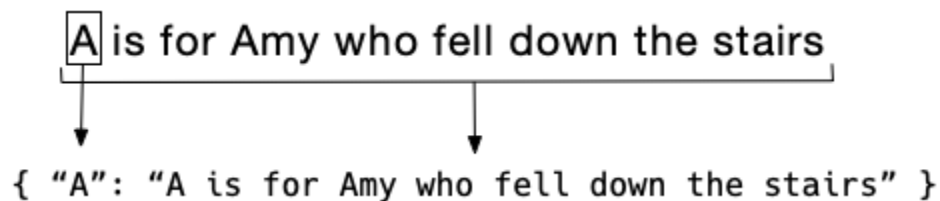
Now think about how you can use the first letter of each line to create an entry into a `dict`. Use the `print` command to look at your dictionary. Figure out how to check if the given letter is *in* (wink, wink, nudge, nudge) your dictionary. If given a value that does not exist in the list of first characters on the lines from the input file (when searched without regard to case), you should print a message:

```
$ ./gashlycrumb.py 3
I do not know "3".
$ ./gashlycrumb.py CH
I do not know "CH".
```



If the given letter is in the dictionary, print the value for it:

```
$ ./gashlycrumb.py a
A is for Amy who fell down the
stairs.
$ ./gashlycrumb.py z
Z is for Zillah who drank too much
gin.
```



Run the test suite to ensure your program meets all the requirements. Read the errors closely and fix your program.

Hints:

- Start with `new.py` and remove everything but the positional and optional `argparse.FileType('r')` parameters.
- A dictionary is a natural data structure that you can use to associate some value like the letter "A" to some phrase like "A is for Amy who fell down the stairs." Create a new, empty `dict`.
- Once you have an open file handle, you can read the file line-by-line with a `for` loop.
- Each line of text is a string. How can you get the first character of a string?
- Using that first character, how can you set the value of a `dict` to be the key and the line itself to be the value?
- Once you have constructed the dictionary of letters to lines, how can you check that the user's `letter` argument is *in* the dictionary?
- Can you solve this without a `dict`?

No skipping ahead to the solution until you have written your own version! If you peek, you might be die a horrible death stamped by kittens.

Solution

```

1  #!/usr/bin/env python3
2  """Lookup tables"""
3
4  import argparse
5
6
7  # -----
8  def get_args():
9      """get command-line arguments"""
10
11     parser = argparse.ArgumentParser(
12         description='Gashlycrumb',
13         formatter_class=argparse.ArgumentDefaultsHelpFormatter)
14
15     parser.add_argument('letter', help='Letter', metavar='str', type=str) 1
16
17     parser.add_argument('-f', 2
18                         '--file',
19                         help='Input file',
20                         metavar='str',
21                         type=argparse.FileType('r'),
22                         default='gashlycrumb.txt')
23
24     return parser.parse_args()
25
26
27  # -----
28  def main():
29      """Make a jazz noise here"""
30
31     args = get_args()
32     letter = args.letter 3
33     lookup = {line[0].upper(): line.rstrip() for line in args.file} 4
34
35     if letter.upper() in lookup: 5
36         print(lookup[letter.upper()]) 6
37     else:
38         print(f'I do not know "{letter}"') 7
39
40
41  # -----
42  if __name__ == '__main__':
43     main()

```

- 1 The required argument (a letter).
- 2 An optional `--file` that must be a readable file, if provided. Defaults to a known value.
- 3 Copy the `args.letter` value into the variable `letter`.
- 4 Build the `lookup` using a dictionary comprehension that reads the given file. Use the `upper` function to disregard case.
- 5 See if the `letter` argument is in the `lookup` dictionary, checking the `upper` value to disregard case.
- 6 If so, print the line of text from the `lookup` for the `letter`.

- 7 Otherwise, print a message that the `letter` is unknown.

Discussion

Did the frightful paws of the kittens hurt much? Let's talk about how I solved this problem. Remember, mine is just one of many possible solutions.

Handling the arguments

I prefer to have all the logic for parsing and validating the command-line arguments in the `get_args` function. In particular, `argparse` can do a fine job verifying tedious things such as an argument being an existing, readable file which is why I use `type=argparse.FileType('r')` for that argument. If the user doesn't supply a valid argument, then `argparse` will throw an error, printing a helpful message along with the short usage and exiting with an error code.

By the time I get to the line `args = get_args()`, I know that I have a valid, open file handle in the `args.file` slot. In the REPL, I can use `open` to get a file handle which I like to usually like to call `fh`. For copyright purposes, I'll use my alternate text:

```
>>> fh = open('alternate.txt')
```

Reading the input file

I know that I want to use a dictionary where the keys are the first letters of each line and the values are the lines themselves. That means I need to start by creating a new, empty dictionary either by using the `dict()` function or by setting a variable equal to an empty set of curly braces (`{}`). I'll call my variable `lookup`:

```
>>> lookup = {}
```

I will use a `for` loop to read each `line` of text. From Crow's Nest, I know I can use `line[0].upper()` to get the first letter of `line` and uppercase it. I can use that as the key into `lookup`. Each `line` of text ends with a newline that I'd like to remove. I can use the `str.rstrip` method to strip whitespace from the right side of the `line` (`rstrip` = *right strip*). The result of that will be the value for my `lookup`:

```
for line in fh:
    lookup[line[0].upper()] = line.rstrip()
```

I'd like to look at the resulting `lookup` dictionary. I can `print` it from the program or type `lookup` in the REPL, but it's going to be hard to read. I encourage you to try it. Luckily there is a lovely module called `pprint` to "pretty print" data structures. Here is how I can import the `pprint` function from the `pprint` module with the alias `pp`:

```
>>> from pprint import pprint as pp
```

Now let's take a peek at `lookup`:

```
from pprint import pprint as pp
```

module function alias

```
>>> pp(lookup)
{'A': 'A is for Alfred, poisoned to death.',
 'B': 'B is for Bertrand, consumed by meth.',
 'C': 'C is for Cornell, who ate some glass.',
 'D': 'D is for Donald, who died from gas.',
 'E': 'E is for Edward, hanged by the neck.',
 'F': 'F is for Freddy, crushed in a wreck.',
 'G': 'G is for Geoffrey, who slit his wrist.',
 'H': 'H is for Henry, who's neck got a twist.',
 'I': 'I is for Ingrid, who tripped down a stair.',
 'J': 'J is for Jered, who fell off a chair,',
 'K': 'K is for Kevin, bit by a snake.',
 'L': 'L is for Lauryl, impaled on a stake.',
 'M': 'M is for Moira, hit by a brick.',
 'N': 'N is for Norbert, who swallowed a stick.',
 'O': 'O is for Orville, who fell in a canyon,',
 'P': 'P is for Paul, strangled by his banyan,',
 'Q': 'Q is for Quintanna, flayed in the night,',
 'R': 'R is for Robert, who died of spite,',
 'S': 'S is for Susan, stung by a jelly,',
 'T': 'T is for Terrange, kicked in the belly,',
 'U': 'U is for Uma, who's life was vanquished,',
 'V': 'V is for Victor, consumed by anguish,',
 'W': 'W is for Walter, who's socks were too long,',
 'X': 'X is for Xavier, stuck through with a prong,',
 'Y': 'Y is for Yoeman, too fat by a piece,',
 'Z': 'Z is for Zora, smothered by a fleece.'}
```

Hey, that looks like a handy data structure. Hooray for us!

Looping with `for` versus a list comprehensions

We used three lines of code to build our `lookup` dictionary. We can actually accomplish that in *one* line by using a "dictionary comprehension." We've created list comprehensions by sticking a `for` inside brackets `[]`. A dictionary comprehension is the same but the `for` loop is inside curlyies `{}`.

```
lookup = {}
for line in fh:
    lookup[line[0].upper()] = line.rstrip()

lookup = { line[0].upper(): line.rstrip() for line in fh }
```


If you are following along by pasting code into the REPL, note that we have exhausted the file handle `fh` just above by reading it. (Refer back to the "Howler" to read about file handles.) I need to `open` it again for this next bit:

```
>>> fh = open('gashlycrumb.txt')
>>> lookup = { line[0].upper(): line.rstrip() for line in fh }
```

If you print it again, you should see the same output as above. It may seem like showing off to write one line of code instead of three, but it really does make a good deal of sense to write compact, idiomatic code. More code always means more chances for bugs, so I usually try to write code that is as simple as possible (but no simpler).

Dictionary lookups

Now that I have a `lookup`, I can ask if some value is `in` the keys. Note that I know the letters are in uppercase and I assume the user could give me lower, so I use `letter.upper()` to only compare that case:



```
>>> letter = 'a'
>>> letter.upper() in lookup
True
>>> lookup[letter.upper()]
'A is for Amy who fell down the stairs.'
```

If the letter is found, I can print the line of text for that letter; otherwise, I can print the message that I don't know that letter:

```
>>> letter = '4'
>>> if letter.upper() in lookup:
...     print(lookup[letter.upper()])
... else:
...     print('I do not know "{}".'.format(letter))
...
I do not know "4".
```

An even shorter way to write that would use the `dict.get` method:

```
def main():
    args = get_args()
    letter = args.letter
    lookup = {line[0].upper(): line.rstrip() for line in args.file}
    print(lookup.get(letter.upper(), f'I do not know "{letter}").')) 1
```

¹ If `lookup.get` will return the value for `letter.upper()` or the phrase "I don't know..."

Solving without using a dictionary

I think the dictionary solution is rather elegant, but it's certainly not the only way to solve this. Given the size of the input file, it might actually be a poor solution. If you were searching for a needle in a haystack, you would stop as soon as you found the needle. In our dictionary solution, we're reading the entire file into a `dict` and only then looking to see if the `letter` was present. In this solution, we check each line as we read it and stop as soon as we find our letter. Note that this solution needs to `import sys`:

```
def main():
    args = get_args()
    letter = args.letter

    for line in args.file:
        if line[0].upper() == letter.upper():
            print(line, end='')
            sys.exit(0)

    print(f'I do not know "{letter}").') 5
```

¹ Read the file line-by-line.

- 2 Check if the uppercased first character of the `line` is equal to the uppercased `letter`.
- 3 If it is, print the `line`. Use `end=''` because `print` wants to add a newline, but the `line` already has one.
- 4 Exit the program.
- 5 If we make it to this line, then we never exited the program which means we never found our letter, so let the user know that their choice was missing.

Note that this is another "low-memory" solution because we read the file line-by-line whereas the `dict` version stores *all* the lines of the input file. For such a small file as this, it's no concern, but beware if you need to process a file that is very large.

Review

- A dictionary comprehension is a way to build a dictionary in a one-line `for` loop.
- Defining file input arguments using `argparse.FileType` saves you time and code.
- Looking up a key in a `dict` is faster than searching for an element in a `list`. The time to search a `list` is proportional to its length while dictionary keys are stored in a way optimized for finding them.
- Python's `pprint` module is used for "pretty printing" complex data structures.

Going Further

- Write an interactive version that takes input directly from the user. Use `while True` to set up an infinite loop and keep using `input` to get the user's next `letter`:

```
$ ./gashlycrumb_interactive.py
Please provide a letter [! to quit]: t
T is for Titus who flew into bits.
Please provide a letter [! to quit]: 7
I do not know "7".
Please provide a letter [! to quit]: !
Bye
```

Last updated 2020-01-16 21:14:08 -0700