

Tiny Python Projects

Learn coding and testing with puzzles and games

Ken Youens-Clark

MEAP



MANNING



MEAP Edition
Manning Early Access Program
Tiny Python Projects
Learn coding and testing with puzzles and games
Version 2

Copyright 2020 Manning Publications

For more information on this and other Manning titles go to
manning.com

welcome

Hey, thanks for buying the MEAP edition of *Tiny Python Projects!* This is as close to an “interactive” book as I could make. You see, I want you to be actively involved in learning how to write Python really well by using a process called “test-driven development” (TDD). So, really, you’ll learn two things from this book: how to write Python, and how to use and write tests for your code.

The first two chapters describes how to structure a Python program into discrete functions, one of which is dedicated to getting command-line arguments and producing documentation for your program. The following chapters each describe a coding challenge that you can try your hand at writing code and the tests that have been provided in a Git repository so you can use the TDD process. Each chapter is dedicated to some central idea, like how to manipulate a string or list, when and how to use a dictionary, how to use random events, how to write functions and algorithms, etc. As the programs get more complex, I encourage you to write your own functions *and the tests for them* so that you learn how to become a self-sufficient tester.

I really look to hearing feedback from readers, suggestions for improvements, and extension exercise. Please be sure to post any questions, comments, or suggestions you have about the book in the [liveBook discussion forum](#).

—Ken Youens-Clark, Tucson, AZ

brief contents

Introduction

- 1 *How to write a Python program*
- 2 *Using argparse*
- 3 *The Crow's Nest: Working with strings*
- 4 *Going on a picnic: Working with lists*
- 5 *Jump the Five: Working with dictionaries*
- 6 *Howler: Working with files and STDOUT*
- 7 *Words Count: Reading files/ STDIN, iterating lists, formatting strings*
- 8 *Gashlycrumb: Looking items up in a dictionary*
- 9 *Apples and Bananas: Find and replace*
- 10 *Dial-A-Curse: Generating random insults from lists of words*
- 11 *Telephone: Randomly mutating strings*
- 12 *Bottles of Beer Song: Writing and testing functions*
- 13 *Ransom: Randomly captilizing text*
- 14 *Twelve Days of Christmas: Algorithm design*
- 15 *Rhymer: Using regular expressions to create rhyming words*
- 16 *The Kentucky Friar: More regular expressions*
- 17 *The Scrambler: Randomly reordering the middles of words*
- 18 *Mad Libs: Using regular expressions*
- 19 *Gematria: Numeric encoding of text using ASCII values*

20 Workout Of the Day (Parsing CSV file, creating text table output)

21 Password Strength

APPENDIXES:

A Setting up your development environment



Introduction

"Codes are a puzzle. A game, just like any other game." - Alan Turing

I believe you can learn serious things through playing. This is a book of programming exercises that explores programming ideas by coding puzzles and games. Each chapter includes a description of a program you should write with examples of how the program should work. Most importantly, each program includes tests so that you know if your program is working.

When you are done with this books you be able to:

- Write command-line Python programs
- Process a variety of command-line arguments, options, and flags
- Write and run tests for your programs and functions
- Manipulate of Python data structures including strings, lists, tuples, sets, dictionaries
- Use higher-order functions like `map` and `filter`
- Write and use regular expressions
- Read, parse, and write text files
- Use and control randomness

The most fundamental idea I teach is how to test your code. How can you be sure that a function or program does what you think it does? While I will introduce guidelines for writing good code as well as tools you can use to format your code and find errors, none of these is a replacement for *testing* your code. Each exercise includes a test suite that you will run repeatedly as you write your programs. In this way, I'm using ideas from the "test-driven development" where tests are written *first* and then the programs that satisfy those tests come afterwards.

In every book, you get a dose of the author's personal biases, and this one is no different. As a matter of personal preference, I avoid writing object-oriented code and instead encourage you to break problems into small, transparent functions. I believe this leads to code that is easier to teach and test. While Python itself is an object-oriented (OO) language, it has good support for a "functional" approach to development and testing. There is so much you can do with Python's basic data types and structures like strings, numbers, lists, tuples, dictionaries. I personally never write in an OO style, preferring instead to adopt ideas from functional programming (FP) languages like Haskell.

The programming techniques in each exercise are not specific to Python. Most every language has variables, loops, functions, and more languages are incorporating FP ideas from Java to JavaScript. After you write your solutions in Python, I would encourage you to write solutions in every other language you know and compare what parts of a different language make it easier or harder to write your programs. If your programs support the same command-line options, you can even use the included tests to verify those programs!

Why Write Python?

Python is an excellent, general-purpose programming language. You can use it to write command-line programs, web servers, adventure games, and business applications. There are Python modules to help you wrangle complex scientific data, explore machine learning algorithms, and generate publication-ready graphics.

Many college-level computer science programs have moved away from languages like C and Java to Python as their introductory language because Python is a relatively easy language to learn. Often Python codes reads like English — the statement `for item in cart` is a way to visit each `item` in the collection called `cart`. While languages like C, Java, Perl, and PHP use semi-colons (`;`) to mark the ends of statements, Python uses a newline (when you hit the Enter key). These same languages also tend to group statements into "blocks" by enclosing them in curly brackets (`{}`), but Python groups statements by how many spaces they are indented (usually four). Many people feel that Python's syntax and lack of excessive punctuation makes it more readable than other languages.

I believe you can teach some very fundamental and powerful ideas from computer science using Python. As I show you ideas like regular expressions and higher-order functions, I hope it piques your interest in other languages and encourages you to study further.

Who Am I?

My name is Ken Youens-Clark. I work as a Senior Scientific Programmer at the University of Arizona. Most of my career has been spent working in bioinformatics using computer science ideas to study biological data. I began my undergraduate degree as a Jazz Studies on the drumset at the University of North Texas in 1990. I changed my major a few times and eventually ended up with a BA in English literature

in 1995. I didn't really have a plan for my career, but I did like computers. Around 1995, I started tinkering with databases and HTML at first job out of college, building the company's mailing list and first website. I was definitely hooked! After that, I managed to learn VisualBasic on Windows 3.1 and, through the next few years, I programmed in several languages and companies before landing in a bioinformatics group at Cold Spring Harbor Laboratory in 2001 led by Lincoln Stein, an early advocate for open software, data, and science. In 2015 I moved to Tucson, AZ, to work at the University of Arizona where I finished my MS in Biosystems Engineering in 2019.

I was a Perl hacker for longest portion of my coding career, from 1998 to around 2016. Perl ruled the early days of the Internet and the world of genomics and bioinformatics. Over time I've also worked with bash, Python, JavaScript, Haskell, Rust, Elm, Prolog, Ruby, R, Julia, SPSS, and anything else that seems interesting. The programming style that I show blends ideas from many of these languages, and Perl's motto of "There Is More Than One Way To Do It" (TIMTOWTDI) is perhaps why I like to show so many different ways to solve problems.

When I'm not coding, I like playing music, riding bikes, cooking, reading, and being with my wife and children.

Who Are You?

"You don't know the power of the command line!" — Darth Vader (probably)

I think my ideal reader is someone who's been trying to learn to code well but isn't quite sure how to level up. Perhaps you are someone who's been playing with Python or some other language that has a similar syntax like Java(Script) or Perl. Or maybe you've cut your teeth on something like Haskell or Scheme and you're wondering how does it go in Python Land? You are looking for interesting challenges with enough structure to help you know when you're moving in the right direction. Maybe you're interested in "test-driven development" and want to know how to write and use tests?

This is a book that will try to teach you well-structured, documented, testable code in Python. The material introduces best-practices from industry such as test-driven development. It encourages you to read documentation and Python Enhancement Proposals (PEPs) to write idiomatic code that other Python programmers would immediately recognize and understand. This book shows you why you'd want to learn about regular expressions and algorithm design and statistics and random events. This book uses simple puzzles and games you already understand and asks you to teach the rules to Python.

This is probably not an ideal book for the absolute beginning programmer. I assume no prior knowledge of the Python language specifically because I'm thinking of someone who is coming from another language. If you've never written a program in *any* language at all, you might do well to come back to this material when you are comfortable with ideas like variables and loops.

Something that makes this book very different is that we write only *parameterized*,

command-line programs because they are much easier to test than, say, interactive programs, web servers, or Jupyter Notebooks. This means you'll will need to access to a command-line interface. Not to fear! Programs like Microsoft's VSCode or Anaconda's Spyder can help! I hope that you'll learn the power of the command.

Why Did I Write This Book?

"The only way to learn a new programming language is by writing programs in it." - Dennis Ritchie

Over the years, I've had many opportunities to help people learn programming, and I always find it rewarding. The structure of this book comes from my own experience in the classroom where I think formal specifications and tests can be useful aids in learning how to break a program into smaller problems that need to be solved to create the whole program.

The biggest barrier to entry I've found when I'm learning a new language is that small concepts of the language are usually presented outside of any useful context. Yes, we all love to write "HELLO, WORLD!", but after that, I usually struggle to write a complete program that will accept some arguments and do something *useful*. For instance, how can I get the name of a file from the user, then read the file and do something with the contents, all the while handling the various errors that arise like not getting an argument from the user, the argument isn't actually a file, the file isn't readable, etc.? This book teaches you exactly how to write Python that accepts and validates arguments, presents usage messages, handles errors, and runs to completion.

More than anything, I think you need to practice. It's like the old joke: "What's the way to Carnegie Hall? Practice, practice, practice." These coding challenges are short enough that you could probably finish each in a few hours to days. This is more material than I could work through in a semester-long university-level class, so I imagine the whole book would take you several months, perhaps even a year or more. I hope you will solve the problems, then think about them and come back later to see if you can solve them differently, maybe using a more advanced technique or making them run faster.

Using test-driven development

"Test-driven development" is described by Kent Beck in his 2002 book by that title as a method to create more reliable programs. The basic idea is that we write tests even before we write code. We run the tests and verify that our code fails. Then we write the code to make each test pass. We always run all the tests so that, as we fix new tests, we ensure we don't break tests that were passing before. When all the tests pass, we have at least some assurances that the code we've written conforms to some manner of specification.

Each program you are asked to write comes with tests that will tell you when the code is working acceptably. The first test in every exercise is whether the expected program exists. The second test checks if the program returns something like a "usage"

statement when run with `-h` or `--help`.

After that, your program will be exercised with various inputs and options. You'll fail a test, and fix it. Then run the tests again. You will probably fail the next test. Fix that test such that you don't break the previously passing one. Keep fixing each test until they all pass. Then you are done.

It doesn't matter if you solved the problem the same way as in the solution I provide. All that matters is that you solve it *on your own!*

Suggestions for writing

As a general rule, I try to write many very small functions and lots of tests. Every function should be as short as possible, but no shorter — generally 50 lines is my upper limit on function length. I have no lower bound, though. Even if a function is just one line of code, I like to give it a name and some tests to ensure it does what I think. This is often called "unit testing" in the world of software engineering, and it's a very good practice. As the challenges get harder, I start suggesting specific functions and tests you should write in your programs. If we know the smaller bits work individually, then we ought to have confidence that they will work in concert.

This is the basic idea of "compositionality" where we try to compose large systems from smaller ones, so I also write tests to check that my programs as a whole do what they should. This is called "integration testing" and is found in the `test.py` programs I've provided for you in each of these exercises where your programs are run with various options and checked that they produce the expected output.

So, in short, always start a program the same way. Always process the arguments the same way. Always have the same structure to your program (e.g., start with `main`). When you think you need a function in your program, write the tests first and then write the code that will pass the tests. Then integrate your function into the greater whole. Write a test suite outside your program that checks that it fails properly as well as works properly. Whenever you make a change to your program, run your test suite to see if you accidentally broke something that used to work!

Lastly, you should always use a source code manager. If you fork and clone the GitHub repo of the exercises, then you've made an important first step in learning how to keep track of your code. You don't have to use `git` or GitHub. Mercurial is fine, too. (It's also written in Python!) Find a tool that works for you, and commit to tracking your changes.

Setting up your environment

To write the programs and run the tests, you will need to install Python and have some sort of command-line access. Windows users will probably need to install Windows Subsystem for Linux or Cygwin in order to get a command line interface suitable for running the programs and the test suite. On a Mac, the default Terminal app is sufficient.

I wrote and tested the programs with the Python version 3.7 but should work with any version 3.6 or newer. Python 2 reached end-of-life at the end of 2019 and should no longer be used. To see what version of Python you have installed, type `python3 --version`. If you need to install Python, see www.python.org/downloads/.

You will need to use a *text editor* to write code — something that writes *plain text*. Something like Microsoft Word will not work. You author prefers vim, but you might like to try Notepad, Sublime, TextWrangler, or Textmate. Some people like to use an integrated development environment (IDE) that combines text editors with tools to help inspect, debug, and run code. Examples include Xcode, PyCharm, Anaconda Spyder, or Microsoft VSCode. Whatever works for you!

I demonstrate how to run the programs and tests from the command line. Each exercise includes a `Makefile` that allows you to type `make test` in the directory to run the tests. This is a shortcut to execute `pytest -xv test.py` which is how the tests are actually run. If you do not have or if you do not wish to use `make`, you can execute this command yourself to run the tests.

The command line and the Python REPL

I will show commands you can execute on the command line preceded with the \$ (dollar sign) and using a **fixed-width font**. The \$ is a common command-line prompt, but your command line may differ. For instance, I might use the command `cat` (for "concatenate") to show you the contents of a file:

```
$ cat inputs/spiders.txt
Don't worry, spiders,
I keep house
casually.
```

If you want to run that command, *do not copy* the leading \$, only the text that follows, otherwise you'll probably get an error like \$: command not found.

Python has a really excellent tool that allows you to interact directly with the language to try out ideas. If you type `python3` on the command line, you will enter a REPL (read-evaluate-print-loop, often pronounced like "repple" in a way that sort of rhymes with "pebble"). These examples will be also shown using a **fixed-width font** with the text preceded by the prompt >>>. Here is what it looks like on my system when I start the Python REPL:

```
$ python3
Python 3.7.3 (default, Mar 27 2019, 16:54:48)
[Clang 4.0.1 (tags/RELEASE_401/final)] :: Anaconda, Inc. on darwin
Type "help", "copyright", "credits" or "license" for more information.
>>>
```

You can type Python statements like `x = 10` to assign the value 10 to the variable `x`:

```
>>> x = 10
```

As with the command-line prompt \$, do not copy the leading >>> or Python will

complain:

```
>>> >>> x = 10
      File "<stdin>", line 1
        >>> x = 10
          ^
SyntaxError: invalid syntax
```

The ipython REPL has a magical %paste mode that removes the leading >>> prompts so that you can copy and paste all the code examples:

```
$ ipython
Python 3.7.3 (default, Mar 27 2019, 16:54:48)
Type 'copyright', 'credits' or 'license' for more information
IPython 7.11.1 -- An enhanced Interactive Python. Type '?' for help.

In [1]: >>> x = 10

In [2]: x
Out[2]: 10
```

You might also enjoy using Jupyter Notebooks as a REPL. Whichever way you choose to interact with Python, I suggest you manually type all the code yourself as this builds muscle memory and forces you to interact with the syntax of the language.

Code formatting and linting

Every program included has been automatically formatted with yapf (Yet Another Python Formatter, github.com/google/yapf), a tool from Google that can be customized with a configuration file. Another popular formatter is black (github.com/psf/black). I encourage you to adopt and regularly use a formatter *after every modification to your program*.

I would also encourage you to look at code "linters" like pylint (www.pylint.org/) and flake8 (flake8.pycqa.org/en/latest/) to find potential errors in your code that the Python interpreter itself will not complain about. The mypy tool (mypy-lang.org/) will also be very helpful as we introduce type hints.

You can install all of these tools using Python's pip module:

```
$ python3 -m pip install yapf black pylint flake8 mypy
```

Getting the code

All the tests and solutions are available at github.com/kyclark/tiny_python_projects. If you like, you can simply git clone that repo and work out the solutions.

If you would like to make a copy of the code, you can "fork" the repository. This would allow you to commit and push your solutions to your own repo. To fork the repo:

1. Create an account on GitHub.com
2. Go to github.com/kyclark/tiny_python_projects

- Click the "Fork" button to make a copy of the repository into your account.

I may update the repo on occasion. If you would like to be able to pull my changes, you should set mine as an "upstream" repo. Inside your local Git repo, execute this command:

```
$ git remote add upstream https://github.com/kyclark/tiny_python_projects.git
```

After that, use `git pull upstream master` to get updates.

Starting new programs with new.py or template.py

I wrote a program called `new.py` that will help you create new Python programs with boilerplate code that will be expected of all the programs. That is, every program is expected to respond to the `-h` and `--help` flags and produce documentation on program parameters. The generated program is over 70 lines of a well-structured and documented program that you can modify to suite your needs.

You will find `new.py` the `bin` directory of the GitHub repository. The program expects a single argument which is the name of the new program to create. It will create the new program in the directory where you execute the command. The programs you write need to exist in the same directory as their tests, so when you are ready to work on the `crowsnest` exercise you should first change to that directory:

```
$ cd ~/tiny_python_projects/crowsnest
```

You can then execute the `new.py` using an absolute path to the program:

```
$ ~/tiny_python_projects/bin/new.py crowsnest.py
Done, see new script "crowsnest.py."
```

Or you can use a *relative* path from the current directory:

```
$ ../bin/new.py crowsnest.py
Done, see new script "crowsnest.py."
```

Alternately, you may need to invoke `python3` directly:

```
$ python3 ../bin/new.py crowsnest.py
Done, see new script "crowsnest.py."
```

If you wish, you can copy `new.py` to a location in your `$PATH`, or you might prefer to alter your `$PATH` variable to include the directory where `new.py` lives.

If you don't want to use `new.py`, you can also copy `template/template.py` to start your new programs:

```
$ cd ~/tiny_python_projects/crowsnest
$ cp ../template/template.py crowsnest.py
```

Both methods are meant to save you time and make it easier to write programs that use `argparse` to handle the program's parameters.

Why Not Notebooks?

Notebooks are great for an interactive and visual exploration of data, but the downsides:

- Stored as JSON not line-oriented text, so no good diff tools
- Not easily shared
- Too easy to run cells out of order
- Hard to test
- No way to pass in arguments

I believe you can better learn how to create testable, *reproducible* software by writing command-line programs that always run from beginning to end and have a test suite. It's difficult to achieve that with Notebooks, but I do encourage you to explore Notebooks on your own.

A Note about the lingo

Often in programming books you will see "foobar" used in examples. The word has no real meaning, but its origin probably comes from the military acronym "FUBAR" (Fouled Up Beyond All Recognition). If I use "foobar" in an example, it's because I don't want to talk about any specific thing in the universe, just the idea of a string of characters. If I need a list of items, usually the first item will be "foo," the next will be "bar." After that, convention uses "baz" and "quux," again because they mean nothing at all. Don't get hung up on "foobar." It's just a shorthand placeholder for something that could be more interesting later.

We also tend to call errors code "bugs." This comes from the days of computing before the invention of transistors. Early machines used vacuum tubes, and the heat from the machines would attract actual bugs like moths that could cause short circuits. The "operators" (the people running the machines) would have to hunt through the machinery to find and remove the bugs, hence the term to "debug."

1

How to write a Python program

Before we start writing the exercises, I want to discuss how to write programs that are documented and tested using the following principles. Specifically, we're going to:

- Process and document command-line arguments using argparse
- Write and run tests for your code with pytest
- Use tools like yapf or black to format your code
- Annotate variables and functions with type hints and then use mypy to check correctness
- Use tools like flake8 and pylint to find problems in your code

1.1 How to write the world's best "Hello, World!"

It's pretty common to write "Hello, World!" as your first program in any language, so let's start there. We're going to start as simply as possible and end with a program you'd be proud to submit for a code review.

In your repo, go into the `hello` directory. There you'll see 14 versions of the "hello" program we'll write along with the `test.py` we'll use to test the final version. Start off by creating a text file called `hello.py` in that directory and add this line:

```
print('Hello, World!')
```

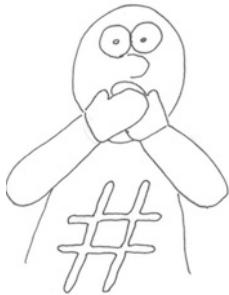
We can run it with the command `python3 hello.py` to have Python version 3 execute the commands in the file called `hello.py`. You should see this:

```
$ python3 hello.py
Hello, World!
```



If that was your first Python program, congratulations! You've done well, and we're all really proud of you.

1.1.1 Comment lines



In Python, the `#` character and anything following it is ignored by Python. This is useful to add comments to your code or to temporarily disable lines of code when testing and debugging. It's always a good idea to document your programs with the purpose of the program and/or the author's name and email address. We can use a comment for that:

```
# Purpose: Say hello
print('Hello, World!')
```

If you run it again, you should see the same output as before because the "Purpose" line is ignored. Note that any text to the left of the `#` is executed, so you can add a comment to the end of a line, if you like.

1.1.2 The shebang

We could write similar programs in other languages like bash, Ruby, and Perl, so it's common to document the language inside the program with a special comment character combination of `#!`. The nickname for this is "shebang" (pronounced "shuh-bang" — I always think of the `#` as "shuh" and the `!` as the "bang!"). As usual, Python will ignore this, but the OS will use it to decide which program to use to interpret the rest of the file.

Here is the shebang you should add:

```
#!/usr/bin/env python3
```

The `env` program will tell you about your "environment." If you run it like `env`, you should see lots of lines of output with lines like `USER=kyclark`. If you run it like `env python3`, it will search for and run the `python3` program. On my systems, I like to use the Anaconda Python distribution, and this is usually installed in my "home" directory. On my Mac laptop, this means that my `python3` is actually `/Users/kyclark/anaconda3/bin/python3`, but on another system like one of my Linux web servers will be an entirely different location. If I hardcoded the Mac location of `python3` in the shebang line, then my program won't work when I run it on Linux, so I always use the `env` program to find `python3`.

Now your program should look like this:

```
#!/usr/bin/env python3 ①
# Purpose: Say hello ②
print('Hello, World!') ③
```

- ① The shebang line telling the operating system to use the `/usr/bin/env` program to find `python3` to interpret this program.
- ② A comment line documenting the purpose of the program.
- ③ A Python command to print some text to the screen.

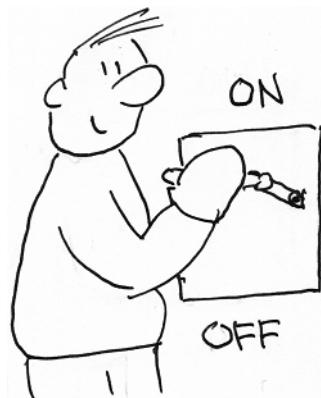
1.1.3 Making a program executable

Now I'd like to "run" the script. On a Unix-like system, we can make any file "executable" with the command `chmod` (*change mode*). Think of it like turning your program "on." Run this command to make `hello.py` executable — to add (+) the *executable* (x) bit:

```
$ chmod +x hello.py
```

Now you can run the program like so:

```
$ ./hello.py
Hello, World!
```



That looks way cooler.

1.1.4 String diagrams

Throughout the book, I'll use "string diagrams" to visualize the inputs and outputs of the programs we'll write. If we create one for our program as it is, there are no inputs and the output is always "Hello, World!"



1.1.5 Adding a parameter

It's not terribly interesting for our program to always say "Hello, World!" Let's have it enthusiastically greet some name that we will pass as an *argument*. That is, we'd like our program to work like this:

```
$ ./hello.py Terra
Hello, Terra!
```

The arguments to a program are available through the Python `sys` (system) module's `argv` (argument vector). We have to add `import sys` to our program to use it. Change your program so it looks like the code below. Don't worry if you don't understand what you're typing — we'll cover all this in detail later!

```
1      #!/usr/bin/env python3
2      # Purpose: Say hello
3
4      import sys          ①
5
6      args = sys.argv[1:] ②
7      if args:            ③
8          name = args[0] ④
9          print(f'Hello, {name}!') ⑤
10     else:              ⑥
11         print('Hello, World!') ⑦
```

- ① Allows us to use code from the `sys` module.
- ② `argv` is a list, slice from index 1.
- ③ If there are any arguments,
- ④ Then get the name from index 0.
- ⑤ Print name using an f-string,
- ⑥ Otherwise,
- ⑦ Print the default greeting.

Line 6 probably looks a bit funny, so let's talk about that. The argument vector `argv` is a record of how the program was run starting with the path to the program itself and followed by any arguments. If you run it like `./hello.py Terra`, then `argv` will look like this:

command	<code>\$./hello.py Terra</code>
sys.argv	<code>['./hello.py', 'Terra']</code>

index



Python, like so many other programming languages, uses the *index 0* for the first

element in a list. The element at `argv[0]` will always be the path to the program itself.

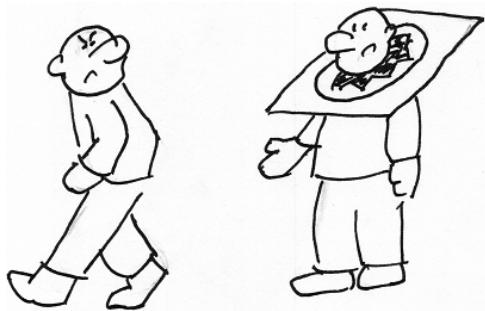
So on line 6, we use `sys.argv[1:]` to say we want a *slice* of that list starting from the index 1 and going to the end of the list and assign that to the variable called `args`. On line 7, we check if there is something in `args` — that is, there were some arguments to the program. If so, on line 8 we assign the variable `name` to the first element in `args`. If there are no `args`, we will print the default greeting on line 11. In the "Picnic" exercise, we'll talk much more about lists and slices.

Now our program will say "Hello, World!" when there are no arguments:

```
$ ./hello.py
Hello, World!
```

Or will extend warm salutations to whatever value is given as an argument:

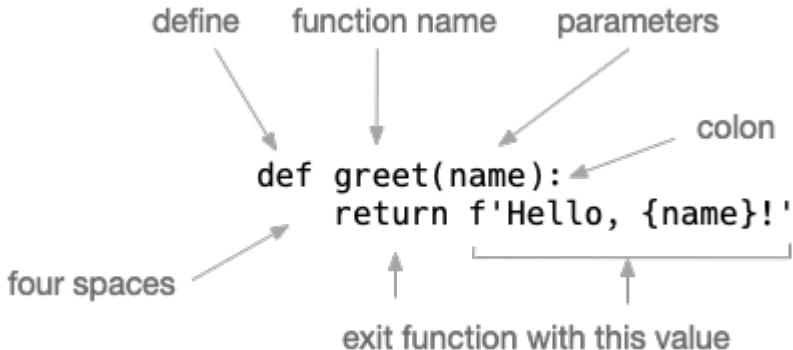
```
$ ./hello.py Universe
Hello, Universe!
```



1.1.6 Writing and testing functions

Our `hello.py` program now behaves differently depending on how it is run. How can we *test* our program to be sure it does the right thing? Let's find the part of our program that is variable and put that into a smaller unit of code that we can test. The unit will be a *function*, and the thing that changes is the greeting that we produce.

Create a function called `greet` with the name as a parameter.

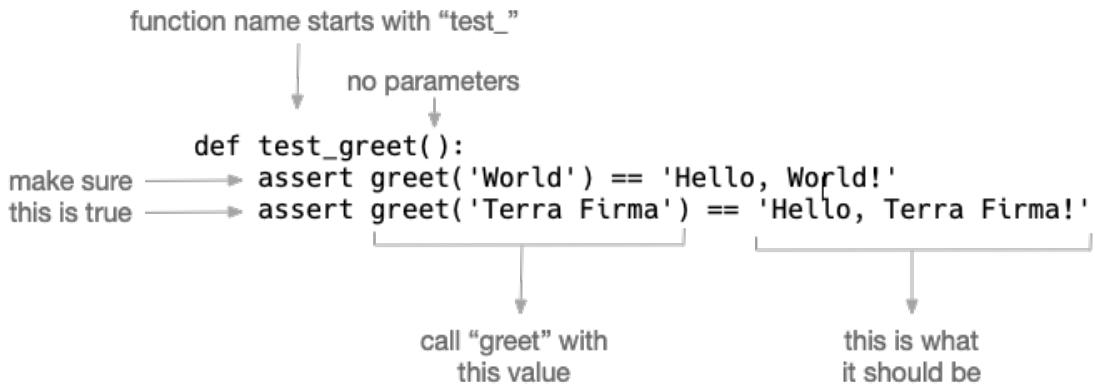


The `def` is to *define* a function, and the function name follows after that. The function's parameters (the things that can change when it is run) go inside the parentheses. Here we have just one, and it's called `name`. The colon (`:`) lets Python know that a "block" of statements will follow, all of which need to be indented to the same level. Indentation can be done with tabs (pressing the `Tab` key) or using spaces, but you are not allowed

to use a mix of tabs and spaces. Whatever you choose, be consistent. I use 4 spaces which seems to be in line with community standards.

The return will leave the function, optionally returning any value that follows. If no value follows the return, then the special value None will be returned to the caller. Some languages will return the last value of the function, but not Python. You must explicitly return what you want. You may return multiple values by separating them with commas, and you are also allowed to have more than one return in a function.

Notice that our new greet function does not print the greeting to the user. Its only job is to create the greeting so that we can write a test for it like so:



Here is what our whole program looks like now:

```

1  #!/usr/bin/env python3
2  # Purpose: Say hello
3
4  import sys
5
6  def greet(name):          ①
7      return f'Hello, {name}!'
8
9  def test_greet():          ②
10     assert greet('World') == 'Hello, World!'
11     assert greet('Terra Firma') == 'Hello, Terra Firma!'
12
13  args = sys.argv[1:]        ③
14  name = args[0] if args else 'World' ④
15  print(greet(name))        ⑤
  
```

- ① A function to create a greeting.
- ② A function to test the greet function.
- ③ The program starts here.
- ④ An if expression is a shorthand way to write if/else that fits on one line. The name will be set to args[0] if is something in args, otherwise use the value 'World'. We'll talk more about these in chapter 2 (Crow's Nest) when we are dealing with *binary* (two) choices.
- ⑤ The print and greet functions are *two separate ideas*. I know that I can count on Python to reliably print, but I need to test my greet function to ensure it works properly.

I personally like to use the `pytest` module to run my tests, but you may prefer to use the `unittest` module, especially if you have a background in Java. `pytest` will execute every function that has a name starting with `test_`, so I call my function `test_greet`. This function will call the `greet` function two times, once with the argument '`World`' and once with the argument '`Terra Firma`' to assert that the value returned from the function is what I would expect.

To see the tests in action, run `pytest -xv hello.py` (some output elided)

```
$ pytest -xv hello.py
=====
 test session starts =====
...
collected 1 item

hello.py::test_greet PASSED [100%]

===== 1 passed in 0.03 seconds =====
```

NOTE

The `-v` flag is to have `pytest` run in the "verbose" mode, and the `-x` flag to `pytest` tells it to stop at the first failing test. It's common to combine short flags in this way (`-xv` or `-vx`, the order doesn't matter). I include a `Makefile` with each exercise so that you can run `make test` to execute this for you. If you don't have `make` or don't want to use it, execute `pytest -xv test.py`.

1.1.7 Adding a main function

Our program is pretty good, but it's not quite idiomatic Python. Do you notice how we have two functions and then some other code just sort of hanging out (lines 13-15) flush left? It's good practice to put *all* your code inside functions, and it's very common to create a function called `main` where your program starts:

```
def main():
    args = sys.argv[1:]
    name = args[0] if args else 'World'
    print(greet(name))
```

We still have to tell Python to run this function, though, and the idiom for that in Python is to put these as the *last* two lines in your program:

```
if __name__ == '__main__':
    main()
```

Python reads your whole program from top to bottom. When it gets to these last lines, it will see if the special variable called `{dbl_name}{dbl_}` (one of the many special "double-under" or "dunder" variables) shows that we are in the "main" namespace. If you use `import` to bring your code into another piece of code, then it would *not* be in the "main" namespace. When you execute the code like a program, then the namespace will be "main" and so the `main` function will be run.

Here is our whole program now:

```
1  #!/usr/bin/env python3
2  # Purpose: Say hello
```

```

3      import sys
4
5      def greet(name):
6          return f'Hello, {name}!'
7
8      def test_greet():
9          assert greet('World') == 'Hello, World!'
10         assert greet('Terra Firma') == 'Hello, Terra Firma!'
11
12     def main():                      ①
13         args = sys.argv[1:]
14         name = args[0] if args else 'World'
15         print(greet(name))
16
17     if __name__ == '__main__':    ②
18         main()                  ③
19

```

- ① A new `main` entry function.
- ② Program starts here.
- ③ Call "main" function if in "main" namespace

Note that the call to run `main()` must occur in the file *after* the definition of the `main` function. Python will throw an exception if you try to call a function that hasn't yet been defined. Try moving these to the *top* of your program and running it again, and you'll see an error:

```
NameError: name 'main' is not defined
```

1.1.8 Why is testing important?



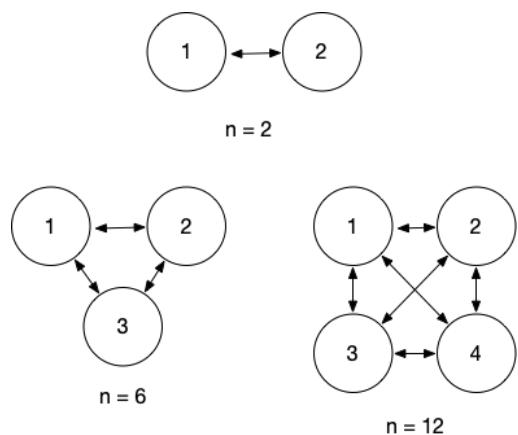
I love to bake, especially breads and cakes. A simple but effective test for cakes is to insert a toothpick in the middle. If the toothpick comes out with gooey batter stuck to it, then the cake is not done. If it comes out clean, it's ready. When I'm cooking bread, I usually insert a thermometer to measure the internal temperature. It's really important to tests like these so you don't serve undercooked food!¹

¹ In Monty Python and the Search for the Holy Grail they have a test to determine if a woman is a witch. Witches are burned, and wood also is burned, therefore witches burn because they are made of wood. Wood floats in water, and so do ducks. "Logically," if a woman weighs the same as a duck, she's made of wood and so should be burned. So we see that just having a test is not quite enough. Tests need to be logical and correct!

Ask yourself if you would you fly on a plane or ride on an elevator if you knew it had never been tested? While lives may not depend so directly on the software you write, you still want to write code that is free from errors.

As programs get longer, the tendency is that they get much harder to maintain. It's like adding a new team member. When there are just two people, there are only two ways to talk. Every new member increases the number of channels exponentially. Each new line of code or function you write interacts with all the other code in ways that become far too complicated for you to remember.

Figure 1.1. Complexity increases quickly!



Tests help us check that our code still does what we think it does. Imagine we want to translate our `hello.py` to Spanish. We need to change the "Hello" to "Hola," but we misspell it "Halo":

```

1      #!/usr/bin/env python3
2      # Purpose: Say hello
3
4      import sys
5
6      def greet(name):
7          return f'Halo, {name}!'                                ①
8
9      def test_greet():
10         assert greet('World') == 'Hola, World!'               ②
11         assert greet('Terra Firma') == 'Hola, Terra Firma!'
12
13     def main():
14         args = sys.argv[1:]
15         name = args[0] if args else 'World'
16         print(greet(name))
17
18     if __name__ == '__main__':
19         main()

```

- ① "Hola" is misspelled as "Halo."
- ② Check for the correct spelling.

If you run `pytest` on this, you will see this failure:

```

=====
 FAILURES =====
 -----
 test_greet

def test_greet():
>     assert greet('World') == 'Hola, World!'

```

```

E      AssertionError: assert 'Halo, World!' == 'Hola, World!'
E          - Halo, World!
E          ?  --
E          + Hola, World!
E          ?  ++
E

hello07.py:10: AssertionError
=====
===== 1 failed in 0.07 seconds =====

```

I would encourage you to write many functions, each of which does a few things as possible. Each function should have a `test_`, either in the same source file (I usually place it immediately after the function it's testing) or in a separate file (which I usually call `unit.py` for "unit" tests).

Often I will even write my `test_` function *before* I write the function itself, and I will imagine how I would want the function to work under a variety of conditions. As a general rule, I usually will test a function with:

- No arguments
- One argument
- Several arguments
- Bad arguments

I will then test if the function works or fails as I expect it. For instance, what would you have `greet` do if given no arguments or given a *list* of arguments instead of one?

1.1.9 Adding type hints

Each variable in Python has a type like "text" or "number." We can execute `python3` on the command line to interact directly with the Python interpreter and see the types. Create a `name` variable set to "World":

```
>>> name = 'World'
```

Now use the `type` command to see how Python represents this data:

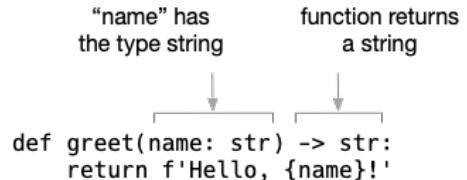
```
>>> type(name)
<class 'str'>
```

Anything in quotes (single or double) is a "string," which in Python is represented by the class `str`. Python has many other types, for instance, the number `10` is an `int` or "integer":

```
>>> type(10)
<class 'int'>
```

We can add "type hints" to our code to ensure that we use the right types of data, like not trying to divide an `int` by a `str` which would cause an *exception* and crash our code.

Shown is how greet looks with type hints. You can read name: str as "name has the type string." Additionally, the greet function itself has been annotated with -> str to indicate that it will return a str value. If a function (like main) returns nothing, add the special None as the return type.



Here is what the entire program looks like with type hints.

Can you find the error?

```

1  #!/usr/bin/env python3
2  # Purpose: Say hello
3
4  import sys
5
6  def greet(name: str) -> str: ①
7      return f'Hello, {name}!'
8
9  def test_greet() -> None: ②
10     assert greet('World') == 'Hello, World!'
11     assert greet('Terra Firma') == 'Hello, Terra Firma!'
12
13 def main() -> None: ③
14     args = sys.argv[1:]
15     print(greet(args))
16
17 if __name__ == '__main__':
18     main()
```

- ① The name argument is a str, and the function returns a str.
- ② The function returns None.
- ③ The function returns None.

Did you find the error? What happens when you run this program? Read on to see how we can find the error.

1.1.10 Verification with mypy

Python itself completely ignores the type hints, but we can use the `mypy` program to find a problem in the code. If you don't have `mypy`, you can install it like so:

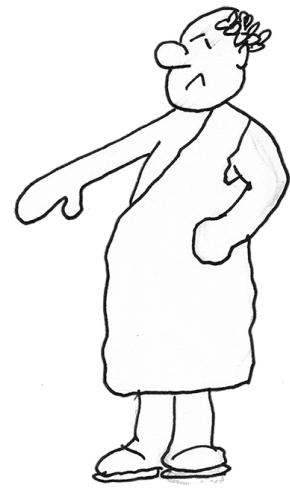
```
$ python3 -m pip install mypy
```

If we run it on `hello.py`, we see that it finds a type error:

```
$ mypy hello.py
hello.py:15: error: Argument 1 to "greet" has incompatible
type "List[str]"; expected "str"
```

Can you see that we tried to pass `args` (which is a list of strings) to `greet` instead of a single `str` value? If you run the above, it will print the entire list `['World']` instead of the value inside the list:

```
$ ./hello.py World
Hello, ['World']!
```



This is not an error that our `test_greet` would have found because it's not a problem with the `greet` function but with *how the function is called*. Here is the correct `main` function:

```
def main() -> None:
    args = sys.argv[1:]
    name = args[0] if args else 'World'
    print(greet(name))
```

1.1.11 Manually validating and documenting arguments

We're now going to change our program to require exactly one argument. If we don't get that, we need to print a "usage" statement that explains how to use the program:

```
$ ./hello.py
usage: hello.py NAME
```

And likewise if run with more than one argument:

```
$ ./hello.py Terra Firma
usage: hello.py NAME
```

When given one argument, it should work as expected:

```
$ ./hello.py "Terra Firma"
Hello, Terra Firma!
```

Here is the new version to make that work:

```
1      #!/usr/bin/env python3
2      # Purpose: Say hello
3
4      import os
5      import sys
```

```

6      def greet(name: str) -> str:
7          return f'Hello, {name}!'
8
9
10     def test_greet() -> None:
11         assert greet('World') == 'Hello, World!'
12         assert greet('Terra Firma') == 'Hello, Terra Firma!'
13
14     def main() -> None:
15         args = sys.argv[1:]
16         if len(args) != 1:
17             prg_name = os.path.basename(sys.argv[0]) ①
18             print(f'usage: {prg_name} NAME') ②
19             sys.exit(1) ③
20         else:
21             print(greet(args[0]))
22
23     if __name__ == '__main__':
24         main()

```

- ① args must have 1 element.
- ② sys.argv[0] has program name. Use basename to remove the path.
- ③ Print "usage" statement.
- ④ Exit with error (not a 0 value).

Here I import os so that I can use the os.path.basename function to get the "basename" of sys.argv[0] which, you'll recall, is the path of the program itself. The "basename" of a path is the filename itself:

```

>>> import os
>>> os.path.basename('/path/to/hello.py')
'hello.py'

```

On line 15, I check if the len (length) of the args is *not* 1. If so, I'll print a "usage" statement on how to run the program. Note that I call sys.exit(1) to indicate a *non-zero exit value* because 0 (which is the default) indicates "zero errors." Any exit value that is not 0 indicates a failure.

Most Unix tools will also respond to -h or --help to show a "usage" statement. What happens when we try that?

```

$ ./hello.py --help
Hello, --help!

```

Well, that didn't work out. Just because we gave --help as an argument didn't mean that our program interpreted it correctly. Our program has been written to process *positional* arguments only, and the --help argument is a "flag," meaning a value that is True when present and False when absent. Here is a pretty ugly way to code this:

```

def main() -> None:
    args = sys.argv[1:]
    val = args[0] if args else ''
    if len(args) != 1 or val == '-h' or val == '--help':
        prg_name = os.path.basename(sys.argv[0])

```

```

    print(f'usage: {prg_name} NAME')
else:
    print(greet(val))

```



Now we are setting a `val` ("value") variable equal to the first argument if one is present. Then we can check `len(args)` or if the `val` is equal to either the string '`-h`' or '`--help`'. I can verify that it works:

```

$ ./hello.py -h
usage: hello.py NAME
$ ./hello.py --help
usage: hello.py NAME

```

That big compound `if` test is pretty ugly. It's what I would call a "code smell" — it's trying to do too many things at once. We can write something much more elegant.

1.1.12 Documenting and validating arguments using argparse

The Python language has a standard module called `argparse` that will parse all the command line arguments, options, and flags. You have to invest a bit of time to learn it, but it will save you so much time in return. Don't worry if you don't understand this. Chapter 2 will show you many examples of how to use `argparse` that you can copy and paste for your own programs. By the end of this book, you will be an `argparse` pro!

Here is a new version that will make our code much cleaner:

```

1  #!/usr/bin/env python3
2  # Purpose: Say hello
3
4  import argparse
5
6  def greet(name: str) -> str:
7      return f'Hello, {name}!'
8
9  def test_greet() -> None:
10     assert greet('World') == 'Hello, World!'
11     assert greet('Terra Firma') == 'Hello, Terra Firma!'
12
13 def main() -> None:
14     parser = argparse.ArgumentParser(description='Say hello') ①
15     parser.add_argument('name', help='Name to greet') ②
16     args = parser.parse_args() ③
17     print(greet(args.name)) ④
18
19 if __name__ == '__main__':
20     main()

```

- ① Create argument parser.
- ② Add the name parameter.
- ③ Get the parsed args.
- ④ Print the greeting given the `args.name` value.



Now instead of directly dealing with `sys.argv`, we describe to `argparse` that we want a single argument called `name` and let it do the hard work of parsing and validating the arguments. This will be a *positional argument* because `name` does *not* start with a dash. We will set `args` equal to the result of our parser doing the work to `parse_args`.

Now if you run the program with no arguments, the "usage" statement will be generated by `argparse`. Though you can't see it directly, the exit value is also set to something other than 0 to indicate failure because we failed to provide the required argument:

```
$ ./hello.py
usage: hello.py [-h] str
hello.py: error: the following arguments are required: str
```

And both the `-h` and `--help` flags will trigger a longer help document that looks like a Unix `man` page.

```
$ ./hello.py --help
usage: hello.py [-h] str

Say hello

positional arguments:
  str          The name to greet

optional arguments:
  -h, --help    show this help message and exit
```

1.1.13 Adding `get_args`

As a matter of personal taste, I like to put all the `argparse` code into its own function that I always call `get_args`. For some of my programs, this can get quite long, and it makes the `main` function stay much shorter if this is separated. Getting the command-line arguments is a functional unit in my mind, and so it belongs by itself. I always put `get_args` as the first function so that I can see it immediately when I read the source code. I usually put `main` right after it. You are, of course, welcome to structure your programs however you like.

Here is how the program looks now:

```
1  #!/usr/bin/env python3
2  # Purpose: Say hello
3
4  import argparse
5
6  def get_args() -> argparse.Namespace: ①
7      parser = argparse.ArgumentParser(description='Say hello')
8      parser.add_argument('name', metavar='str', help='The name to greet')
```

```

9         return parser.parse_args()
10
11     def main() -> None:
12         args = get_args()                      ②
13         print(greet(args.name))
14
15     def greet(name: str) -> str:
16         return f'Hello, {name}!'
17
18     def test_greet() -> None:
19         assert greet('World') == 'Hello, World!'
20         assert greet('Terra Firma') == 'Hello, Terra Firma!'
21
22     if __name__ == '__main__':
23         main()

```

- ① The `get_args` function dedicated to getting the command-line arguments
 ② Call `get_args` function to get parsed arguments.

1.1.14 Checking style and errors



We now have the bones of a pretty respectable program. I like to use the `flake8` and `pylint` tools to give me suggestions on how to improve my programs. This is a process called "linting," and the tools are called "linters." You can use the `pip` module to install them like so:

```
$ python3 -m pip install flake8 pylint
```

I find that `flake8` is unhappy with readability as it tells me to put 2 lines after each function definition:

```
$ flake8 hello.py
hello.py:6:1: E302 expected 2 blank lines, found 1
hello.py:11:1: E302 expected 2 blank lines, found 1
hello.py:15:1: E302 expected 2 blank lines, found 1
hello.py:18:1: E302 expected 2 blank lines, found 1
hello.py:22:1: E305 expected 2 blank lines after class or function
definition, found 1
```

The `pylint` program has other things to complain about, namely that my functions are missing documentation ("docstrings"):

```
$ pylint hello.py
*****
Module hello
hello.py:1:0: C0111: Missing module docstring (missing-docstring)
hello.py:6:0: C0111: Missing function docstring (missing-docstring)
hello.py:11:0: C0111: Missing function docstring (missing-docstring)
hello.py:15:0: C0111: Missing function docstring (missing-docstring)
hello.py:18:0: C0111: Missing function docstring (missing-docstring)

-----
Your code has been rated at 6.67/10 (previous run: 6.67/10, +0.00)
```

A docstring is a string that occurs just after the `def` of the function. It can be a single line of text enclosed in single or double quotes. It's also common to use several lines,

and Python allows you to use triple-quotes for strings that have line breaks. You can assign them to a variable:

```
>>> multi_line = """
... I should have been a pair of ragged claws.
... Scuttling across the floors of silent seas.
... """
```

Or you can use them in place of the `#` for multi-line comments. It's common practice to write a long docstring just after the shebang to document the overall purpose of the function. I like to include my name, email address, the purpose of the script, and the date so that any future person using my program will know who wrote it and how to get in touch with me if they have problems.

To fix the formatting issues, I ran the whole program through the formatting program called yapf (Yet Another Python Formatter). Another popular formatter is `black`. It really doesn't matter which one you choose, just choose one and use it consistently.

Here is a version that will silence all of our critics. Note that I like to add comments followed by long lines of dashes as visual separators between my functions, but you can omit these. I think this program is nicely formatted and more readable than previous versions:

```
#!/usr/bin/env python3
"""

Purpose: Say hello
Author: Ken Youens-Clark
"""

import argparse

# -----
def get_args() -> argparse.Namespace:
    """Get command-line arguments"""
    parser = argparse.ArgumentParser(description='Say hello')
    parser.add_argument('name', metavar='str', help='The name to greet')
    return parser.parse_args()

# -----
def main() -> None:
    """Start here"""

    args = get_args()
    print(greet(args.name))

# -----
def greet(name: str) -> str:
    """Create a greeting"""

    return f'Hello, {name}!'

# -----
```

```

def test_greet() -> None:
    """Test greet"""

    assert greet('World') == 'Hello, World!'
    assert greet('Terra Firma') == 'Hello, Terra Firma!'

# -----
if __name__ == '__main__':
    main()

```

- ① Triple-quoted, multi-line docstring for program/module.
- ② A big horizontal "line" to help me see the functions.
- ③ Triple-quotes can be used on a single line, too.

1.1.15 Making the argument optional

Let's return to making the argument to our program optional so that we can run with and without an argument. We'd like to run the program with no argument and have it default to using "World" for the name like so:

```
$ ./hello.py
Hello, World!
```

If we make name an *optional* argument, it can no longer be a *positional* argument. We create the option --name that will be followed with the *name* to greet:

```
$ ./hello.py --name Cleveland
Hello, Cleveland!
```

The "short" name is a single dash and a single letter like -n, and the "long" name is two dashes followed by a longer name like --name. Here is what the usage looks like now:

```

$ ./hello.py -h
usage: hello.py [-h] [-n str]

Say hello

optional arguments:
  -h, --help            show this help message and exit
  -n str, --name str   The name to greet (default: World)

```

And here is the code. Note the use of the `formatter_class` to have `argparse` show the default values for arguments in the "usage" output:

```

#!/usr/bin/env python3
"""

Purpose: Say hello
Author: Ken Youens-Clark
"""

import argparse

# -----
def get_args() -> argparse.Namespace:
    """Get command-line arguments"""

```

```

parser = argparse.ArgumentParser(
    description='Say hello',
    formatter_class=argparse.ArgumentDefaultsHelpFormatter) ①

parser.add_argument('-n',
                    '--name',
                    default='World',
                    metavar='str',
                    help='The name to greet') ②③④⑤

return parser.parse_args()

# -----
def main() -> None:
    """Start here"""

    args = get_args()
    print(greet(args.name))

# -----
def greet(name: str) -> str:
    """Create a greeting"""

    return f'Hello, {name}!'

# -----
def test_greet() -> None:
    """Test greet"""

    assert greet('World') == 'Hello, World!'
    assert greet('Terra Firma') == 'Hello, Terra Firma!'

# -----
if __name__ == '__main__':
    main()

```

- ① Show default values in "usage."
- ② Short option name
- ③ Long option name
- ④ The default value.
- ⑤ Hint to user of the data type.

1.1.16 Testing `hello.py`

Included in the `hello` directory is a `test.py` that will run the program to ensure it creates the correct output. We can run it either using `make test` or `pytest -xv test.py`² I will elide some of the output:

² The `make` program looks for a file called `Makefile` in the directory where we run `make`. Inside that file is a "target" called `test` that says the command that should be run is `pytest -xv test.py`. `Makefiles` are a way to create workflows of actions, each of which can depend on each other. I'm using `make` here to create a shortcut for running the test, but it's not necessary. You can run `pytest` directly.

```
$ make test
pytest -xv test.py
=====
...                                          test session starts =====
collected 4 items

test.py::test_exists PASSED [ 25%] ①
test.py::test_usage PASSED [ 50%] ②
test.py::test_default PASSED [ 75%] ③
test.py::test_input PASSED [100%] ④

=====
4 passed in 0.40s =====
```

- ① The first test in every `test.py` checks to see if the expected program file exists. Here it is checking if `hello.py` exists in the same directory as the `test.py` program.
- ② The second test always runs the program with `-h` and `--help` to see if it produces anything that looks like a "usage."
- ③ This test runs the program with no arguments to see if it produces the default output, Hello, World!
- ④ This test runs the program with the `-n` and `--name` options and different values to see if it accepts a name option and creates the correct output.

In the parlance of testing, the `test_greet` function that we wrote to test the `greet` function might be called a "unit test" as it tests one "unit" (a function) of code, while the `test.py` might be called an "integration test" because it's looking at the program as a whole from the outside to see if it performs correctly.

It may be a bit overwhelming to look at the `test.py`, but it will be instructive as the goal will be for you to learn how to write your own tests for your programs:

```
#!/usr/bin/env python3
"""tests for hello.py"""

import os
from subprocess import getstatusoutput, getoutput

prg = './hello.py' ①

# -----
def test_exists(): ②
    """exists"""

    assert os.path.isfile(prg)

# -----
def test_usage(): ③
    """usage"""

    for flag in ['-h', '--help']:
        rv, out = getstatusoutput(f'{prg} {flag}')
        assert rv == 0
        assert out.lower().startswith('usage')

# -----
```

```

def test_default(): ④
    """Says 'Hello, World!' by default"""

    out = getoutput(prg)
    assert out.strip() == 'Hello, World!'

# -----
def test_input(): ⑤
    """test for input"""

    for val in ['Universe', 'Multiverse']:
        for option in ['-n', '--name']:
            rv, out = getstatusoutput(f'{prg} {option} {val}')
            assert rv == 0
            assert out.strip() == f'Hello, {val}!'

```

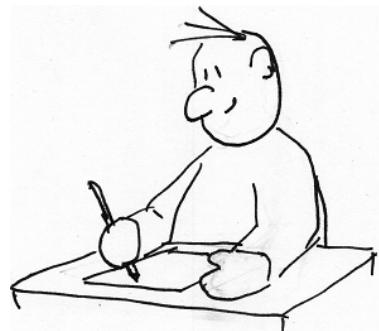
- ① This is a name of the expected program.
- ② Tests are run in the order that they are defined in the program. This will be the first test to check if the prg exists.
- ③ Run the program with the -h and --help flags to see if the output looks like a "usage" statement.
- ④ This test will run the program with no arguments to see if it prints Hello, World!
- ⑤ Run the program using both -n and --name with two different values to see if the program accepts the name option and prints the correct value.

In addition to running the integration `test.py` programs I have provided, I will suggest you include and run unit tests in your programs. I will also suggest writing your own functions and unit tests as well as extending the provided `test.py` to cover functionality that you add to your programs. I encourage you to study the `test.py` programs as you will learn as much or more from them as from the exercises you write!

1.1.17 Starting a new program with `new.py`

In my own practice, I almost never start writing a Python program from an empty page. I created a Python program called `new.py` that helps me start writing new Python programs. As most of my programs need to take parameters, I always use the `argparse` to interpret the command-line options. I have put my `new.py` program into the `bin` ("binaries" even though these are just text files) directory of the GitHub repo, and I recommend you start every new program with this program.

A central tenet of this book is to create *documented* and *testable* programs. Anything that *can* change about a program should be passed as an *argument* to the program when it is run. For instance, if a program will read an input file for data, the name of that file should be an argument like `-f input.txt`. Then the fact that the program takes an input file is now visible through the `--help` and we can pass in different files and test if the program processes the files correctly. It is *not* a requirement that you use `new.py` and `argparse`, however. As long as your programs process command-line arguments in the same way as `argparse` and always produce a usage on `-h` or `--help`, your programs should pass the test suites. The template that `new.py` provides is meant only to make it faster and more convenient to create new programs.



Here is how I would use `new.py` to create a new `hello.py` program:

```
$ new.py hello.py
Done, see new script "hello.py."
```

This is what will be produced:

```
#!/usr/bin/env python3
"""
Author : Ken Youens-Clark <kyclark@gmail.com>
Date   : 2019-10-21
Purpose: Rock the Casbah
"""

import argparse
import os
import sys

# -----
def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Rock the Casbah',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('positional',
                        metavar='str',
                        help='A positional argument')

    parser.add_argument('-a',
                        '--arg',
                        help='A named string argument',
                        metavar='str',
```

```

        type=str,
        default='')

parser.add_argument('-i',
                    '--int',
                    help='A named integer argument',
                    metavar='int',
                    type=int,
                    default=0)

parser.add_argument('-f',
                    '--file',
                    help='A readable file',
                    metavar='FILE',
                    type=argparse.FileType('r'),
                    default=None)

parser.add_argument('-o',
                    '--on',
                    help='A boolean flag',
                    action='store_true')

return parser.parse_args()

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()
    str_arg = args.arg
    int_arg = args.int
    file_arg = args.file
    flag_arg = args.on
    pos_arg = args.positional

    print('str_arg = "{}".format(str_arg)')
    print('int_arg = "{}".format(int_arg)')
    print('file_arg = "{}".format(file_arg.name)')
    print('flag_arg = "{}".format(flag_arg)')
    print('positional = "{}".format(pos_arg)')

# -----
if __name__ == '__main__':
    main()

```

The arguments that this program will accept are:

1. A single positional argument of the type `str`. *Positional* means it is not preceded by a flag to name it but has meaning because of its position.
2. An automatic `-h` or `--help` flag that will cause `argparse` to print the usage.
3. A named string argument called either `-a` or `--arg`
4. A named interger argument called `-i` or `--int`
5. A named file argument called `-f` or `--file`
6. A boolean (off/on) flag called `-o` or `--on`

Each option here has both a "short" and "long" name. It is not a requirement to have both, but it is common and tends to make your program more readable.

After you use `new.py` to start your new program, you should open it with your editor and modify the argument names and types to suit the needs of your program. For instance, in the "Crow's Nest" chapter, you can delete everything but the positional argument which you should rename from `'positional'` to something like `'word'` (because the argument is going to be a word).

Note that you can control the name and `email` values that are used by `new.py` by creating a file called `.new.py` (note the leading dot!) in your home directory. Here is mine:

```
$ cat ~/.new.py
name=Ken Youens-Clark
email=kyclark@gmail.com
```

If you don't want to use `new.py`, then I have included a sample of the above program as `template/template.py` that you can copy. For instance, in the "Crow's Nest" chapter you should create the program `crowsnest/crowsnest.py`. (That is, from the root directory of the repository, go into the `crowsnest` directory and create a file called `crowsnest.py`.)

Either you can do this with `new.py`:

```
$ cd crowsnest
$ new.py crowsnest.py
```

Or the `cp` (copy) command:

```
$ cp template/template.py crowsnest/crowsnest.py
```

The main point is that I don't want you to have to start every program from scratch! I think it's much easier to start from a complete, working program and modify it.

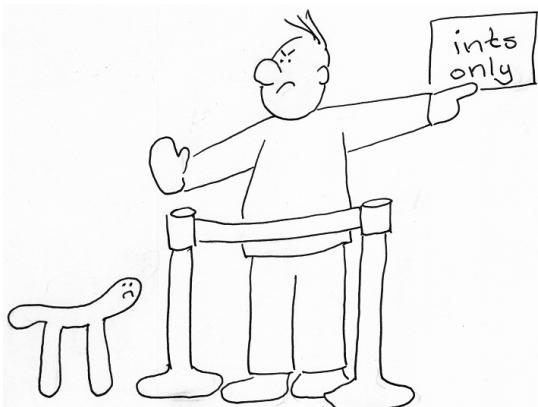
1.2 Summary

- The `argparse` module will help you document and parse all the parameters to your program. You can validate the types and numbers of arguments which can be positional, optional, or flags. The usage will be automatically generated.
- You should write short functions with accompanying `test_` functions which can be run by `pytest`.
- It is often helpful to write the tests before you write the functions to imagine how they ought to work.
- You should run your tests after any change to your program to ensure that everything still works.
- Code formatters like `yapf` and `black` will automatically format your code to community standards, making it easier to read and debug.
- The `mypy` tool can check for errors in your code that has been annotated with type hints.
- Code linters like `pylint` and `flake8` can help you correct both programmatic and stylistic problems.

- You can use the `new.py` program to generate new Python programs that use `argparse`.

Using argparse

Often getting the right data into your program is a real chore. The `argparse` module can really make your life much easier by validating and rejecting bad arguments from the user. It's like our program's "bouncer," only allowing the right kinds of values into our program. Often half or more of the programs in this book can be handled simply by defining the arguments properly with `argparse`!



In Chapter 1, we ended up writing a very flexible program that could extend warm salutations to an optionally named entity such as the "World" or "Universe":

```
$ ./hello.py
Hello, World!
$ ./hello.py --name Universe
Hello, Universe!
```

The program would respond to the `-h` and `--help` flags with helpful documentation:

```
$ ./hello.py -h
usage: hello.py [-h] [-n str]

Say hello

optional arguments:
```

```
-h, --help      show this help message and exit
-n str, --name str  The name to greet (default: World)
```

The `argparse` module helped us define a parser for the parameters and generate the usage, saving us loads of time and making our program look professional. Every program in this book is tested on different inputs, so you'll really understand how to use this module by the end. I would recommend you look over the documentation (docs.python.org/3/library/argparse.html). Now let's dig further into what this module can do for us. In this chapter, we will:

- Learn how to use `argparse` to handle positional parameters, options, and flags.
- Set default values for options.
- Use `type` to force the user to provide values like numbers or files.
- Use `choices` to restrict the values for an option.

2.1 Types of arguments

As we saw in Chapter 1, command-line arguments can be classified as follows:

- **Positional arguments:** The order and number of the arguments is what determines their meaning. Some programs might expect, for instance, a file name as the first argument and an output directory as the second. Typically positional arguments are always required. Making them optional is difficult — how would you write a program that accepts 2 or 3 arguments where the second and third ones are independent and optional? In all the versions of `hello.py` up until the last one, the argument (a name to greet) was positional.
- **Named options:** Standard Unix format allows for a "short" name like `-f` (one dash and a single character) or a "long" name like `--file` (two dashes and a string of characters) followed by some value like a file name or a number. Named options allow for arguments to be provided in any order, so their *position* is not relevant; hence they are the right choice when the user is not required to provide them (they are "options," after all). It's good to provide reasonable default values for options. We changed the required, position name argument of `hello.py` to the optional `--name`. Note that some languages like Java might define "long" names with a single dash like `-jar`.
- **Flags:** A "Boolean" value like "yes"/"no" or `True/False` is indicated by something that starts off looking like a named option but there is no value after the name, for example, `-d` or `--debug` to turn on debugging. Typically the presence of the flag indicates a `True` value for the argument; therefore, its absence would mean `False`, so `--debug` turns *on* debugging while its absence means it is off. If you run `ls -s`, the `-s` is the flag to show that you want the files sorted by size.

2.2 Using argparse

Let's dig a bit deeper into the `get_args` function which is defined like this:

```
def get_args():
```

```
"""Get command-line arguments"""


```

The `def` keyword defines a new function. The arguments to the function are listed in the parentheses. Even though the `get_args` function takes no arguments, the parentheses are still required. The triple-quoted line after the function `def` is the "docstring" which serves as a bit of documentation for the function. Docstrings are not required, but they are good style and `pylint` will complain if you leave them out.

2.2.1 Creating the parser

The following line creates a parser that will deal with the arguments from the command line. To "parse" here means to infer some meaning from the order and syntax of the bits of text provided as arguments:

```
parser = argparse.ArgumentParser(
    description='Argparse Python script',
    formatter_class=argparse.ArgumentDefaultsHelpFormatter) ①
②
③
```

- ① Call the `argparse.ArgumentParser` method to create a new parser.
- ② A short summary of your program's purpose.
- ③ The `formatter_class` argument tells `argparse` to show the default values in usage. See the documentation for other values you can use.

2.2.2 A positional parameter

The following line will create a new *positional* parameter:

```
parser.add_argument('positional',
                   metavar='str',
                   help='A positional argument') ①
②
③
```

- ① The lack of leading dashes makes this a positional parameter, not the name "positional."
- ② A hint to the user for the data type. By default, all arguments are strings.
- ③ A brief description of the parameter for the usage.

Remember that the parameter is not positional because the *name* is "positional." That's just there to remind you that it *is* a positional parameter. The `argparse` interprets the string 'positional' as such because it is not preceded with any dashes.

2.2.3 An optional string parameter

The following line creates an *optional* parameter with a short name of `-a` and a long name of `--arg` that will be a `str` with a default value of '' (the empty string). Note that you can leave off either the short or long name in your own programs, but it's good form to provide both. Most of the tests for the exercises will use both short and long option names.

```
parser.add_argument('-a',
                   '--arg',
                   help='A named string argument',
                   metavar='str') ①
②
③
④
```

```
    type=str,
    default='')
```

(5)
(6)

- (1) The short name.
- (2) The long name.
- (3) Brief description for the usage.
- (4) Type hint for usage.
- (5) The actual Python data type (note the lack of quotes around str).
- (6) The default value.

If you wanted to make this a required, named parameter, you would remove the default and add required=True.

2.2.4 An optional numeric parameter

The following line creates the option called -i or --int that accepts an int (integer) with a default value of 0. If the user provides anything that cannot be interpreted as an integer, the argparse module will stop processing the arguments and will print an error message and a short usage statement:

```
parser.add_argument('-i',
                    '--int',
                    help='A named integer argument',
                    metavar='int',
                    type=int,
                    default=0)
```

(1)
(2)
(3)
(4)
(5)
(6)

- (1) The short name.
- (2) The long name.
- (3) Brief description for usage.
- (4) Type hint for usage.
- (5) Python data type that the string must be converted to. You can also use float for a floating point value (a number with a fractional component like 3.14).
- (6) The default value.

One of the big reasons to define numeric arguments in this way is that argparse will convert the input to the correct type. That is, all values coming from the command are strings. It's the job of the program to convert the value to an actual numeric value. If you tell argparse that the option should be type=int, then when you ask the parser for the value, it will have already been converted to an actual int value. If the value provided by the user cannot be converted to an int, then the value will be rejected. That saves you a lot of time and effort!

2.2.5 An optional file parameter

The following line creates an option called -f or --file that will only accept a valid, readable file. This argument alone is worth the price of admission as it will save you oodles of time validating the input from your user. Note that pretty much every exercise that has a file input will have tests that pass *invalid* file arguments to ensure

that your program rejects them.

```
parser.add_argument('-f',  
                  '--file',  
                  help='A readable file',  
                  metavar='FILE',  
                  type=argparse.FileType('r'),  
                  default=None)
```

- ① The short name.
- ② The long name.
- ③ Brief usage.
- ④ Type suggestion.
- ⑤ Says that the argument must name a readable ('r') file.
- ⑥ Default value.

2.2.6 A flag

The flag option is slightly different in that it does not take a value like an string or integer. It's just the name part itself. Flags are either present or not. A common flag is `--debug` to turn *on* debugging statements or `--verbose` to print extra messages to the user. When `--debug` is not present, the default value is "off" (or `False`), which is why the action for this argument is `store_true`. It's not necessary to set a default value as it is automatically set to `False`.

```
parser.add_argument('-o',  
                  '--on',  
                  help='A boolean flag',  
                  action='store_true')
```

- ① Short name.
- ② Long name.
- ③ Brief usage.
- ④ What to do when this is present. The default value is `False`, so `True` will be stored if present.

2.2.7 Returning from get_args

The final statement in `get_args` is to return the result of having the `parser` object parse the arguments. That is, the code that calls `get_args` will receive this value back:

```
return parser.parse_args()
```

This could fail because `argparse` finds that the user provided invalid arguments, for example, a string value when it expected a `float` or perhaps a misspelled filename. If the parsing succeeds, then we will have a way in our code to access all the values the user provided. Additionally, those values will be of the *types* that we indicated. That is, if we indicate that the `--int` argument should be an `int`, then when we ask for `args.int`, it will already be an `int`. If we define a file argument, we'll get an *open file handle*. That may not seem impressive now, but it's really enormously helpful.

2.2.8 Manually checking arguments

It's also possible to manually validate arguments before you return from `get_args`. For instance, we can define that `--int` should be an `int` but how can we require that it must be between 1 and 10? One fairly simple way to do this is to manually check the value. If there is a problem, you can use the `parser.error` function to halt execution of the program, print an error message along with the short usage, and then exit with an error:

```
#!/usr/bin/env python3
"""Manually check an argument"""

import argparse

# -----
def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Manually check an argument',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('-v',
                        '--val',
                        help='Integer value between 1 and 10',
                        metavar='int',
                        type=int,
                        default=5)

    args = parser.parse_args()
    if not 1 <= args.val <= 10: ①
        parser.error(f"--val '{args.val}' must be between 1 and 10") ② ③

    return args ④

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()
    print(f'val = "{args.val}"')

# -----
if __name__ == '__main__':
    main()
```

- ① Parse the arguments.
- ② Check if the `args.int` value is *not* between 1 and 10.
- ③ Call `parser.error` with an error message. The entire program will stop, the error message and the brief usage will be shown to the user.
- ④ If we get here, then everything was OK, and the program will continue as normal.

If we provide a good `--val`, all is well:

```
$ ./manual.py -v 7
val = "7"
```

If we run this program with a value like 20, we get an error message:

```
$ ./manual.py -v 20
usage: manual.py [-h] [-v int]
manual.py: error: --val "20" must be between 1 and 10
```

It's not possible to tell here, but the `parser.error` also caused the program to exit with a non-zero status. In the Unix world, an exit status of 0 indicates "zero errors," so anything not 0 is considered an error. You may not realize just yet how wonderful that is, so just trust me. It is.

2.3 Examples using argparse

Many of the program tests can be satisfied by learning how to use `argparse` effectively to validate the arguments to your programs. I think of the command line as the boundary of your program, and you need to be judicious about what you let in. You should always expect and defend against every argument being wrong.³ Our `hello.py` program was an example of a single, positional argument and then a single, optional argument. Let's look at some more examples of how you can use `argparse`.

2.3.1 Two different positional arguments

Imagine you want two *different* positional arguments, like the `color` and `size` of an item to order. The color should be a `str`, and the size should be an `int` value. When you define them positionally, the order in which you declare them is the order in which the user must supply the arguments.

Here we define `color` first and then `size`:

```
#!/usr/bin/env python3
"""
Two positional arguments
"""

import argparse

# -----
def get_args():
    """
    get args
    """

    parser = argparse.ArgumentParser(
        description='Two positional arguments',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('color', ①
                        metavar='str',
                        type=str,
                        help='Color')

    parser.add_argument('size', ②
```

³ I always think of the kid who will type "fart" for every input.

```

        metavar='int',
        type=int,
        help='Size')

    return parser.parse_args()

# -----
def main():
    """main"""

    args = get_args()
    print('color =', args.color) ③
    print('size =', args.size)

# -----
if __name__ == '__main__':
    main()

```

- ① This will be the first of the positional arguments because it is defined first.
- ② This will be the second of the position arguments.
- ③ The argument is accessed via the name of the parameter color.

Again, the user must provide exactly two positional arguments. No arguments triggers a short usage:

```
$ ./two_args.py
usage: two_args.py [-h] str int
two_args.py: error: the following arguments are required: str,
int
```

Just one won't cut it. We are told that "size" is missing:

```
$ ./two_args.py blue
usage: two_args.py [-h] str int
two_args.py: error: the following arguments are required: int
```



If we give two arguments, the second of which can be interpreted as an int, all is well:

```
$ ./two_args.py blue 4
color = blue
size = 4
```



Remember that *all* the arguments coming from the command line are strings. The shell (here bash) doesn't require quotes around the blue or the 4. To the shell, these are both strings, and they are passed to Python as strings. When we tell argparse that the second argument needs to be an int, then argparse will do the work to attempt the conversion of the string '4' to the integer 4. If you provide 4.1, that will be rejected,

too:

```
$ ./two_args.py blue 4.1
usage: two_args.py [-h] str int
two_args.py: error: argument int: invalid int value: '4.1'
```

Positional arguments have the problem that the user is required to remember the correct order. In the case of switching a `str` and `int`, `argparse` will detect invalid values:

```
$ ./two_args.py 4 blue
usage: two_args.py [-h] COLOR SIZE
two_args.py: error: argument SIZE: invalid int value: 'blue'
```

Imagine, however, a case of two strings or two numbers which represent two *different* values like a car's make and model or a person's height and weight. How could you detect that the arguments are reversed? Generally speaking, I only ever create programs that take exactly one positional argument or one or more of *the same thing* like a list of files to process.

2.3.2 Two of the same positional arguments

If you were writing a program that adds two numbers, you could define them as two positional arguments, like `number1` and `number2`. Since they are the same kinds of arguments (two numbers that we will add), it might make more sense to use the `nargs` option to tell `argparse` that you want exactly two of some thing:

```
#!/usr/bin/env python3
"""nargs=2"""

import argparse

# -----
def get_args():
    """get args"""

    parser = argparse.ArgumentParser(
        description='nargs=2',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('numbers',
                        metavar='int',
                        nargs=2,          ①
                        type=int,         ②
                        help='Numbers')

    return parser.parse_args()

# -----
def main():
    """main"""

    args = get_args()
    n1, n2 = args.numbers           ③
    print(f'{n1} + {n2} = {n1 + n2}') ④
```

```
# -----
if __name__ == '__main__':
    main()
```

- ① The nargs=2 will require exactly 2 values.
- ② Each value must be parseable as an integer value or the program will error out.
- ③ Since we defined that there are exactly two values for numbers, we can copy them into two variables.
- ④ Because these are actual int values, the result of + will be numeric addition and not string concatenation.

The help indicates we want two numbers:

```
$ ./nargs2.py
usage: nargs2.py [-h] int int
nargs2.py: error: the following arguments are required: int
```

On line 29, you see we can unpack the two numbers into n1 and n2 and use them in the next line for addition:

```
$ ./nargs2.py 3 5
3 + 5 = 8
```



It's completely safe to unpack numbers in this way because we would never get to line 29 if the user hadn't provided exactly two arguments, both of which can be converted to int values. Also, notice that the n1 and n2 values were actually integers. If they had been strings, then our program would print 35 instead of 8 for the arguments 3 and 5 because the + operator in Python both adds numbers and concatenates strings!

```
>>> 3 + 5
8
>>> '3' + '5'
'35'
```

2.3.3 One or more of the same positional arguments

You could expand your 2-number adder into one that sums as many numbers as you provide. When you want *one or more* of some argument, you can use nargs='+' :

```
#!/usr/bin/env python3
"""nargs=+"""

import argparse

# -----
def get_args():
    """get args"""
```

```

parser = argparse.ArgumentParser(
    description='nargs=+',
    formatter_class=argparse.ArgumentDefaultsHelpFormatter)

parser.add_argument('numbers',
                    metavar='INT',
                    nargs='+',
                    type=int,
                    help='Numbers')

return parser.parse_args()

# -----
def main():
    """main"""

    args = get_args()
    numbers = args.numbers

    print('{} = {}'.format(' + '.join(map(str, numbers)), sum(numbers))) ④

# -----
if __name__ == '__main__':
    main()

```

- ① The + will make nargs accept one or more values.
- ② The int means that all the values must be integer values.
- ③ numbers will be a list with at least one element.
- ④ Don't worry if you don't understand this line. You will by the end of the book!

Note that this will mean `args.numbers` is always a list. Even if the user provides just one argument, `args.numbers` will be a list containing that one value:

```
$ ./nargs+.py 5
5 = 5
$ ./nargs+.py 1 2 3 4
1 + 2 + 3 + 4 = 10
```

2.3.4 Restricting values using choices

Sometimes you want to limit the values of an argument. Maybe you offer shirts in only primary colors. You can pass in a list of valid values using the `choices` option. Here we restrict the color to one of "red," "yellow," or "blue."

```

#!/usr/bin/env python3
"""Choices"""

import argparse

# -----
def get_args():
    """get args"""

    parser = argparse.ArgumentParser()

```

```

        description='Choices',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

parser.add_argument('color',
                    metavar='str',
                    help='Color',
                    choices=['red', 'yellow', 'blue']) ①

return parser.parse_args()

# -----
def main():
    """main"""

    args = get_args()
    print('color =', args.color) ②

# -----
if __name__ == '__main__':
    main()

```

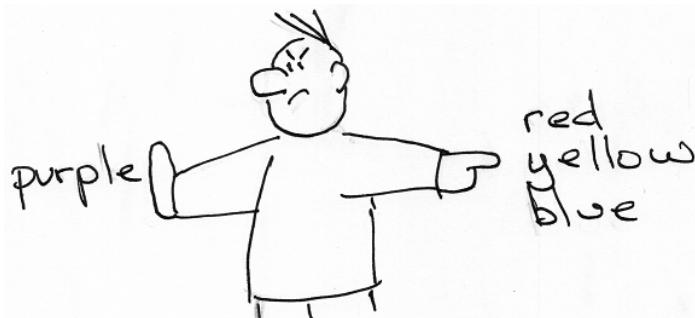
- ① The choices option takes a list of values. argparse will error out if the user fails to supply one of these.
- ② If we make it to this point, we know that args.color will definitely be one of those values. If the value was rejected, the program will never get to this point.

Any value not present in the list will be rejected and the user will be shown the valid choices. Again, no value is rejected:

```
$ ./choices.py
usage: choices.py [-h] str
choices.py: error: the following arguments are required: str
```

If we provide "purple," it will be rejected because it is not in choices we defined. The error message that argparse produces tells the user the problem ("invalid choice") and even lists the acceptable colors!

```
$ ./choices.py purple
usage: choices.py [-h] str
choices.py: error: argument str: invalid choice: 'purple' (choose from 'red', 'yellow',
'blue')
```



That's really quite a bit of error checking and feedback that you never have to write. The best code is code you don't write!

2.3.5 File arguments

So far we've seen that we can define that an argument should be of a type like `str` (which is the default), `int`, or `float`. There are many exercises that require a file as input, and you can use the type of `argparse.FileType('r')` to indicate that an argument must be a *file* which is *readable* (the '`r`' part).

Here is an example showing an implementation in Python of the command `cat -n` where `cat` will *concatenate* files and the `-n` says to *number* the lines of output (which is the command I use to create the following numbered line view of the program — how meta):

```
#!/usr/bin/env python3
"""Python version of `cat -n`"""

import argparse

# -----
def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Python version of `cat -n`',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('file',
                        metavar='FILE',
                        type=argparse.FileType('r'), ①
                        help='Input file')

    return parser.parse_args()

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()

    for i, line in enumerate(args.file, start=1): ②
        print(f'{i:6} {line}', end='')

# -----
if __name__ == '__main__':
    main()
```

① The argument will be rejected if it does not name a valid, readable file.

② The value of `args.file` is an open file handle that we can directly read. Again, don't worry if you don't understand this code. We'll talk all about file handles in the exercises!

When I define an argument as `type=int`, I get back an actual `int` value. Here, I define the `file` argument as a file type, and so I receive an *open file handle*. If I had defined the `file` argument as a string, I would have to manually check if it were a file and then use `open` to get a file handle:

```
file = args.file
if not os.path.isfile(file):
    print(f'{file} is not a file')
    sys.exit(1)

fh = open(file)
```

- ① Get whatever the user passed in for the `file`.
- ② Check if this is *not* a file.
- ③ Print an error message.
- ④ Exit the program with a non-zero value.
- ⑤ Proceed to open the `file`.



With the file type definition, you don't have to write any of this code.

2.3.6 Automatic help

When you define a program's parameters using `argparse`, the `-h` and `--help` flags will be reserved for generating help documentation. You do not need to add these nor are you allowed to use these flags for other purposes.

I think of this documentation like a door to your program. Doors are how we get into buildings and cars and such. Have you ever come across a door that you can't figure out how to open? Or one that requires a "PUSH" sign when clearly the handle is design to "pull"? The book *The Design of Everyday Things* by Don Norman uses the term "affordances" to describe the interfaces that objects present to us which do or do not inherently describe how we should use them.



The usage statement of your program is like the handle of the door. It should let me know exactly how to use it. When I encounter a program I've never used, I either run it with no arguments or with `-h` or `--help`. I *expect* to see some sort of usage statement. The only alternative would be to open the source code itself and study how to make the program run and how I can alter it, and this is a truly unacceptable way to write and

distribute software!

When you start a new program with `new.py foo`, this is the help that will be generated:

```
$ ./foo.py -h
usage: foo.py [-h] [-a str] [-i int] [-f FILE] [-o] str

Argparse Python script

positional arguments:
  str                  A positional argument

optional arguments:
  -h, --help            show this help message and exit
  -a str, --arg str     A named string argument (default: )
  -i int, --int int     A named integer argument (default: 0)
  -f FILE, --file FILE A readable file (default: None)
  -o, --on              A boolean flag (default: False)
```

Without writing a single line of code, you have

1. an executable Python program
2. that accepts command line arguments
3. and generates a standard and useful help message

This is the "handle" to your program.

2.4 Summary

- Positional parameters typically are always required. If you have more than two or more positional parameters representing different ideas, it would be better to make them into named options.
- Optional parameters can be named like `--file fox.txt` where `fox.txt` is the value for the `--file` option. It is recommended to always define a default value for options.
- `argparse` can enforce many types for arguments including numbers like `int` and `float` or even files.
- Flags like `--help` do not have an associated value. They are considered `True` if present and `False` if not.
- The `-h` and `--help` flags are reserved for use by `argparse`. If you use `argparse`, then your program will automatically respond to these flags with a usage statement.

The Crow's Nest: Working with strings



Avast, you corny-faced gollumpus! Ye are barrelman for this watch. D'ye ken what I mean, ye addle pated blunderbuss?! Ah, land lubber ye be! OK, then, you are the lookout in the crow's nest — the little bucket attached to the top of a mast of a sailing ship. Your job is to keep a lookout for interesting or dangerous things, like a ship to plunder or an iceberg to avoid. When you see something like a "narwhal," you are supposed to cry out, "Ahoy, Captain, **a narwhal** off the larboard bow!" If you see an octopus, you'll shout "Ahoy, Captain, **an octopus** off the larboard bow!" (We'll assume everything is "off the larboard bow" for this exercise. It's a great place for things to be.)



From this point on, I will describe a coding challenge that you should write on your own. I will discuss key ideas you'll need to solve the problems as well as how to use the provided tests to help you know when your program is correct. You should have a copy of the Git repository locally (see the setup instructions). You should write your program in the chapter's directory, like this program should be written in the `crowsnest` directory where the tests for the program live.

In this chapter, we're going to start off working with strings. By the end, you will be able to:

- Create a program that accepts a positional argument and produces usage documentation
- Create a new output string depending on the inputs to the program
- Run a test suite

Your program should be called `crowsnest.py`. It will accept a single positional argument and will print the given argument inside the "Ahoy" bit along with the word "a" or "an" depending on whether the argument starts with a consonant or a vowel.

That is, if given "narwhal," it should do this:

```
$ ./crowsnest.py narwhal
Ahoy, Captain, a narwhal off the larboard bow!
```

And if given "octopus":

```
$ ./crowsnest.py octopus
Ahoy, Captain, an octopus off the larboard bow!
```

This means we're going to need to write a program that accepts some input on the command line, decides on the proper article ("a" or "an") for the input, and prints out a new string that puts those two values into the "Ahoy" phrase.

3.1 Getting started

You're probably ready to start writing the program! Well, hold on just a minute longer, ye duke of limbs. We need to discuss how we'll use the tests to know when our program is working and how we might get started programming.

3.1.1 How to use the tests

"The greatest teacher, failure is." — Yoda

In the code repository, I've included tests that will guide you in the writing of your program. Before you even write the first line of code, I'd like you to run `make test` or `pytest -xv test.py` so you can see how the first test fails. Be sure you are in the `crowsnest` directory for this!

Among all the output, you'll notice this line:

```
test.py::test_exists FAILED [ 16%]
```

If you read more, you'll see lots of other output all trying to convince you that the expected file, `crowsnest.py` does not exist. Learning to read the test output is a skill in itself! It takes quite a bit of practice to learn to read test output, so try not to feel overwhelmed. In my terminal (iTerm on a Mac), the output from `pytest` shows colors and bold print to highlight key failures. The text in bold, red letters is usually where I start, but your terminal may behave differently.

3.1.2 Creating programs with new.py

In order to pass this test, we need to create a file called `crowsnest.py` inside

the `crowsnest` directory where `test.py` is located. While it's perfectly fine to start writing from scratch, I suggest you use the `new.py` program to print some useful boilerplate code that you'll need in every exercise:

```
$ new.py crowsnest.py
Done, see new script "crowsnest.py."
```

If you don't want to use `new.py`, you could copy the `template/template.py` program:

```
$ cp template/template.py crowsnest/crowsnest.py
```

At this point you should have the outline of a working `crowsnest.py` program that accepts command-line arguments. If you run your new `crowsnest.py` with no arguments, it will print a short usage statement like the following (notice how "usage" is the first word of the output):

```
$ ./crowsnest.py
usage: crowsnest.py [-h] [-a str] [-i int] [-f FILE] [-o] str
crowsnest.py: error: the following arguments are required: str
```

Those are not the correct parameters for our program, just the default examples given to you by `new.py`. You will need to modify them to suit this program.

If you run your tests again, you will pass the first *two* tests that check:

1. Does the program exist?
2. Does the program print something that looks like "usage."

And then you will fail the third test. There are more tests after this, but that's all you see if you run `pytest -xv test` because the `-x` flag tells `pytest` to stop at the first failing test. Note that we can combine the `-x` and `-v` flags into `-xv` and that the order doesn't matter, so `-vx` is fine, too:

```
test.py::test_exists PASSED
test.py::test_usage PASSED
test.py::test_consonant FAILED
```

[16%]	(1)
[33%]	(2)
[50%]	(3)

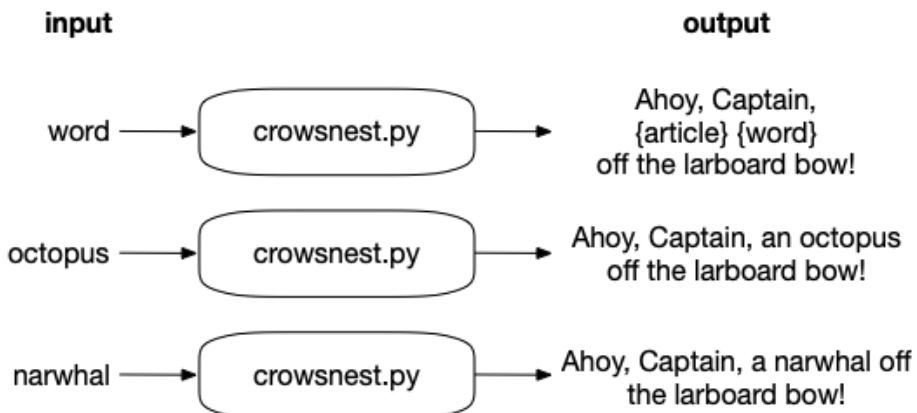
- (1) This test checks if the file `crowsnest.py` exists.
- (2) This test runs `crowsnest.py` to see if it produces a "usage" statement.
- (3) This test runs the program and passes it a word starting with a consonant to see if it produces the expected output. It has FAILED, and so all testing stops because of the `-x` flag.

Now we have a working program that accepts some arguments (but not the right ones). Next we need to make our program accept the "narwhal" or "octopus" value that needs to be announced, and we'll use command-line arguments to do that.

3.1.3 Defining your arguments

Here is a diagram sure to shiver your timbers showing the inputs (or *parameters*) and output of the program. We'll use these throughout the book to imagine how code and data work together. In this program, some "word" is the input, and a phrase incorporating that word with the correct article is the output.

Figure 3.1. The input to the program is a word, and the output is that word plus its proper article (and some other stuff).



We need to modify the part of the program that gets the arguments — the aptly named `get_args` function. This function uses the `argparse` module to parse the command-line arguments. Refer to the `argparse` chapter, particularly the section "A single, positional argument." The default `get_args` names the first argument '`positional`', and that's the only one you need. Remember that positional arguments are defined by their position and don't have names that start with dashes. You can delete all the arguments except for the positional `word`.

Modify the `get_args` part of your program until it will print this usage:

```
$ ./crowsnest.py
usage: crowsnest.py [-h] str
crowsnest.py: error: the following arguments are required: str
```

Likewise, it should print a longer usage for the `-h` or `--help` flag:

```
$ ./crowsnest.py -h
usage: crowsnest.py [-h] str

Crow's Nest -- choose the correct article

positional arguments:
  str      A word ①

optional arguments:
  -h, --help show this help message and exit ②
```

① You need to define a `word` parameter. Notice that it is listed as a "positional" argument.

② The `-h` and `--help` flags are created automatically by `argparse`. You are not allowed to use these as options. They are used to create the documentation for your program.

When your program prints the correct usage, you can get the `word` argument inside the `main` function like so:

```
def main():
    args = get_args()
    word = args.word
```

Make your program print the given word:

```
$ ./crowsnest.py narwhal
narwhal
```

And now run your tests. You should still be passing two and failing the third. Let's read the test failure:

```
===== FAILURES =====
test_consonant

def test_consonant():
    """brigatine -> a brigatine"""

    for word in consonant_words:
        out = getoutput('{} {}'.format(prg, word))
>       assert out.strip() == template.format('a', word)
E           AssertionError: assert 'brigatine' == 'Ahoy, Captain, a brigatine off the
larboard bow!' (2)
E               - brigatine
E               + Ahoy, Captain, a brigatine off the larboard bow!
```

- ① The line starting with > shows code that produced an error. The output of the program is compared to an expected string. Since it didn't match, the assert produces an exception.
- ② This line starts with E to indicate the "error."
- ③ The line starting with a - is what the test got when it ran with the argument 'brigatine' — it got back just the word "brigatine."
- ④ The line starting with the + is what the test expected, "Ahoy, Captain, a brigatine off the larboard bow!"

So, we need to get the word into the "Ahoy" phrase. How can we do that?

3.1.4 Concatenating strings

Putting strings together is called "concatenating" or "joining" strings. To demonstrate, I'm going to enter some code directly into the Python interpreter. I want you to type along. No, really! Type everything you see, and try it for yourself.

Open a terminal and type `python3` or `ipython` to start a REPL, a "Read-Evaluate-Print-Loop" because Python will *read* each line of input, *evaluate* and *print* the results in a *loop*. Here's what it looks like on my system:

```
$ python3
Python 3.7.3 (default, Mar 27 2019, 16:54:48)
[Clang 4.0.1 (tags/RELEASE_401/final)] :: Anaconda, Inc. on darwin
Type "help", "copyright", "credits" or "license" for more information.
>>>
```

You may also like to use Python's IDLE (integrated development and learning environment) program, or you can use Jupyter Notebooks to interact with the language. I'll stick to the `python3` REPL for showing code examples. To exit the REPL, either

type `quit()` or CTRL-d (the Control key plus the d).

The `>>>` is a prompt where you can type code. Let's start off by assigning the variable `word` to the value "narwhal." In the REPL, type `word = 'narwhal'<Enter>`:

```
>>> word = 'narwhal'
```

Note that you can put as many (or no) spaces around the `=` as you like, but convention and readability (and tools like `pylint` or `flake8` that help you find errors in your code) would ask you to use exactly one space on either side. If you type `word<Enter>`, Python will print the current value of `word`:

```
>>> word
'narwhal'
```

Now type `werd<Enter>`:

```
>>> werd
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
NameError: name 'werd' is not defined
```

WARNING

There is no `werd` variable because we haven't set `werd` to be anything. Using an undefined variable causes an exception that will crash your program. Python will happily create a `werd` for you when you assign it a value.

In Python, there are many ways we can concatenating strings. The `+` operator can be used to join strings together:

```
>>> 'Ahoy, Captain, a ' + word + ' off the larboard bow!'
'Ahoy, Captain, a narwhal off the larboard bow!'
```

If you change your program to print that instead of just the `word`, you should be able to four tests:

<code>test.py::test_exists</code>	PASSED	[16%]
<code>test.py::test_usage</code>	PASSED	[33%]
<code>test.py::test_consonant</code>	PASSED	[50%]
<code>test.py::test_consonant_upper</code>	PASSED	[66%]
<code>test.py::test_vowel</code>	FAILED	[83%]

If we look closely at the failure, you'll see this:

```
E           - Ahoy, Captain, a aviso off the larboard bow!
E           + Ahoy, Captain, an aviso off the larboard bow!
E           ?           +
```

So we hard-coded the "a" before the word, but we really need to figure out whether to put "a" or "an" depending on whether the word starts with a vowel. How can we do that?

3.1.5 Variable types

Before we go much further, I need to take a small step back and point out that

our word variable is a "string." Every variable in Python has a "type" that describes the kind of data they hold. Because we put the value for word in quotes ('narwhal'), the word holds a "string" which Python represents with a class called str. (A "class" is a collection of code and functions that we can use.)

The type function will tell us what kind of data Python thinks this is:

```
>>> type(word)
<class 'str'>
```

Whenever you put a value in single (' ') or double quotes (""), Python will interpret it as a str:

```
>>> type("submarine")
<class 'str'>
```

WARNING

If you forget the quotes, then Python will look for some variable or function by that name. If there is no variable or function by that name, it will cause an exception.

```
>>> word = narwhal
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
NameError: name 'narwhal' is not defined
```

Exceptions are bad, and we will try to write code that avoids them or at least knows how to handle them gracefully.

3.1.6 Getting just part of a string

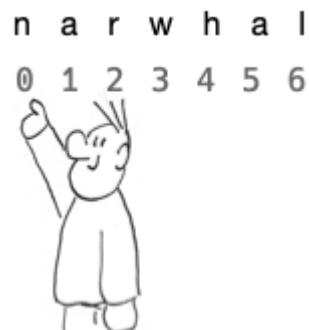
Back to our problem! We need to put either "a" or "an" in front of the word we're given based on whether the first character of word is a vowel or a consonant. In Python, we use square brackets and an *index* to get an individual character from a string. The index is the numeric position of an element in a sequence, and we must remember that indexing starts at 0.

```
>>> word[0]
'n'
```

Or directly on a string:

```
>>> 'narwhal'[0]
'n'
```

You can use this with a variable:



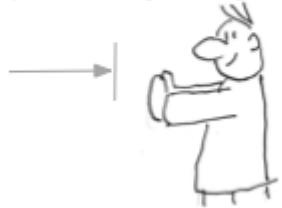
n a r w h a l
0 1 2 3 4 5 6



You can also use negative numbers to count backwards from the end, so the last index is also -1:

```
>>> word[-1]
'1'
```

n a r w h a l
0 1 2 3 4 5 6



This means that the last index is *one less than the length*, which is often confusing. The length of "narwhal" is 7, but the last character is found at index 6:

```
>>> word[6]
'1'
```

n a r w h a l
-6 -5 -4 -3 -2 -1



You can use the "slice" notation [start:stop] to get a range of characters. Both start and stop are optional. The default value for start is 0 (the beginning of the string), and the stop value is *not inclusive*:

```
>>> word[:3]
'nar'
```

And the default value for stop is the end of the string:

```
>>> word[3:]
'whal'
```

In the next chapter, we'll see that this is the same syntax for slicing lists. A string is (sort of) a list of characters, so this isn't too strange.

3.1.7 Finding help in the REPL

The class `str` has a ton of functions we can use to handle strings, but what are they? A

large part of programming is knowing how to ask questions and where to look for answers. A common refrain you may hear is "RTFM" — Read the Fine Manual. The Python community has created reams of documentation which are all available at docs.python.org/3/. You will need to refer to the documentation constantly to remind yourself how to use certain functions.

The docs for the string class are here:

docs.python.org/3/library/string.html

I prefer to read the docs directly inside the REPL by typing `help(str)`:

```
>>> help(str)
```

Inside the help, you move up and down in the text using the up and down cursor arrows on your keyboard. You can also press the <Space> bar or CTRL-f to jump forward to the next page, and CTRL-b to jump backward. You can search through the documentation by pressing / and then the text you want to find. If you press n (for "next") after a search, you will jump to the next place that string is found. To leave the help, press q (for "quit").



3.1.8 String methods



Now that we know `word` is a string (`str`), we have all these incredibly useful *methods* we can call on the variable. (A "method" is a function that belongs to a variable like `word`.) For instance, if I wanted to shout about the fact that we have a "narwhal," I could print it in UPPERCASE LETTERS. If I search through the help, I see there is a function called `upper`. Here is how to call it:

```
>>> word.upper()
'NARWHAL'
```

You must include the parentheses () or else you're talking about the *function itself*:

```
>>> word.upper
<built-in method upper of str object at 0x10559e500>
```

That will actually come in handy later when we use functions like `map` and `filter`, but for now we want Python to *execute* or *call* the `upper` function on the variable `word`, so we add the parens. Note that the function returns an uppercase version of the word

but *does not* change the value of `word` itself:

```
>>> word
'narwhal'
```

There is another `str` function with "upper" in the name called `isupper`. The name helps you know that this will return a True/False type answer. Let's try it:

```
>>> word.isupper()
False
```

We can chain methods together like so:

```
>>> word.upper().isupper()
True
```

That makes sense. If I convert the `word` to uppercase, then `isupper` is `True`.

I find it odd that the `str` class does not include a method to get the length of a string. For that, we use a separate function called `len`, short for "length":

```
>>> len('narwhal')
7
```

Are you typing all this into Python yourself? I recommend you do! Find other methods in the `str` help and try them out.



3.1.9 String comparisons

So now you know how to get the first letter of `word` by using `word[0]`. Let's assign it to the variable `char`:

```
>>> word = 'octopus'
>>> char = word[0]
>>> char
'o'
```

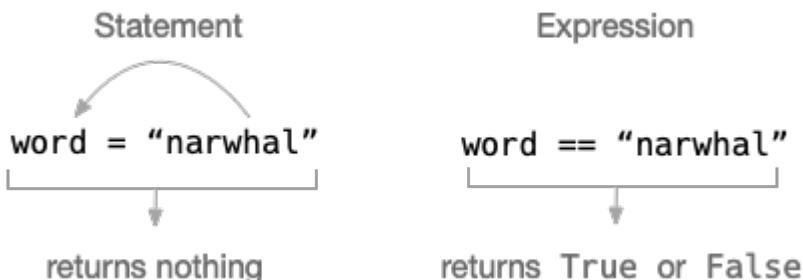
Now we need to figure out if `char` is a vowel or a consonant. We'll say that letters "a," "e," "i," "o," and "u" make up our set of "vowels." You can use `==` to compare strings:

```
>>> char == 'a'
False
>>> char == 'o'
True
```

NOTE

Be careful to always use one equal sign (=) when *assigning a value* to a variable, like `word = 'narwhal'` and two equal signs (==, which, in my head, I say "equal-equal") when you *compare two values* like `word == 'narwhal'`. The first is a statement that changes the value of `word`, and the second is an *expression* that returns True or False.

Figure 3.2. An expression returns a value. A statement does not.



We need to compare our `char` to *all* the vowels. You can use `and` and `or` in such comparisons and they will be combined according to standard Boolean algebra:

```
>>> char == 'a' or char == 'e' or char == 'i' or char == 'o' or char == 'u'
True
```

What if the word is "Octopus" or "OCTOPUS"?

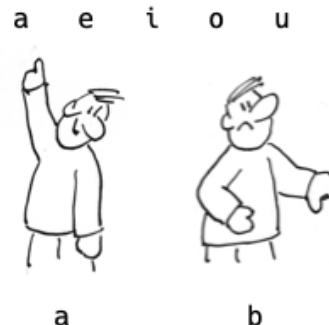
```
>>> word = 'OCTOPUS'
>>> char = word[0]
>>> char == 'a' or char == 'e' or char == 'i' or char == 'o' or char == 'u'
False
```

Do we have to make 10 comparisons in order to check the uppercase versions, too? What if we were to lowercase `word[0]`? Remember, that `word[0]` returns a `str`, and so we can chain other `str` methods onto that:

```
>>> word = 'OCTOPUS'
>>> char = word[0].lower()
>>> char == 'a' or char == 'e' or char == 'i' or char == 'o' or char == 'u'
True
```

An easier way to determine if `char` is a vowel would be to use Python's `x in y` construct where we want to know if the value `x` is in the collection `y`. We can ask if the letter '`a`' is in the longer string '`aeiou`':

```
>>> 'a' in 'aeiou'
True
```



But the letter '`b`' is not:

```
>>> 'b' in 'aeiou'
False
```

Let's use that to test the first character of the lowercased word (which is '`o`'):

```
>>> word = 'OCTOPUS'
>>> word[0].lower() in 'aeiou'
True
```

3.1.10 Conditional branching

Once you have figured out if the first letter is a vowel, you will need to select an article. We'll use a very simple rule that, if the word starts with a vowel, choose "`an`," otherwise choose "`a`." This misses exceptions like when the initial "`h`" in a word is silent, for instance, we say "`a hat`" but "`an honor`". Nor will we consider when an initial vowel has a consonant sound as in "`union`" where the "`u`" sounds like a "`y`."

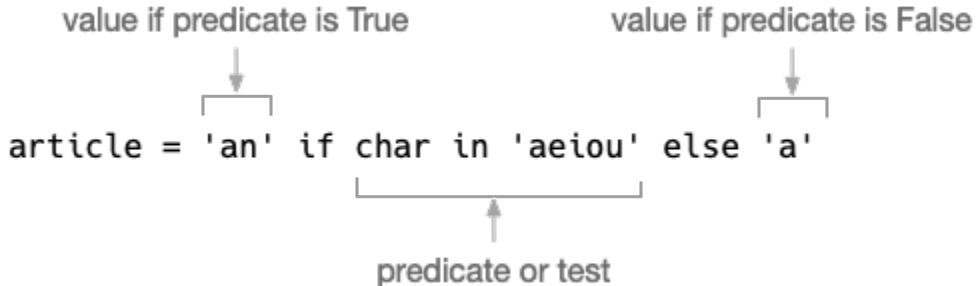
We can create a new variable called `article` that we will set to the empty string and then use an `if/else` statement to figure out what to put in it:

```
>>> article = ''                                ①
>>> if word[0].lower() in 'aeiou': ②
...     article = 'an'                           ③
... else:                                         ④
...     article = 'a'                            ⑤
... 
```

- ① Initialize `article` to the empty string.
- ② Check if the first, lowercased character of `word` is a vowel.
- ③ If it is, set `article` to '`an`'
- ④ Otherwise,
- ⑤ Set `article` to '`a`'.

Here is a much shorter way to write that with an `if expression` (expressions return values, statements do not). The `if` expression is written a little backwards. First comes the value if the test (or "predicate") is `True`, then the predicate, then the value if the

predicate is False.



This way is also safer because the if expression is *required* to have the else. There's no chance that we could forget to handle both cases:

```
>>> article = 'an' if char in 'aeiou' else 'a'
```

Let's verify that we have the correct article:

```
>>> article
'an'
```

3.1.11 String formatting

Now we have two variables, `article` and `word` that need to be incorporated into our "Ahoy!" phrase. We saw earlier that we can use the plus sign (+) to concatenate strings. Another method to create new strings from other strings is to use the `str.format` method. To do so, you create a string template with curly brackets {} that indicate placeholders for values. The values that will be substituted go as arguments to the `format`, and they are substituted in the same order that the {} appear:

'Ahoy, Captain, {} {} off the larboard bow!'.format(article, word)

Here it is in code:

```
>>> 'Ahoy, Captain, {} {} off the larboard bow!'.format(article, word)
'Ahoy, Captain, an octopus off the larboard bow!'
```

Another method uses the special "f-string" where you can put the variables directly into the {} brackets. It's a matter of taste which one you choose.

```
>>> f'Ahoy, Captain, {article} {word} off the larboard bow!'
'Ahoy, Captain, an octopus off the larboard bow!'
```

Python variables are very variable

A note that in some programming languages, you have to declare the variable's name and what *type* of data it will hold. If a variable is declared to be a number, then it can never hold a value of a different type like a string. This is called *static typing* because the type of the variable can never change. Python is a *dynamically typed* language because you do not have to declare a variable or what kind of data the variable will hold. You can change the value and type of data at any time. This could be either great or terrible news. As Hamlet says, "There is nothing either good or bad, but thinking makes it so."



Hints:

- Start your program with `new.py` and fill in the `get_args` with a single position argument called `word`.
- You can get the first character of the word by indexing it like a list, `word[0]`.
- Unless you want to check both upper- and lowercase letters, you can use either the `str.lower` or `str.upper` method to force the input to one case for checking if the first character is a vowel or consonant.
- There are fewer vowels (five, if you recall) than consonants, so it's probably easier to check if the first character is one of those.
- You can use the `x in y` syntax to see if the element `x` is in the collection `y` where "collection" here is a `list`.
- Use the `str.format` or f-strings to insert the correct article for the given word into the longer phrase.
- Run `make test` (or `pytest -xv test.py`) *after every change to your program* to ensure your program compiles and is on the right track.

Now go write the program before you turn the page and study a solution! Look alive, you ill-tempered shababoon!

3.2 Solution

```
#!/usr/bin/env python3
"""Crow's Nest"""

import argparse

# -----
def get_args():
    """Get command-line arguments"""
    (1)
```

```

parser = argparse.ArgumentParser(
    description="Crow's Nest -- choose the correct article",
    formatter_class=argparse.ArgumentDefaultsHelpFormatter) ②
    ③
    ④

parser.add_argument('word', metavar='str', help='A word') ⑤

return parser.parse_args() ⑥

# -----
def main(): ⑦
    """Make a jazz noise here"""

    args = get_args() ⑧
    word = args.word ⑨
    article = 'an' if word[0].lower() in 'aeiou' else 'a' ⑩

    print(f'Ahoy, Captain, {article} {word} off the larboard bow!') ⑪

# -----
if __name__ == '__main__': ⑫
    main() ⑬

```

- ① Defines the function `get_args` to handle the command-line arguments. I like put this first so I can see it right away when I'm reading the code.
- ② The `parser` will do the work of parsing the arguments.
- ③ The `description` shows in the usage to describe what the program does.
- ④ Show the default values for each parameter in the usage.
- ⑤ Define a positional argument called `word`.
- ⑥ The result of parsing the arguments will be returned to line 24.
- ⑦ Defines the `main` function where the program will start.
- ⑧ `args` contains the return value from the `get_args` function.
- ⑨ Put the `args.word` value from the arguments into the variable `word`.
- ⑩ Choose the correct `article` using an `if` expression to see if the lowercased, first character of `word` is or is not in the set of vowels.
- ⑪ Print the output string using an f-string to interpolate the `article` and `word` variables inside the string.
- ⑫ Check if we are in the "main" namespace, which means the program is *running*.
- ⑬ If so, call the `main()` function to make the program start.

3.3 Discussion

I'd like to stress that the preceding is *a* solution, not *the* solution. There are many ways to express the same idea in Python. As long as your code passes the test suite, it is correct.

That said, I created my program with `new.py` which automatically gives me two functions:

1. `get_args` where I define the arguments to the program
2. `main` where the program starts

Let's talk about these two functions.

3.3.1 Defining the arguments with get_args

I prefer to put the `get_args` function first so that I can see right away what the program expects as input. You don't have to define this as a separate function. You could put all this code inside `main`, if you prefer. Eventually our programs are going to get longer, though, and I think it's nice to keep this as a separate idea. Every program I present will have a `get_args` function that will handle defining and validating the input.

Our program specifications (the "specs") say that the program should accept one positional argument. I changed the '`positional`' argument name to '`word`' because I'm expecting a single word:

```
parser.add_argument('word', metavar='str', help='Word')
```

I would really recommend you never leave the "positional" argument named '`positional`' because it is an entirely undescriptive term. Naming your variables *what they are* will make your code more readable. Since the program doesn't need any of the other options created by `new.py`, you can delete the rest of the `parser.add_argument` calls. The `get_args` function will return the result of parsing the command line arguments which I put into the variable `args`:

```
return parser.parse_args()
```

If `argparse` is not able to parse the arguments — for example, there are none — it will never return from `get_args` but will instead print the "usage" for the user and exit with an error code to let the operating system know that the program exited without success. (In the Unix world, an exit value of 0 means there were 0 errors. Anything other than 0 is considered an error.)

3.3.2 The main thing

Many programming languages will automatically start from the `main` function, so I always define a `main` function and start my programs there. This is not a requirement, just how I like to write programs. Every program I present will start with the `main` function which will first call `get_args` to get the program's inputs:

```
args = get_args()
```

I can now access the `word` by call `args.word`. Note the lack of parentheses. It's not `args.word()` because is not a function call. Think of `args.word` like a slot where the value of the "word" lives:

```
word = args.word
```

I like to work through my ideas using the REPL, so I'm going to pretend that `word` has been set to "octopus":

```
>>> word = 'octopus'
```

3.3.3 Classifying the first character of a word

To figure out whether the article I choose should be a or an, I need to look at the first character of the word which we can get like so. In the introduction, we used this:

```
>>> word[0]
'o'
```

I can check if the first character is in the string of vowels, both lower- and uppercase:

```
>>> word[0] in 'aeiouAEIOU'
True
```

I can make this shorter, however, if I use `word.lower` function so I'd only have to check the lowercase vowels:

```
>>> word[0].lower() in 'aeiou'
True
```

Remember that the `x in y` form is a way to ask if element `x` is in the collection `y`. You can use it for letters in a longer string (like the vowels):

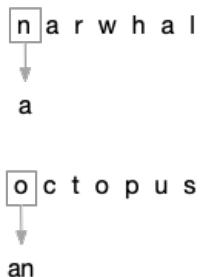
```
>>> 'a' in 'aeiou'
True
```

Or for a string in list of other strings:

```
>>> 'tanker' in ['yatch', 'tanker', 'vessel']
True
```

We can use membership in the "vowels" as a condition to choose "an," otherwise we choose "a": As mentioned in the introduction, the `if` expression is the shortest and safest for a "binary" choice (where there are only two possibilities):

```
>>> article = 'an' if word[0].lower() in 'aeiou' else
'a'
>>> article
'an'
```



The safety comes from the fact that Python will not even run this program if you forget the `else`. We can change the `word` to "galleon" and check that it still works:

```
>>> word = 'galleon'
>>> article = 'an' if word[0].lower() in 'aeiou' else 'a'
>>> article
'a'
```

3.3.4 Printing the results

Finally we need to print out the phrase with our `article` and `word`. As noted in the introduction, you can use `str.format`:

```
>>> article = 'a'
>>> word = 'ketch'
>>> print('Ahoy, Captain, {} {} off the larboard bow!'.format(article, word))
Ahoy, Captain, a ketch off the larboard bow!
```

Python's f-strings will *interpolate* any code inside the {} placeholders, so variables get turned into their contents:

```
>>> print(f'Ahoy, Captain, {article} {word} off the larboard bow!')
Ahoy, Captain, a ketch off the larboard bow!
```

However you chose to print out the article and word is fine as long as it passes the tests.

3.3.5 Running the test suite

"A computer is like a mischievous genie. It will give you exactly what you ask for, but not always what you want. - Joe Sondow"

Computers are a bit like bad genies. They will do exactly what you tell them but not necessarily what you *want*. In an episode of *The X-Files*, the character Mulder wishes for peace on Earth and a genie removes all humans but him. Tests are what we can use to verify that our programs are doing what we *actually* want them to do. Tests they can never prove that our program is truly free from errors, only that the bugs we imagined or found while writing the program no longer exist. Still, we write and run tests because they are really quite effective and much better than not doing so.

This is the idea behind "test-driven development":

- We can write tests *even before* we write the software.
- We run the tests to verify that our as-yet-unwritten software definitely fails to deliver on some task.
- Then we write the software to fulfill the request.
- Then we run the test to check that it now *does* work.
- We keep running all the tests to ensure that, when we add some new code, we do not break existing code.

3.3.6 Passing Tests

I would encourage you to look at the `test.py` program to see how it is testing your program. Eventually I'll recommend you write your own tests, but for now just see what's being expected of your code. I use the `pytest` module to write tests. There are other testing frameworks in Python, but I find `pytest` to be relatively easy to use. The `pytest` module will run any functions that begin with `test_` in the order they are found in the source code.

The first `test_` function is `test_exists` that uses the `assert` function to check if the `crowsnest.py` program exists. This is why your program must be named '`crowsnest.py`'. It must exist as this name so that we can run it and check the output:

```
prg = './crowsnest.py'

def test_exists():
    """exists"""

    assert os.path.isfile(prg)
```

The next is `test_usage` to check if the program will print something that looks like "usage" when run with `-h` and `--help` flags:

```
def test_usage():
    """usage"""

    for flag in ['-h', '--help']:
        rv, out = getstatusoutput(f'{prg} {flag}')
        assert rv == 0
        assert out.lower().startswith('usage')
```

Inside `test.py`, there are two lists of words starting with consonants and others starting with vowels.

```
consonant_words = [
    'brigantine', 'clipper', 'dreadnought', 'frigate', 'galleon', 'haddock',
    'junk', 'ketch', 'longboat', 'mullet', 'narwhal', 'porpoise', 'quay',
    'regatta', 'submarine', 'tanker', 'vessel', 'whale', 'xebec', 'yatch',
    'zebrafish'
]
vowel_words = ['aviso', 'eel', 'iceberg', 'octopus', 'upbound']
```

There is also a string template for what the program should print:

```
template = 'Ahoy, Captain, {} {} off the larboard bow!'
```

The `test_consonant` test runs through each of the `consonant_words` and checks if the program puts an "a" in front of the word.

```
def test_consonant():
    """brigantine -> a brigantine"""

    for word in consonant_words:
        out = getoutput(f'{prg} {word}')
        assert out.strip() == template.format('a', word)
```

The next function does the same thing but uses a capitalized version of the consonant word. The next two tests then use the `vowel_words`, checking both lower- and uppercase versions.

When all tests are passing, this is the output you should see (some output elided):

```
$ make test
pytest -xv test.py
=====
test session starts =====
...
collected 6 items

test.py::test_exists PASSED [ 16%]
test.py::test_usage PASSED [ 33%]
test.py::test_consonant PASSED [ 50%]
```

```
test.py::test_consonant_upper PASSED [ 66%]
test.py::test_vowel PASSED [ 83%]
test.py::test_vowel_upper PASSED [100%]

===== 6 passed in 2.28 seconds =====
```

3.4 Summary

- All Python's documentation is available on docs.python.org/3/ and with the `help` command in the REPL.
- Variables in Python are dynamically typed according to whatever value you assign them and come into existence when you assign a value to them.
- Strings have methods like `upper` and `isupper` that you can call to alter them or get information.
- You can get parts of a string by using square brackets and indexes like `[0]` for the first letter or `[-1]` for the last.
- You can concatenate strings with the `+` operator.
- The `str.format` method allows you to create a template with `{}` placeholders that get filled in with the arguments.
- F-strings like `f'{article} {word}'` allow variables and code to go directly inside the brackets.
- The `x in y` expression will report if the value `x` is present in the collection `y`.
- Statements like `if/else` do not return a value while expressions like `x if y else z` do return a value.
- Test-driven development is a way to ensure programs meet some minimum criteria of correctness. Every feature of a program should have tests, and writing and running test suites should be an integral part of writing programs.

3.5 Going Further

- Have your program match the case of the incoming word, e.g., "an octopus" and "An Octopus." Copy an existing `test_` function in the `test.py` to verify that your program works correctly while still passing all the other tests. Try writing the test first, then make your program pass the test. That's *test-driven development*!
- Accept a new parameter that changes "larboard" (the left side of the boat) to "starboard" (the right side ⁴ [\[a\]](#)). You could either make an option called `--side` that defaults to "larboard," or you could make a `--starboard` flag that, if present, changes the side to "starboard."



⁴ "Starboard" has nothing to do with stars but with the "steering board" or a rudder which typically would be on the right-side of the boat for right-handed sailors!

- The provided tests only give you words that start with an actual alphabetic character. Expand your code to handle words that start with numbers or punctuation. Should your program reject these? Add more tests to ensure that your program does what you intend.

Going on a picnic: Working with lists



Writing code makes me hungry! Let's write a program to consider some tasty foods we'd like to eat. So far we've handled *one* of something like a name to say "hello" to or a nautical-themed object to point out. In this program, we want to eat one or more foods which we will store in a *list*, a variable that can hold any number of items. We use lists all the time in life. Maybe it's your top-five favorite songs, your birthday wish-list, or a bucket list of the best types of buckets.

In this exercise, we're going on a picnic, and we want to print a list of items to bring. You will learn to:

- Write a program that accepts multiple positional arguments.
- Use `if`, `elif`, and `else` to handle conditional branching with three or more options.
- Find and alter items in a list.
- Sort and reverse lists.
- Format a list into a new string.

The items will be passed as positional arguments. When there is only one item, you'll print that:

```
$ ./picnic.py salad
You are bringing salad.
```



What? Who just brings salad on a picnic? When there are two items, you'll put "and" in between them:

```
$ ./picnic.py salad chips
You are bringing salad and chips.
```

Hmm, chips. That's an improvement. When there are three or more items, you will separate the items with commas:

```
$ ./picnic.py salad chips cupcakes
You are bringing salad, chips, and
cupcakes.
```



There's one other twist. You will also need to accept a `--sorted` argument that will require you to sort the items before you print them, but we'll deal with that in a bit. So, your Python program must:

- Store one or more positional arguments in a `list`
- Count the number of arguments
- Possibly modify the list like maybe to sort the items
- Use the list to print a new a string that formats the arguments according to how

many items there are.

How should we begin?

4.1 Starting the program

I will always recommend you start programming either by running new.py or by copying template/template.py to the program name. This time the program should be called picnic.py, and we need to create it in the picnic directory:

```
$ cd picnic
$ new.py picnic.py
Done, see new script "picnic.py."
```

Now run make test or pytest -xv test.py. You should pass the first two tests (program exists, program creates usage), and fail the third:

test.py::test_exists PASSED	[14%]
test.py::test_usage PASSED	[28%]
test.py::test_one FAILED	[42%]

The rest of the output is complaining about the fact that the test expected "You are bringing chips" but got something else:

```
===== FAILURES =====
test_one

def test_one():
    """one item"""

    out = getoutput(f'{prg} chips')
>   assert out.strip() == 'You are bringing chips.' ①
E   assert 'str_arg = "...nal = "chips"' == 'You are bringing chips.'
E       + You are bringing chips. ②
E       - str_arg = "" ③
E       - int_arg = "0" ④
E       - file_arg = ""
E       - flag_arg = "False"
E       - positional = "chips"

test.py:31: AssertionError
===== 1 failed, 2 passed in 0.56 seconds =====
```

- ① The program is being run with the argument "chips."
- ② This is the line that is causing the error. The output is tested to see if it is equal (==) to the string "You are bringing chips."
- ③ The line starting with a + sign is showing what was expected.
- ④ The lines starting with the - sign are showing what was returned by the program.

Let's run the program with the argument "chips" and see what it gets:

```
$ ./picnic.py chips
str_arg = ""
int_arg = "0"
file_arg = ""
```

```
flag_arg = "False"
positional = "chips"
```

Right, that's not correct at all! Remember, the template doesn't have the *correct* arguments, just some examples, so the first thing we need to do is to fix the `get_args` function. Here is what your program should print a usage statement if given *no arguments*:

```
$ ./picnic.py
usage: picnic.py [-h] [-s] str [str ...]
picnic.py: error: the following arguments are required: str
```

And here is the usage for the `-h` or `--help` flags:

```
$ ./picnic.py -h
usage: picnic.py [-h] [-s] str [str ...]

Picnic game

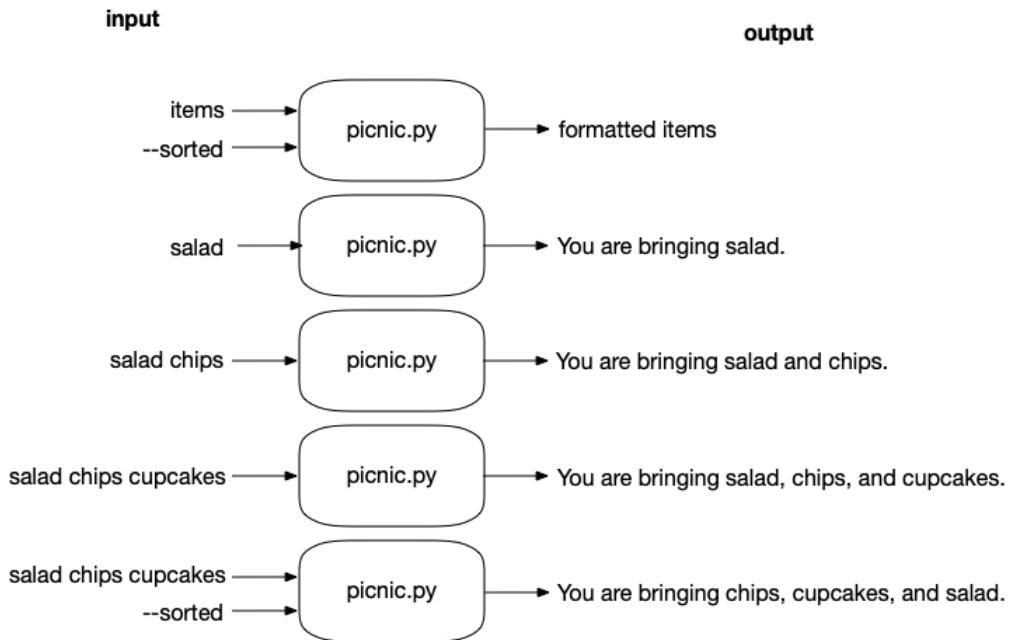
positional arguments:
  str          Item(s) to bring

optional arguments:
  -h, --help    show this help message and exit
  -s, --sorted  Sort the items (default: False)
```

We need a single positional argument and an optional flag called `--sorted`. Modify your `get_args` until it produces the above output. Note that there should be one or more of the `item` parameter, so you should define it with `nargs='+'`. Refer to the section "One or more of the same positional arguments" in the chapter 2.

4.2 Writing `picnic.py`

Here is a tasty diagram of the inputs and ouputs for the `picnic.py` program we'll write:



The program should accept one or more "positional" arguments as the items to bring on a picnic as well as a `-s` or `--sorted` flag to indicate whether or not to sort the items. The output will be "You are bringing" and then the list of items formatted according the following rules:

1. If one item, state the item:

```
$ ./picnic.py chips
You are bringing chips.
```

2. If two items, put "and" in between the items. Note that "potato chips" is just *one string* that happens to contain *two words*. If we leave out the quotes, then there would be three arguments to the program. Note that it doesn't matter here if we use single or double quotes:

```
$ ./picnic.py "potato chips" salad
You are bringing potato chips and salad.
```

3. If three or more items, place a comma and space between each item and the word "and" before the final element. Don't forget the comma before the "and" (sometimes called the "Oxford comma") because your author was an English lit major and, while I may have finally stopped using two spaces after the end of a sentence, you can pry the Oxford comma from my cold, dead hands:

```
$ ./picnic.py "potato chips" salad soda cupcakes
You are bringing potato chips, salad, soda, and cupcakes.
```

Be sure to sort if given the `-s` or `--sorted` flag:

```
$ ./picnic.py --sorted salad soda cupcakes
You are bringing cupcakes, salad, and soda.
```

In order to figure out how many items we have, how to sort and slice them, and how to format the output string, we need to talk about the `list` type in Python in order to solve this problem.

4.3 Stuff about lists

We briefly touched on the idea of a `list` in the `hello.py` program when we looked at `sys.argv`, the "argument vector" for a program. If we run this:

```
$ ./picnic.py salad chips cupcakes
```

Then the arguments `salad` `chips` `cupcakes` would be available in `sys.argv` as the list of the strings

```
['salad', 'chips', 'cupcakes']
```

Note that they would be in the same order as they were provided on the command line. Lists always keep their order!

Let's go into the REPL and create a variable called `items` to hold some tasty foods we plan to bring on our picnic. I really want you to type these commands yourself, too, whether in the `python3` REPL or `ipython` or a Jupyter Notebook. It's very important to interact in real time with the language!

To create a new, empty list, we can either use the `list()` function:

```
>>> items = list()
```

Or use empty square brackets:

```
>>> items = []
```

Check what Python says for the type. Yep, it's a `list`:

```
>>> type(items)
<class 'list'>
```

One of the first things we need to know is how many `items` we have for our picnic. Like a `str`, we can use `len` (length) to get the number of elements in `items`:

```
>>> len(items)
0
```

The length of an empty `list` is `0`.

4.3.1 Adding one element to a list

An empty list is not very useful. Let's see how we can add new items. We used `help(str)` in the last chapter to read the documentation about the string *methods*,

the functions that belong to every `str` in Python. Here I want you to use `help(list)` to find the `list` methods:

```
>>> help(list)
```

You'll see lots of "double-under" methods like `{dbl_}len{dbl_}`. Skip over those, and the first method we can find is `append`, which we can use to add items to the end of the list. If we evaluate our `items`, we see that the empty brackets tell us that it's empty:

```
>>> items
[]
```

Let's add "sammiches" to the end:

```
>>> items.append('sammiches')
```

Nothing happened, so how do we know that worked? Let's check the length. It should be 1:

```
>>> len(items)
1
```

Hooray! That worked. In the spirit of testing, use the `assert` method to verify that the length is 1. The fact that nothing happens is good. When an assertion fails, it triggers an exception that results in a lot of messages. Here, no news is good news:

```
>>> assert len(items) == 1
```

If you type `items``<Enter>` in the REPL, Python will show you the contents:

```
>>> items
['sammiches']
```

Cool, we added one element.

4.3.2 Adding many elements to a list

Let's try to add "chips" and "ice cream" to the `items`:

```
>>> items.append('chips', 'ice cream')
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: append() takes exactly one argument (2 given)
```

Here is one of those pesky exceptions, and these will cause your programs to *crash*, something we want to avoid at all costs. We see that `append()` takes exactly one argument, and we gave it two. If you look at the `items`, you'll see that nothing was added:

```
>>> items
['sammiches']
```

OK, so maybe we were supposed to give it a `list` of items to add? Let's try that:

```
>>> items.append(['chips', 'ice cream'])
```

Well, that didn't cause an exception, so maybe it worked? We would expect there to be 3 `items`, so let's use an assertion to check that:

```
>>> assert len(items) == 3
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
AssertionError
```

Another exception because `len(items)` is not 3! What is the length?

```
>>> len(items)
2
```

Only 2? Let's look at `items`:

```
>>> items
['sammiches', ['chips', 'ice cream']]
```

Check that out! Lists can hold any type of data like strings and numbers and even other lists. We asked `append` to add `['chips', 'ice cream']`, which is a list, and that's just what it did. Of course, it's not what we wanted.



Let's reset `items` so we can fix this. We can either use the `del` command to delete the element at index 1:

```
>>> del items[1]
```

Or reassign it a new value:

```
>>> items = ['sammiches']
```

If you read further into the help, you will find the `extend` method:

```
| extend(self, iterable, /)
|     Extend list by appending elements from the iterable.
```

Let's try that:

```
>>> items.extend('chips', 'ice cream')
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: extend() takes exactly one argument (2 given)
```

Well that's frustrating! Now Python is telling us that `extend()` takes exactly one argument which, if you refer to the help, should be an iterable. A list is something you can iterate (travel over from beginning to end), so that will work:

```
>>> items.extend(['chips', 'ice cream'])
```

Nothing happened. No exception, so maybe that worked? Let's check the length. It *should* be 3:

```
>>> assert len(items) == 3
```

Yes! Let's look at the `items` we've added:

```
>>> items
['sammiches', 'chips', 'ice cream']
```

Great! This is sounding like a pretty delicious outing.

If you know everything that will go into the `list`, you can create it like so:

```
>>> items = ['sammiches', 'chips', 'ice cream']
```

The `append` and `extend` methods add new elements to the *end* of a given `list`. The `insert` method allows you to place new items at any position by specifying the index. I can use the index `0` to put a new element at the beginning of `items`:

```
>>> items.insert(0, 'soda')
>>> items
['soda', 'sammiches', 'chips', 'ice cream']
```

In addition to `help(list)`, you can also find lots of great documentation here:

docs.python.org/3/tutorial/datastructures.html

I recommend you read over all the `list` functions so you get an idea of just how powerful this data structure is!

4.3.3 Indexing lists

So now we have a `list` of `items`. We know how to use `len` to find how many `items` there are, and now we need to know how to get parts of the `list` to format. Indexing a `list` in Python looks exactly the same as indexing a `str`. (This actually makes me a bit uncomfortable, so I tend to imagine a `str` as a `list` of characters and then I feel somewhat better.)



0	1	2	3
['soda', 'sammiches', 'chips', 'ice cream']			
-4	-3	-2	-1

All indexing in Python is zero-offset, so the first element of `items` is at index `items[0]`:

```
>>> items[0]
'soda'
```

If the index is negative, Python starts counting backwards from the end of the list. The index `-1` is the last element of the list:

```
>>> items[-1]
'ice cream'
```

You should be very careful when using indexes to reference elements in a list. This is unsafe code:

```
>>> items[10]
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
IndexError: list index out of range
```

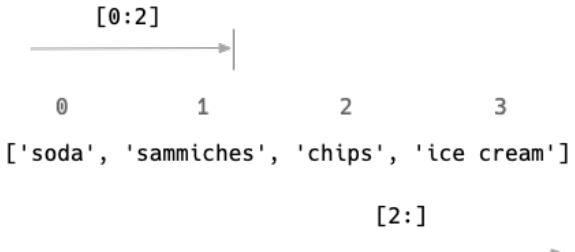
WARNING

Referencing an index that is not present will cause an exception.

We'll soon learn how to safely *iterate* or travel through a list so that we don't have to use indexes to get at elements.

4.3.4 Slicing lists

You can extract "slices" (sub-lists) of a list by using `list[start:stop]`. To get the first two items elements, you use `[0:2]`. Remember that the 2 is actually the index of the *third* element but *it's not inclusive*:



```
>>> items[0:2]
['soda', 'sammiches']
```

If you leave out `start`, it will be `0`, so this does the same thing:

```
>>> items[:2]
['soda', 'sammiches']
```

If you leave out `stop`, it will go to the end of the list:

```
>>> items[2:]
['chips', 'ice cream']
```

Oddly, it is completely *safe* for slices to use list indexes that don't exist. Here I can ask for all the elements from index `10` to the end even though there is nothing at index `10`. Instead of an exception, we get an empty list:

```
>>> items[10:]
[]
```

For our exercise, you're going to need to get the word "and" into the list if there are three or more elements. Could you use a list index to do that?

4.3.5 Finding elements in a list

Did we remember to pack the chips?! Often we want to know if some items is in a list. The `index` method will return the location of an element in a list:

```
>>> items.index('chips')
2
```

WARNING

`list.index` is unsafe code because it will cause an exception if the argument is not present in the list!

See what happens if we check for the fog machine:

```
>>> items.index('fog machine')
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ValueError: 'fog machine' is not in list
```

You should never use `index` unless you have first verified that an element is present. The `x in y` that we used in "Crow's Nest" to see if a letter was in the list of vowels can also be used for lists. We get back a `True` value if `x` is in the collection of `y`:

```
>>> 'chips' in items
True
```

I hope they're salt and vinegar chips.

The same returns `False` if it is not present:

```
>>> 'fog machine' in items
False
```

We're going to need to talk to the planning committee. What's a picnic without a fog machine?



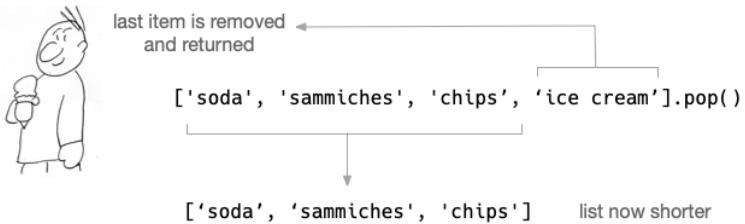
4.3.6 Removing elements from a list

We've seen that we can use the `del` function to delete a list element by index. The `list.pop` method will remove *and return* the element at the index. By default it will remove the *last* item (-1):

```
>>> items.pop()
'soda'
```

If we look at the list, we will see it's shorter by one:

```
>>> items
['soda', 'sammiches',
'chips']
```

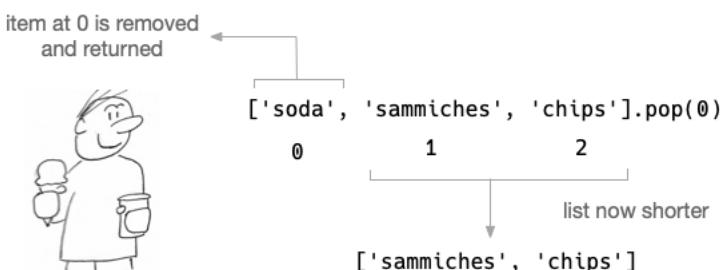


You can use an item's index to remove an element at a particular location. For instance, we can use `0` to remove the first element:

```
>>> items.pop(0)
'soda'
```

And now our list is shorter still:

```
>>> items
['sammiches', 'chips']
```



You can also use the `list.remove` method to remove the first occurrence of a given item:

```
>>> items.remove('chips')
>>> items
['sammiches']
```

item is removed
and returned



['sammiches', 'chips'].remove('chips')

list now shorter

['sammiches']

WARNING | The `list.remove` will cause an exception if the element is not present.

If we try to remove the chips again, we get an exception:

```
>>> items.remove('chips')
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ValueError: list.remove(x): x not in list
```

So don't use this code unless you've verified that a given element is in the list:

```
item = 'chips'
if item in items:
    items.remove(item)
```



4.3.7 Sorting and reversing a list

If the `--sorted` flag is present, we're going to need to sort the `items`. You might notice in the help documentation that two methods, `list.reverse` and `list.sort` stress that they work *IN PLACE*. That means that the `list` itself will be either reversed or sorted and *nothing will be returned*. So, given this list:

```
>>> items = ['soda', 'sammiches', 'chips', 'ice cream']
```

The `sort` method will return nothing:

```
>>> items.sort()
```

But if you inspect `items`, you will see they have been sorted alphabetically:

```
>>> items
['chips', 'ice cream', 'sammiches', 'soda']
```

None ← `items.sort()`

↓
items are sorted,
and nothing is returned

Note that Python will sort a `list` of numbers *numerically*, so we've got that going for us, which is nice:

```
>>> sorted([4, 2, 10, 3, 1])
[1, 2, 3, 4, 10]
```

As with `list.sort`, we see nothing on the `list.reverse` call:

```
>>> items.reverse()
```

But the `items` are now in the opposite order:

```
>>> items
['soda', 'sammiches', 'ice cream', 'chips']
```

The `list.sort` and `list.reverse` methods are easily confused with the sorted and reversed functions. The sorted function accepts a list as an argument and returns a new list:

```
>>> items = ['soda', 'sammiches', 'chips', 'ice cream']
>>> sorted(items)
['chips', 'ice cream', 'sammiches', 'soda']
```

It's crucial to note that the sorted function *does not alter* the list:

```
>>> items
['soda', 'sammiches', 'chips', 'ice cream']
```

The `list.sort` method is a function that belongs to the list. It can take arguments that affect the way the sorting happens. Let's look at the help(`list.sort`):

```
sort(self, /, *, key=None, reverse=False)
    Stable sort *IN PLACE*.
```

So we could also sort the `items` in reverse like so:

```
>>> items.sort(reverse=True)
```

And now they look like this:

```
>>> items
['soda', 'sammiches', 'ice cream', 'chips']
```



The reversed function works a bit differently:

```
>>> reversed(items)
<list_reverseiterator object at 0x10e012ef0>
```

I bet you were expecting to see a new list with the `items` in reverse? This is an example of a *lazy* function in Python. The process of reversing a list might take a while, so Python is showing that it has generated an "iterator object" that will provide the reversed list just as soon as we actually need the elements. We can do that in the REPL by using the `list` function to evaluate the iterator:

```
>>> list(reversed(items))
['ice cream', 'chips', 'sammiches', 'soda']
```

As with the sorted function, the original `items` itself remain unchanged:

```
>>> items
['soda', 'sammiches', 'chips', 'ice cream']
```

If you use the `list.sort` method instead of the `sorted` function, you might end up deleting your data. Imagine you wanted to set your `items` equal to the sorted list of `items` like so:

```
>>> items = items.sort()
```

What is in `items` now? If you print the `items` in the REPL, you won't see anything useful, so inspect the type:

```
>>> type(items)
<class 'NoneType'>
```

It's no longer a `list`. We set it equal to the result of called `items.sort()` method that works on the `list` *in-place* and returns `None`. I would note that I tend to not use the `sort/reverse` methods because I don't generally like to mutate my data. I would tend to do something like this:

```
>>> items = ['soda', 'sammiches', 'chips', 'ice cream']
>>> sorted_items = sorted(items)
>>> sorted_items
['chips', 'ice cream', 'sammiches', 'soda']
```

Now I have explicitly named a `sorted_items` list, and the original `items` has not been altered.

If the `--sorted` flag is given to your program, you will need to sort your items in order to pass the test. Will you use `list.sort` or the `sorted` function?

4.3.8 Lists are mutable

As we've seen, we can change a `list` quite easily. The `list.sort` and `list.reverse` methods change the whole list, but you can also change any single element by referencing it by index. Maybe we make our picnic slightly healthier by changing out the chips for apples:

```
>>> items
['soda', 'sammiches', 'chips', 'ice cream']
>>> if 'chips' in items:
...     idx = items.index('chips') ①
...     items[idx] = 'apples'        ②
...
...
```

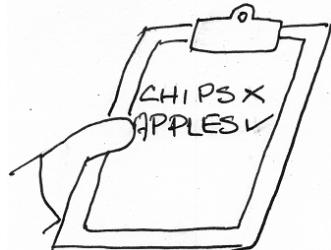
- ① See if the string 'chips' is in the list of `items`.
- ② Assign the index of 'chips' to the variable `idx`.
- ③ Use the index `idx` to change the element to 'apples'.

Let's look at `items` to verify:

```
>>> items
['soda', 'sammiches', 'apples', 'ice cream']
```

We can also write a couple of tests:

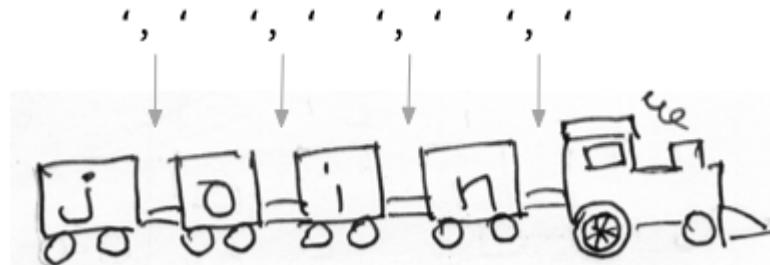
```
>>> assert 'chips' not in items ①
>>> assert 'apples' in items ②
① Make sure "chips" are no longer on the menu.
② Check that we now have some "apples."
```



4.3.9 Joining a List

In our exercise, you'll need to print a string based on the number of elements in the given list. The string will intersperse some other string like ', ' in between all the elements of the list. Oddly, this is the syntax to join a list on the string made of the comma and a space:

```
>>> ', '.join(items)
'soda, sammiches, chips, ice cream'
```



Here we use the `str.join` method and pass the `list` as an argument. It always feels backwards to me, but that's the way it goes. The result of `str.join` is a *new string*:

```
>>> type(', '.join(items))
<class 'str'>
```

The original `list` remains unchanged:

```
>>> items
['soda', 'sammiches', 'chips', 'apples']
```

There is quite a bit more that we can do with Python's `list`, but that should be enough for you to solve this problem.

4.4 Conditional branching with `if/elif/else`

You need to use the conditional branching based on the number of items to correctly format the output. In the "Crow's Nest" exercise, there were two conditions (a "binary" choice) — either a vowel or not — so we used `if/else` statements. Here we have three options to consider, so you will have to use `elif` (`else-if`). For instance, we want to

classify someone by their age by three options:

1. If their age is greater than 0, it is valid.
2. If their age is less than 18, they are a minor.
3. Otherwise they are 18 years or older, which means they can vote:

Here is how I could write that code:

```
>>> age = 15
>>> if age < 0:
...     print('You are impossible.')
... elif age < 18:
...     print('You are a minor.')
... else:
...     print('You can vote.')
...
You are a minor.
```

See if you can use that to figure out how to write the three options for `picnic.py`. That is, first write the branch that handles one item. Then write the branch that handles two items. Then write the last branch for three or more items. Run the tests *after every change to your program*.

Now go write the program yourself before you continue to look at my solution.

Hints:

- Go into your `picnic` directory and run `new.py` `picnic.py` to start your program. Then run `make test` (or `pytest -xv test.py`) and you should pass the first two tests.
- Next work on getting your `--help` usage looking like the above. It's very important to define your arguments correctly. For the `items` argument, look at `nargs` in `argparse` as discussed in chapter 1's "One or more of the same positional arguments" section.
- If you use `new.py` to start your program, be sure to leave the "boolean flag" and modify it for your `sorted` flag.
- Solve the tests in order! First handle one item, then handle two items, then handle three. Then handle the sorted items.

You'll get the best benefit from this book if you try writing the program and passing the tests before reading the solution!

4.5 Solution

```
#!/usr/bin/env python3
"""Picnic game"""

import argparse

# -----
def get_args():
    """Get command-line arguments"""
    ①
```

```

parser = argparse.ArgumentParser(
    description='Picnic game',
    formatter_class=argparse.ArgumentDefaultsHelpFormatter)

parser.add_argument('item',                                ②
                    metavar='str',
                    nargs='+',
                    help='Item(s) to bring')

parser.add_argument('-s',                                 ③
                    '--sorted',
                    action='store_true',
                    help='Sort the items')

return parser.parse_args()                                ④

# -----
def main():                                              ⑤
    """Make a jazz noise here"""

    args = get_args()                                     ⑥
    items = args.item                                    ⑦
    num = len(items)                                    ⑧

    if args.sorted:                                     ⑨
        items.sort()                                    ⑩

    bringing = ''                                       ⑪
    if num == 1:                                         ⑫
        bringing = items[0]                            ⑬
    elif num == 2:
        bringing = ' and '.join(items)                ⑭
    else:
        items[-1] = 'and ' + items[-1]                 ⑮
        bringing = ', '.join(items)                     ⑯

    print('You are bringing {}'.format(bringing)) ⑯

# -----
if __name__ == '__main__':                               ⑳
    main()                                              ㉑

```

- ① The `get_args` function is placed first so that we can easily see what the program accepts when we read it. Note that the function order here is not important to Python, only to us, the reader.
- ② The `item` argument uses the `nargs='+'` so that it will accept *one or more* positional arguments which will be strings.
- ③ The dashes in the short (`-s`) and long (`--sort`) names make this an *option*. There is no value associated with this argument. It's either present (in which case it will be `True`) or absent (or `False`).
- ④ Process the command-line arguments and return them to the caller.
- ⑤ The `main` function is where the program will start.
- ⑥ Call the `get_args` function and put the returned value into the variable `args`. If there is a problem parsing the arguments, the program will fail before the values are returned.
- ⑦ Copy the `item` list from the `args` into the new variable `items`.

- ⑧ Use the length function `len` to get the number of items in the list. There can never be zero arguments because we defined the argument using `nargs='+'` which always requires at least one value.
- ⑨ The `args.sorted` value will either be `True` or `False` because it was defined as a "flag."
- ⑩ If we are supposed to sort the items, call the `items.sort()` method to sort them *in-place*.
- ⑪ Initialize the variable `bringing` with an empty string. We'll put the items we're bringing into this.
- ⑫ If the number of items is 1...
- ⑬ Then we will assign the one item to `bringing`.
- ⑭ If the number of items is 2...
- ⑮ Put the string '`and`' in between the items.
- ⑯ Otherwise...
- ⑰ Alter the last element in `items` to append the string '`and`' before whatever is already there.
- ⑱ Join the `items` on the string '`,` '.
- ⑲ Print the output string, using the `str.format` method to interpolate the `bringing` variable.
- ⑳ When Python runs the program, it will read all the lines to this point but will not run anything. Here we look to see if we are in the "main" namespace.
- ㉑ If we are, call the `main` function to make the program begin.

4.6 Discussion

How did it go? Did it take you long to write your version? How different was it from mine? Let's talk about my solution. It's perfectly fine if yours is really different from mine, just as long as you pass the tests!

4.6.1 Defining the arguments

This program can accept a variable number of arguments which are all the same thing (strings). In my `get_args`, I define an `item` like so:

```
parser.add_argument('item',  
                   metavar='str',  
                   nargs='+',  
                   help='Item(s) to bring') ④
```

- ① A single, required, positional (because no dashes in name) argument called `item`.
- ② An indicator to the user in the usage that this should be a string.
- ③ The number of arguments where '+' means *one or more*.
- ④ A longer help description that appears for the `-h` or `--help` options.

This program also accepts `-s` and `--sorted` arguments. Remember that the leading dashes makes them optional. They are "flags," which typically means that they are `True` if they are present and `False` if absent.

```
parser.add_argument('-s',  
                   '--sorted',  
                   action='store_true',  
                   help='Sort the items') ④
```

- ① The short flag name.
- ② The long flag name.

- ③ If the flag is present, store a True value. The default value will be False.
- ④ The longer help description.

4.6.2 Assigning and sorting the items

To get the arguments, in `main` I call `get_args` and assign them to the `args` variable. Then I create the `items` variable to hold the `args.item` value(s):

```
args = get_args()
items = args.item
```

If `args.sorted` is `True`, then I need to sort my `items`. I chose the *in-place* sort method here:

```
if args.sorted:
    items.sort()
```

Now I have the items, sorted if needed, and I need to format them for the output.

4.6.3 Formatting the items

I suggested you solve the tests in order. There are 4 conditions we need to solve:

1. Zero items
2. One item
3. Two items
4. Three or more items

The first test is actually handled by `argparse` — if the user fails to provide any arguments, they get a usage:

```
$ ./picnic.py
usage: picnic.py [-h] [-s] str [str ...]
picnic.py: error: the following arguments are required: str
```

Since `argparse` handles the case of no arguments, we have to handle the other three conditions. Here's one way to do that:

```
bringing = ''
if num == 1:                                ①
    bringing = items[0]                      ②
elif num == 2:                                ③
    bringing = ' and '.join(items)          ④
else:                                         ⑤
    items[-1] = 'and ' + items[-1]           ⑥
    bringing = ', '.join(items)              ⑦
                                                ⑧
```

- ① Initialize a variable for what we are bringing.
- ② Check if the number of items is one.
- ③ If there is one item, then `bringing` is the one item.
- ④ Check if the number of items is two.
- ⑤ If two items, join the `items` on the string ' `and` '.
- ⑥ Otherwise...

- ⑦ Insert the string 'and ' before the last item
- ⑧ Join all the items on the string ', '.

Can you come up with any other ways?

4.6.4 Printing the items

Finally to print the output, I can use a format string where the {} indicates a placeholder for some value like so:

```
>>> print('You are bringing {}'.format(bringing))
You are bringing salad, soda, and cupcakes.
```

Or, if you prefer, you can use an f''-string:

```
>>> print(f'You are bringing {bringing}.')
You are bringing salad, soda, and cupcakes.
```

They both get the job done, so whichever you prefer.

4.6.5 Testing

For this exercise, I've written the tests for you. Can you see how this is similar to the `hello.py` where we created a `greet` function inside the program and added a `test_greet` function to test it? If I were writing this code for myself, I would do the same by creating a function to format the items and placing the unit test for that inside the program itself.

I believe there is an art to testing. It's up to you to figure out how best to test your own code. Try copying some of the relevant `test_` functions from `test.py` into your `picnic.py` and then running `pytest picnic.py` to see how it works. What do you prefer?

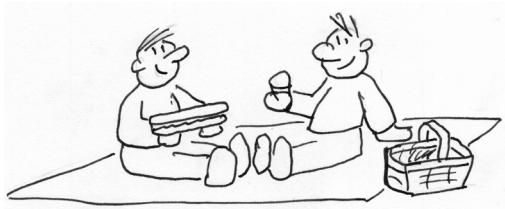
4.7 Summary

- Python lists are ordered sequences of other Python data types such as strings and numbers.
- There are methods like `append` and `extend` to add elements to a list and `pop` and `remove` to remove them.
- You can use `x in y` to ask if element `x` is in the list `y`. You could also use `list.index` to find the index of an element, but this will cause an exception if the element is not present.
- Lists can be sorted and reversed, and elements within lists can be modified. Lists are useful when the order of the elements is important.
- Strings and lists share many features such as using `len` to find their lengths, using zero-based indexing where `0` is the first element and `-1` is the last, and using slices to extract smaller pieces from the whole.
- The `str.join` method can be used to make a new `str` from a `list`.

- `if/elif/else` can be used to branch code depending on conditions.

4.8 Going Further

- Add an option so the user can choose not to print with the Oxford comma (even though that is a morally indefensible option).
- Add an option to separate items with some character passed in by the user (like a semicolon if the list of items needed to contain commas).



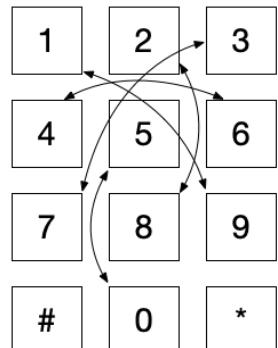
Jump the Five: Working with dictionaries



"When I get up, nothing gets me down."

- D. L. Roth

In an episode of the television show *The Wire*, drug dealers encode telephone numbers that they text in order to obscure them from the police who they assume are intercepting their messages. They use an algorithm we'll call "Jump The Five" where a number is changed to the one that is opposite on a US telephone pad if you jump over the 5. You feel me?



If we start with "1" and jump across the 5, we get to "9," then "6" jumps the 5 to become "4," and so forth. The numbers "5" and "0" will swap with each other. In this exercise, we're going to write a Python program called `jump.py` that will take in some text as a positional argument. Each number in the text will be encoded using this algorithm. All non-number will pass through unchanged, for example:

```
$ ./jump.py 867-5309
243-0751
$ ./jump.py 'Call 1-800-329-8044 today!'
Call 9-255-781-2566 today!
```

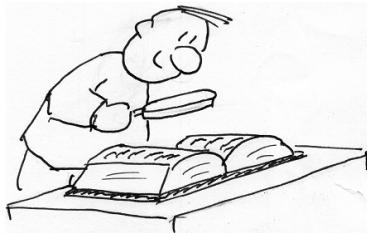
We will need some way to inspect each character in the input text and identify the numbers. We will learn how to use a `for` loop for this and how that relates to a "list comprehension." Then we will need some way to associate a number like "1" with the number "9," and so on for all the numbers. We'll learn about a data structure in Python called a "dictionary" type that allows us to do exactly that.

In this chapter, you will learn to:

- Create a dictionary.
- Use a `for` loop to process text character-by-character.
- Check if items exist in a dictionary.
- Retrieve values from a dictionary.
- Print a new string with the numbers substituted for their encoded values.

Before we get started with the coding, we need to learn about Python's dictionaries.

5.1 Dictionaries



A Python `dict` allows us to relate some *thing* (a "key") to some other *thing* (a "value"). An actual dictionary does this. If we look up a word like "quirky" in a dictionary (www.merriam-webster.com/dictionary/quirky), we can find a definition. We can think of the word itself as the "key" and the definition as the "value."

quirky ➔ unusual, esp. in an interesting or appealing way

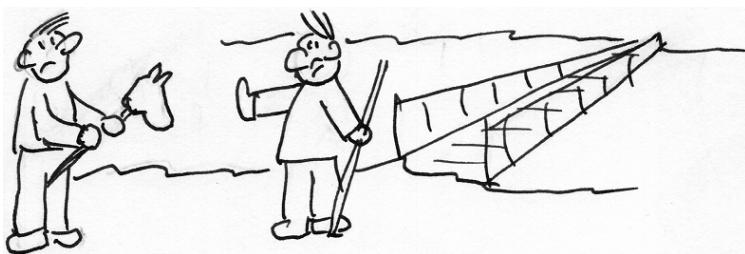
Dictionaries actually provide quite a bit more information such as pronunciation, part of speech, derived words, history, synonyms, alternate spellings, etymology, first known use, etc. (I really love dictionaries.) Each of those attributes has a value, so we could also think of the lookup itself as a dictionary:

definition	➔ unusual, esp. in an interesting or appealing way
pronunciation	➔ 'kwər-kē
part of speech	➔ adjective

Let's see how we can use Python's dictionaries to go beyond word definitions.

5.1.1 Creating a dictionary

In the film *Monty Python and the Holy Grail*, King Arthur and his knights must cross The Bridge of Death. Anyone who wishes to cross must correctly answer three questions from the Keeper. Those who fail are cast into the Gorge of Eternal Peril.



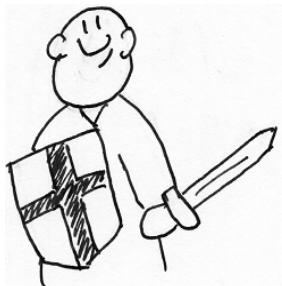
Let us ride to CAMEL...no, sorry, let us create and use a dict to keep track of the questions and answers as key/value pairs. Once again, I want you to fire up your python3/ipython REPL or Jupyter Notebook and type these out for yourself!

Lancelot goes first. We can use the dict() function to create an empty dictionary for his answers.

```
>>> answers = dict()
```

Or we can use empty curly brackets:

```
>>> answers = {}
```



The Keeper's first question: "What is your name?" Lancelot answers "My name is Sir Lancelot of Camelot." We can add the key "name" to our answers by using square brackets ([]) — not curly braces! — and the literal string 'name':

```
>>> answers['name'] = 'Sir Lancelot'
```

If you type `answers<Enter>` in the REPL, Python will show you a structure in curly braces to indicate this is a dict:

```
>>> answers
{'name': 'Sir Lancelot'}
```

{'name': 'Sir Lancelot'}
curly brackets
mean "dictionary"

You can verify with the type function:

```
>>> type(answers)
<class 'dict'>
```

Next the Keeper asks, "What is your quest?" to which Lancelot answers "To seek the

Holy Grail." Let's add "quest" to the answers:

```
>>> answers['quest'] = 'To seek the Holy Grail'
```

There's no return value to let us know something happened, so type `answers` to inspect the variable again to ensure our new key/value was added:

```
>>> answers
{'name': 'Sir Lancelot', 'quest': 'To seek the Holy Grail'}
```

Finally the Keeper asks "What is your favorite color?," and Lancelot answers "blue."

```
>>> answers['favorite_color'] = 'blue'
>>> answers
{'name': 'Sir Lancelot', 'quest': 'To seek the Holy Grail',
 'favorite_color': 'blue'}
```

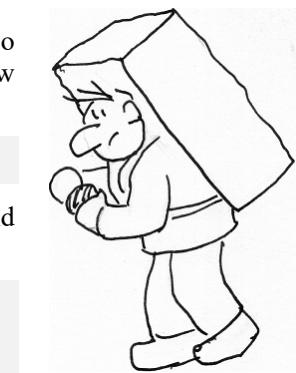
If you knew all the answers beforehand, you could create answers using the `dict()` function with this syntax where you do *not* have to quote the keys and the keys are separate from the values with equal signs:

```
>>> answers = dict(name='Sir Lancelot', quest='To seek the Holy Grail',
favorite_color='blue')
```

Or this syntax using curlies {} where the keys must be quoted and are followed by a colon (:):

```
>>> answers = {'name': 'Sir Lancelot', 'quest': 'To seek the Holy Grail',
'favorite_color': 'blue'}
```

It might be helpful to think of the dictionary `answers` as a box that inside holds the key/value pairs that describe Lancelot's answers just the way the "quirky" dictionary holds all the information about that word.



`answers`

<code>name</code>	⇒ Sir Lancelot
<code>quest</code>	⇒ To seek the Holy Grail
<code>favorite_color</code>	⇒ blue

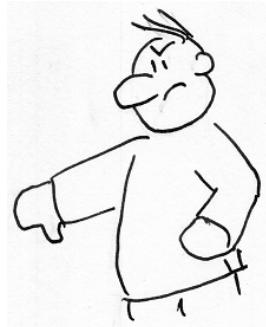
5.1.2 Accessing dictionary values

To retrieve the values, you use the key name inside square brackets ([]). For instance, I can get the name like so:

```
>>> answers['name']
'Sir Lancelot'
```

Let's request his "age":

```
>>> answers['age']
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
KeyError: 'age'
```



WARNING

You will cause an exception if you ask for a dictionary key that doesn't exist!

Just as with lists, you can use the `x in y` to first see if a key exists in the dict:

```
>>> 'quest' in answers
True
>>> 'age' in answers
False
```



The `dict.get` method is a *safe* way to ask for a value:

```
>>> answers.get('quest')
'To seek the Holy Grail'
```

When the requested key does not exist in the dict, it will return the special value `None`:

```
>>> answers.get('age')
```

That doesn't print anything because the REPL won't print a `None`, but we can check the type:

```
>>> type(answers.get('age'))
<class 'NoneType'>
```



There is an optional second argument you can pass to `dict.get` which is the value to return *if the key does not exist*:

```
>>> answers.get('age', 'NA')
'NA'
```

5.1.3 Other dictionary methods

If you want to know how "big" a dictionary is, the `len` (length) function on a `dict` will tell you how many key/value pairs are present:

```
>>> len(answers)
3
```

The `dict.keys` method will give you just the keys:

```
>>> answers.keys()
dict_keys(['name', 'quest', 'favorite_color'])
```

And `dict.values` will give you the values:

```
>>> answers.values()
dict_values(['Sir Lancelot', 'To seek the Holy Grail', 'blue'])
```

Often we want both together, so you might see code like this:

```
>>> for key in answers.keys():
...     print(key, answers[key])
...
name Sir Lancelot
quest To seek the Holy Grail
favorite_color blue
```

An easier way to write this would be to use the `dict.items` method which will return a list of the key/value pairs:

```
>>> answers.items()
dict_items([('name', 'Sir Lancelot'), ('quest', 'To seek the Holy Grail'),
('favorite_color', 'blue')])
```

The above `for` loop could also be written using the `dict.items` method:

```
>>> for key, value in answers.items(): ①
...     print(f'{key:15} {value}')      ②
...
name           Sir Lancelot
quest          To seek the Holy Grail
favorite_color blue
```

- ① For each key/value pair, unpack them into the variables `key` and `value`. Note that you don't have to call them `key` and `value`. You could use `k` and `v` or `question` and `answer`.
- ② Print the `key` in a left-justified field 15 characters wide. The `value` is printed normally.



```
for key, value in [('name', 'Sir Lancelot'), ...]:
```

In the REPL you can execute `help(dict)` to see all the methods available to you like `pop` to remove a key/value or `update` to merge with another `dict`.

Each key in the `dict` is unique. That means if you set a value for a given key twice:

```
>>> answers = {}
>>> answers['favorite_color'] = 'blue'
>>> answers
{'favorite_color': 'blue'}
```

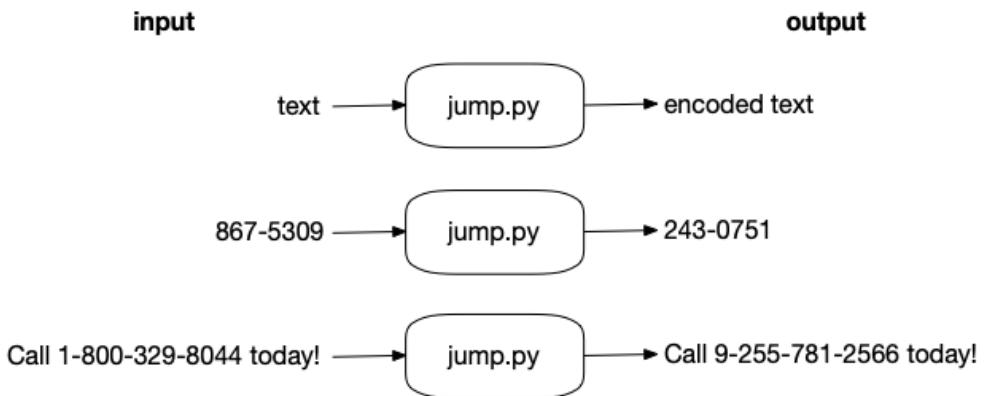
You will not have two entries but one entry with the *second* value:

```
>>> answers['favorite_color'] = 'red'
>>> answers
{'favorite_color': 'red'}
```

Keys don't have to be strings — you can also use numbers like `int` and `float`. Whatever value you use must be immutable. For instance, a `list` could not be a key because it is mutable.

5.2 Writing `jump.py`

Now let's get started with writing our program. Here is a diagram of the inputs and outputs. Note that your program will only affect the numbers in the text. Anything that is *not* a number is unchanged:



When run with no arguments, `-h`, or `--help`, your program should print a usage:

```
$ ./jump.py -h
usage: jump.py [-h] str

Jump the Five

positional arguments:
  str      Input text

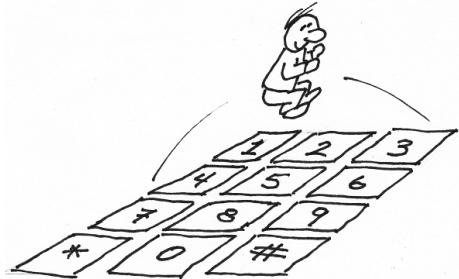
optional arguments:
  -h, --help  show this help message and exit
```

Here is the substitution table for the numbers:

```

1 => 9
2 => 8
3 => 7
4 => 6
5 => 0
6 => 4
7 => 3
8 => 2
9 => 1
0 => 5

```



How would you represent this using a dict? Try creating a dict called `jump` in the REPL and then using a test. Remember that `assert` will return nothing if the statement is True:

```

>>> assert jumper['1'] == '9'
>>> assert jumper['5'] == '0'

```

Next, you will need a way to visit each character in the given text. I suggest you use a `for` loop like so:

```

>>> for char in 'ABC123':
...     print(char)
...
A
B
C
1
2
3

```

Now, rather printing the `char`, print the value of `char` in the `jumper` table or print the `char` itself. Look at the `dict.get` method! Also, if you read `help(print)`, you'll see there is an `end` option to change the newline that gets stuck onto the end to something else.

Hints:

- The numbers can occur anywhere in the text, so I recommend you process the input character-by-character with a `for` loop.
- Given any one character, how can you look it up in your table?
- If the character is in your table, how can you get the value (the translation)?
- If how can you print the translation or the value without printing a newline? Look at `help(print)` in the REPL to read about the options to `print`.
- If you read `help(str)` on Python's `str` class, you'll see that there is a `replace` method. Could you use that?

Now spend the time to write the program on your own before you look at the solutions! Use the tests to guide you.

5.3 Solution

```

#!/usr/bin/env python3
"""Jump the Five"""

import argparse

# -----
def get_args(): ①
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Jump the Five',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('text', metavar='str', help='Input text') ②

    return parser.parse_args()

# -----
def main(): ③
    """Make a jazz noise here"""

    args = get_args() ④
    jumper = {'1': '9', '2': '8', '3': '7', '4': '6', '5': '0',
              '6': '4', '7': '3', '8': '2', '9': '1', '0': '5'} ⑤

    for char in args.text: ⑥
        print(jumper.get(char, char), end='') ⑦
    print() ⑧

# -----
if __name__ == '__main__': ⑨
    main()

```

- ① Define the `get_args` function first so it's easy to see when we read the program.
- ② We define one positional argument called `text`.
- ③ Define a `main` function where the program starts.
- ④ Get the command-line args.
- ⑤ Create a dict for the lookup table.
- ⑥ Process each character in the text.
- ⑦ Print either the value of the char in the `jumper` table or the char if it's not present, making sure *not* to print a newline by adding `end=''`.
- ⑧ Print a newline.
- ⑨ Call the `main()` function if the program is in the "main" namespace.

5.4 Discussion

5.4.1 Defining the arguments

If you look at the solution, you'll see that the `get_args` function is defined first. Our program needs to define one positional argument. Since it's some "text" I expect, I call

the argument 'text' and then assign that to a variable called `text`.

```
parser.add_argument('text', metavar='str', help='Input text')
```

While all that seems rather obvious, I think it's very important to name things *what they are*. That is, please don't leave the name of the argument as 'positional' — that does not describe what it *is*. It may seem like overkill to use `argparse` for such a simple program, but it handles the validation of the correct *number* and *type* of arguments as well as the generation of help documentation, so it's well worth the effort.

5.4.2 Using a dict for encoding

I suggested you could represent the substitution table as a dict where each number key has its substitute as the value in the dict. For instance, I know that if I jump from 1 over the 5 I should land on 9:

```
>>> jumper = {'1': '9', '2': '8', '3': '7', '4': '6', '5': '0',
...             '6': '4', '7': '3', '8': '2', '9': '1', '0': '5'}
>>> jumper['1']
'9'
```

Since there are only 10 numbers to encode, this is probably the easiest way to write this. Note that the numbers are written with quotes around them, so they are actually of the type `str` and not `int` (integers). I do this because I will be reading characters from a `str`. If we stored them as actual numbers, I would have to coerce the `str` types using the `int` function:

```
>>> type('4')
<class 'str'>
>>> type(4)
<class 'int'>
>>> type(int('4'))
<class 'int'>
```

5.4.3 Method 1: Using a for loop to print each character

As suggested in the introduction, I can process each character of the `text` using a `for` loop. To start, I might first see if each character of the `text` is in the `jumper` table using the `x in y` construct.

```
>>> text = 'ABC123'
>>> for char in text:
...     print(char, char in jumper)
...
A False
B False
C False
1 True
2 True
3 True
```

NOTE

When print is given more than one argument, it will put a space in between each of bit of text. You can change that with the sep argument. Read help(print) to learn more.

Now let's try to translate the numbers. I could use an if expression where I print the value from the jumper table if char is present, otherwise print the char:

```
>>> for char in text:
...     print(char, jumper[char] if char in jumper else char)
...
A A
B B
C C
1 9
2 8
3 7
```

It's a bit laborious to check for every character. The dict.get method allows us to safely ask for a value if it is present. For instance, the letter "A" is not in jumper. If we try to retrieve that value, we'll get an exception:

```
>>> jumper['A']
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
KeyError: 'A'
```

But if we use jumper.get, there is no exception:

```
>>> jumper.get('A')
```

When a key doesn't exist in the dictionary, the special None value is returned:

```
>>> for char in text:
...     print(char, jumper.get(char))
...
A None
B None
C None
1 9
2 8
3 7
```

We can provide a second, optional argument to get that is the default value to return when the key does not exist. In our case, if a character does not exist in the jumper, we want to print the character itself. If we had "A," then we'd want to print "A":

```
>>> jumper.get('A', 'A')
'A'
```

But if we have, "5" then we want to print "0":

```
>>> jumper.get('5', '5')
'0'
```

So we can use that to process all the characters:

```
>>> for char in text:
```

```

...     print(jumper.get(char, char))
...
A
B
C
9
8
7

```

I don't want that newline printing after every character, so I can use `end=''` to tell Python to put the empty string at the end instead of a newline. When I run this in the REPL, the output is going to look funny because I have to hit <Enter> after the `for` loop for it to run, then I'll be left with ABC987 with no newline and then the `>>>` prompt:

```

>>> for char in text:
...     print(jumper.get(char, char), end='')
...
ABC987>>>

```

And so in your code you have to add another `print()`. Mostly I wanted to point out a couple things you maybe didn't know about `print`. It's useful that you can change what is added at the end, and that you can `print()` with no arguments to print a newline. There are several other really cool things `print` can do, so I'd encourage you to read `help(print)` and try them out!

5.4.4 Method 2: Using a for loop to build a new string

There are several other ways we could solve this. I don't like all the `print` statements in the first solution, so here's another take:

```

def main():
    args = get_args()
    jumper = {'1': '9', '2': '8', '3': '7', '4': '6', '5': '0',
              '6': '4', '7': '3', '8': '2', '9': '1', '0': '5'}
    new_text = ''                                     ①
    for char in args.text:                           ②
        new_text += jumper.get(char, char)           ③
    print(new_text)                                  ④

```

- ① In this alternate solution, you create an empty `new_text` variable.
- ② Same for loop...
- ③ Append either the encoded number or the original char to the `new_text`
- ④ Print the `new_text`.

While it's fun to explore all the things we can do with `print`, that code is a bit ugly. I think it's cleaner to create a `new_text` variable and call `print` once with that. To do this, we start off by setting a `new_text` equal to the empty string:

```
>>> new_text = ''
```

And we use our same `for` loop to process each character in the `text`. Each time through the loop, we use `+=` to append the right-hand side of the equation to the left-

hand side. The `+=` adds the value on the right to the variable on the left:

```
>>> new_text += 'a'
>>> assert new_text == 'a'
>>> new_text += 'b'
>>> assert new_text == 'ab'
```

On the right, we're using the `jumper.get` method. Each character will be appended to the `new_text`:

```
>>> new_text = ''
>>> for char in text:
...     new_text += jumper.get(char,
char)
... 
```

Now we can call `print` one time with our new value:

```
>>> print(new_text)
ABC987
```

`new_text += jumper.get(char, char)`

the result of `jumper.get`
is appended to `new_text`

5.4.5 Method 3: Using a for loop to build a new List

This method is the same as above, but, rather than `new_text` being a `str`, it's a `list`:

```
def main():
    args = get_args()
    jumper = {'1': '9', '2': '8', '3': '7', '4': '6', '5': '0',
              '6': '4', '7': '3', '8': '2', '9': '1', '0': '5'}
    new_text = []①
    for char in args.text:②
        new_text.append(jumper.get(char, char))③
    print(''.join(new_text))④
```

- ① Initialize `new_text` as an empty list.
- ② Iterate through each character of the text.
- ③ Append the results of the `jumper.get` call to the `new_text` variable.
- ④ Join the `new_text` on the empty string to create a new `str` to print.

As we go through the book, I'll keep reminding you how Python treats strings and lists similarly. It's easy to go back and forth between those two types. Here I'm using `new_text` exactly the same as above, starting off with it empty and then making it longer for each character. We could actually use the exact same `+=` syntax instead of the `list.append` method:

```
for char in args.text:
    new_text += jumper.get(char, char)
```

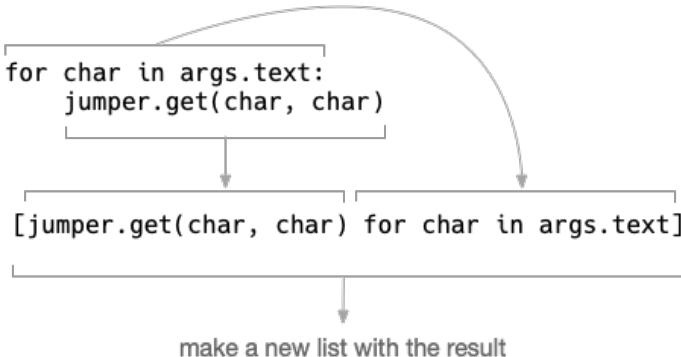
After the `for` loop is done, we have all the new characters that need to be put back together into a new string to print.

5.4.6 Method 4: List comprehension

A shorter solution uses a "list comprehension" which is basically a one-line `for` loop

inside square brackets ([]) which results in a new list:

```
def main():
    args = get_args()
    jumper = {'1': '9', '2': '8', '3': '7', '4': '6', '5': '0',
              '6': '4', '7': '3', '8': '2', '9': '1', '0': '5'}
    print(''.join([jumper.get(char, char) for char in args.text]))
```



A list comprehension is read backwards from a for loop, but it's all there. It's one line of code instead four!

```
>>> text = '867-5309'
>>> [jumper.get(char, char) for char in text]
['2', '4', '3', '-', '0', '7', '5', '1']
```

You can join that new list on the empty string to create a new string you can print:

```
>>> print(''.join([jumper.get(char, char) for char in text]))
243-0751
```

5.4.7 Method 5: str.translate

This last method uses a really powerful method from the str class to change all the characters in one step:

```
def main():
    args = get_args()
    jumper = {'1': '9', '2': '8', '3': '7', '4': '6', '5': '0',
              '6': '4', '7': '3', '8': '2', '9': '1', '0': '5'}
    print(args.text.translate(str.maketrans(jumper)))
```

The argument to str.translate is a translation table that defines how each character should be translated. That's exactly what our jumper does!

```
>>> text = 'Jenny = 867-5309'
>>> text.translate(str.maketrans(jumper))
'Jenny = 243-0751'
```

We'll revisit this function in the "Apples and Bananas" exercise where I'll explain it in greater detail.

5.4.8 (Not) using `str.replace`

Note that you could *not* use `str.replace` to change each number 0-9. Watch how we start off with this string:

```
>>> text = '1234567890'
```

When you change "1" to "9," now you have two 9s:

```
>>> text = text.replace('1', '9')
>>> text
'9234567890'
```

Which means when you try to change all the 9s to 1s, you end up with two 1s:

```
>>> text = text.replace('9', '1')
>>> text
'1234567810'
```

You might try to write it like so:

```
>>> text = '1234567890'
>>> for n in jumper.keys():
...     text = text.replace(n, jumper[n])
...
>>> text
'1234543215'
```

But the correctly encoded string is "9876043215", which is exactly why `str.translate` exists!

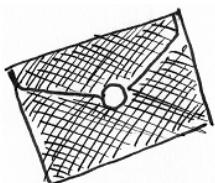
5.5 Summary

- You create a new dictionary using the `dict()` function or with empty curly brackets `({})`.
- Dictionary values are retrieved using their keys inside square brackets or by using the `dict.get` method.
- For a `dict` called `x`, you can use `'key' in x` to determine if a key exists.
- You can use a `for` loop to iterate the characters of a `str` just like you can iterate through the elements of a `list`. You can think of strings as lists of characters.
- The `print` function takes optional keyword arguments like `end=' '` which we can use to print a value to the screen without a newline.

5.6 Going Further

- Try creating a similar program that encodes the numbers with strings, e.g., "5" becomes "five", "7" becomes "seven."
- What happens if you feed the output of the program back into itself. For example, if you run `./jump.py 12345`, you should get `98760`. If you run `./jump.py 98760`, do you recover the original numbers? This is called "round-tripping," and it's a common operation with algorithms that encode and decode text.

Howler: Working with files and `STDOUT`



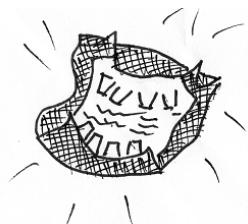
In Harry Potter, a "Howler" is a nasty-gram that arrives by owl at Hogwarts. It will tear itself open, shout a blistering message at the recipient, and then combust. In this exercise, we're going to write a program that will transform text into a rather mild-mannered version of a Howler by MAKING ALL THE LETTERS UPPERCASE. The text that we'll process will be given as a single, positional argument.

For instance, if our program is given the input, "How dare you steal that car!", it should scream back "HOW DARE YOU STEAL THAT CAR!" Remember spaces on the command line delimit arguments, so multiple words need to be enclosed in quotes to be considered one argument:

```
$ ./howler.py 'How dare you steal that car!'
HOW DARE YOU STEAL THAT CAR!
```

The argument to the program may also name a file, in which case we need to read the file for the input:

```
$ ./howler.py ../inputs/fox.txt
THE QUICK BROWN FOX JUMPS OVER THE LAZY DOG.
```



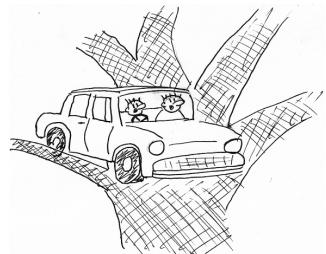
Our program will also accept an `-o` or `--outfile` option that names an output file into which the output text should be written. In that case, *nothing* will be printed on the command line:

```
$ ./howler.py -o out.txt 'How dare you steal that car!'
```

And there should now be a file called `out.txt` that has the output:

```
$ cat out.txt
HOW DARE YOU STEAL THAT CAR!
```

In this exercise, you will learn to:



- Accept text input from the command line or from a file.
- Change strings to uppercase.
- Print output either to the command line or to a file that needs to be created.
- Make plain text behave like a file handle.

6.1 *Reading files*

This will be our first exercise that will involve reading files. The argument to the program will be some text. That text might name an input file in which case you will open and read the file. If it's not the name of a file, then you'll use the text itself.

The built-in `os` (operating system) module has a method for detecting if a string is the name of a file. To use it, you must import the `os` module. For instance, there's probably not a file called "blargh" on your system:

```
>>> import os
>>> os.path.isfile('blargh')
False
```

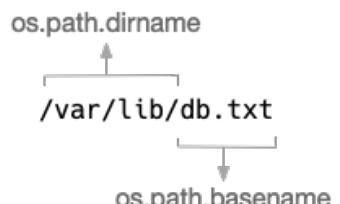


The `os` module contains loads of useful submodules and functions. Consult the documentation at docs.python.org/3/library/os.html or use `help(os)` in the REPL. For instance, the `os.path` module has functions like `basename` and `dirname` for getting a file's name or directory from a path, for example:

```
>>> path = '/var/lib/db.txt'
>>> os.path.dirname(path)
'/var/lib'
>>> os.path.basename(path)
'db.txt'
```

In the `inputs` directory of the source code repo, there are several files. I'll use a file called `inputs/fox.txt`. Note you will need to be in the main directory of the repo for this to work:

```
>>> file = 'inputs/fox.txt'
>>> os.path.isfile(file)
True
```



Once you've determined that the argument is the name of a file, you must open it

to read it. The return from open is a *file handle* and is what we use to read the file. I usually call this variable `fh` to remind me that it's a file handle. If I have more than one open file handle like both input and output handles, I may call them `in_fh` and `out_fh`.

```
>>> fh = open(file)
```

WARNING

If you try to open a file that does not exist, you'll get an exception.

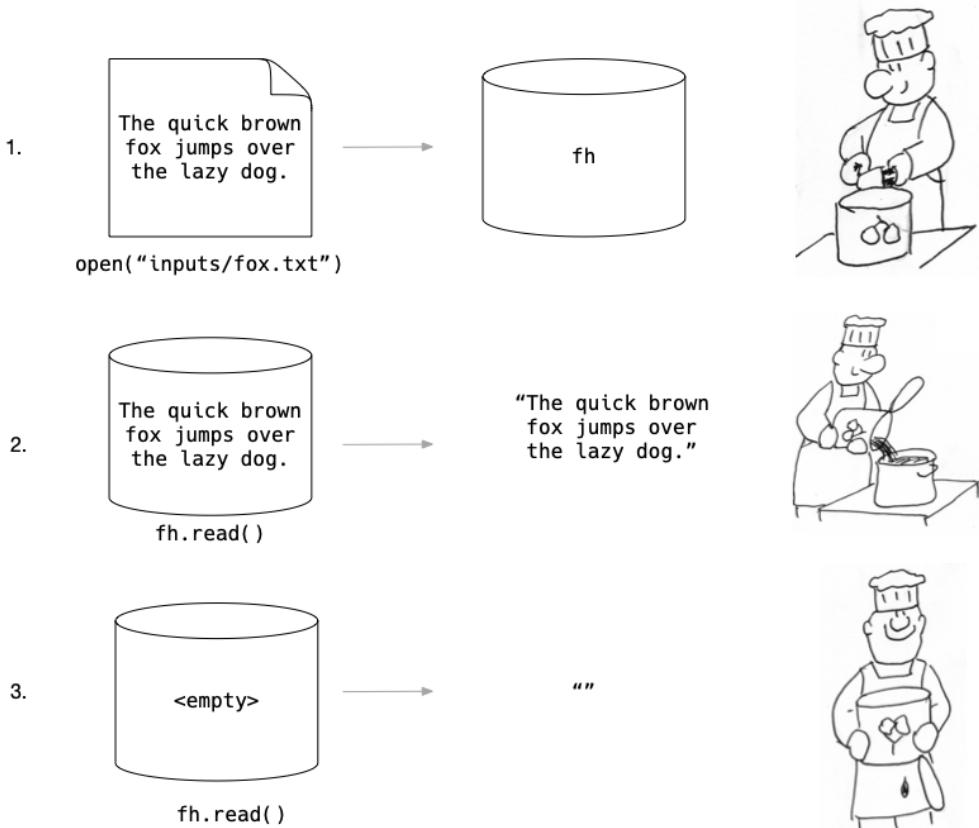
This is unsafe code:

```
>>> file = 'blargh'  
>>> open(file)  
Traceback (most recent call last):  
  File "<stdin>", line 1, in <module>  
FileNotFoundException: [Errno 2] No such file or directory: 'blargh'
```

Always check that the file exists!

```
>>> file = 'inputs/fox.txt'  
>>> if os.path.isfile(file):  
...     fh = open(file)
```

I think of the `file` here ("inputs/fox.txt") as the *name of the file* on disk. It's a bit like a can of tomatoes.



1. The file handle (`fh`) is a mechanism I can use to get at the contents of the file. To get at the tomatoes, we need to open the can.
2. The `fh.read` method returns what is inside the file. With the can opened, we can get at the contents.
3. Once the file handle has been read, there's nothing left.

Let's see what type the `fh` is:

```
>>> type(fh)
<class '_io.TextIOWrapper'>
```

In computer lingo, "io" means "input/output." The `fh` object is something that handles I/O operations. You can use `help(fh)` (using the name of the variable itself) to read the docs on the class `TextIOWrapper`. The two methods you'll use quite often are `read` and `write`. Right now, we care about `read`. Let's see what that gives us:

```
>>> fh.read()
'The quick brown fox jumps over the lazy dog.\n'
```

Do me a favor and execute that line one more time. What do you see?

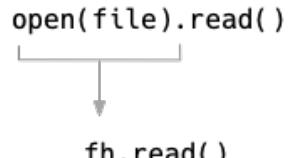
```
>>> fh.read()
''
```



A file handle is different from something like a `str`. Once you `read` a file handle, it's empty! It's like pouring the tomatoes out of the can. Now the can is empty, and you can't empty it again.

We can actually compress the `open` and `read` into one line of code by *chaining* those methods together. The `open` returns a file handle that can be used for the call to `read`. Run this:

```
>>> open(file).read()
'The quick brown fox jumps over the lazy dog.\n'
```



And now run it again:

```
>>> open(file).read()
'The quick brown fox jumps over the lazy dog.\n'
```

Each time you `open` the `file`, you get a fresh file handle to `read`.

If you want to preserve the contents, you need to copy them into a variable.

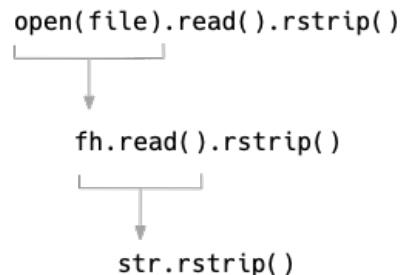
```
>>> text = open(file).read()
>>> text
'The quick brown fox jumps over the lazy dog.\n'
```

The type of the result is a `str`:

```
>>> type(text)
<class 'str'>
```

If I want, I can chain a `str` method onto the end of that. For instance, maybe I want to remove the trailing newline. The `str.rstrip` method will remove any whitespace (which includes newlines) from the *right* end of a string.

```
>>> text = open(file).read().rstrip()
>>> text
'The quick brown fox jumps over the lazy dog.'
```





Once you have your input text — whether it is from the command line or from a file — you need to UPPERCASE it. The `str.upper()` is probably what you want.

6.2 Writing files

The output of the program should either appear on the command line or be written to a file. Command-line output is also called "standard out" or `STDOUT`. (It's the *standard* or normal place for *output* to occur.) There's also an option to the program to write the output to a file, so let's look at how to do that. You still need to open a file handle, but we have to use an optional second argument, the string '`w`', that instructs Python to open it for *writing*. Other modes include '`r`' for *reading* (the default) and '`a`' for *Appending*.

Table 6.1. File writing modes

Symbol	Meaning
<code>w</code>	write
<code>r</code>	read
<code>a</code>	append

You can additionally describe the kind of content, whether '`t`' for *text* (the default) or '`b`' for *binary*:

Table 6.2. File content modes

Symbol	Meaning
<code>t</code>	text
<code>b</code>	bytes

You can combine these two tables like '`rb`' to *read* a *binary* file or '`at`' to *append* to a *text* file. Here we will use '`wt`' to *write* a *text* file. I'll call my variable `out_fh` to remind me that this is the "output file handle":

```
>>> out_fh = open('out.txt', 'wt')
```

If the file does not exist, it will be created. If the file does exist, then it will be *overwritten* which means that all the previous data will be lost! It's possible to open in a mode that will append text to an existing file. For this exercise, '`wt`' will suffice.



You can use the `write` method of the file handle to put text into the file. Whereas

the `print` method will append a newline (`\n`) unless you instruct it not to, the `write` method will *not* add a newline, so you have to explicitly add one.

If you use the `fh.write` method in the REPL, you will see that it return the number of bytes written. Here each character is a byte, and remember that the newline (`\n`) is included:

```
>>> out_fh.write('this is some text\n')
18
```

You can check that this is correct:

```
>>> len('this is some text\n')
18
```

Most code tends to ignore this return value; that is, we don't bother to capture the results into a variable or check that we got a non-zero return. If `write` fails, there's usually some much bigger problem with your system. It's also possible to use the `print` function with the optional `file` argument. Notice that I don't include a newline with `print` because it will add one:

```
>>> print('this is some more text', file=out_fh)
```

When you are done writing to a file handle, you should `close` it so that Python can clean up the file and release the memory associate with it. This function returns no value:

```
>>> out_fh.close()
```

We can verify that our text made it:

```
>>> open('out.txt').read()
'this is some text\nthis is some more text\n'
```

When you `print` on an open file handle, the text will be appended to any previously written data. Look at this code, though:

```
>>> print("I am what I am an' I'm not ashamed", file=open('hagrid.txt', 'wt'))
```

If you run that line twice, will the file called "hagrid.txt" have the line once or twice? Let's find out:

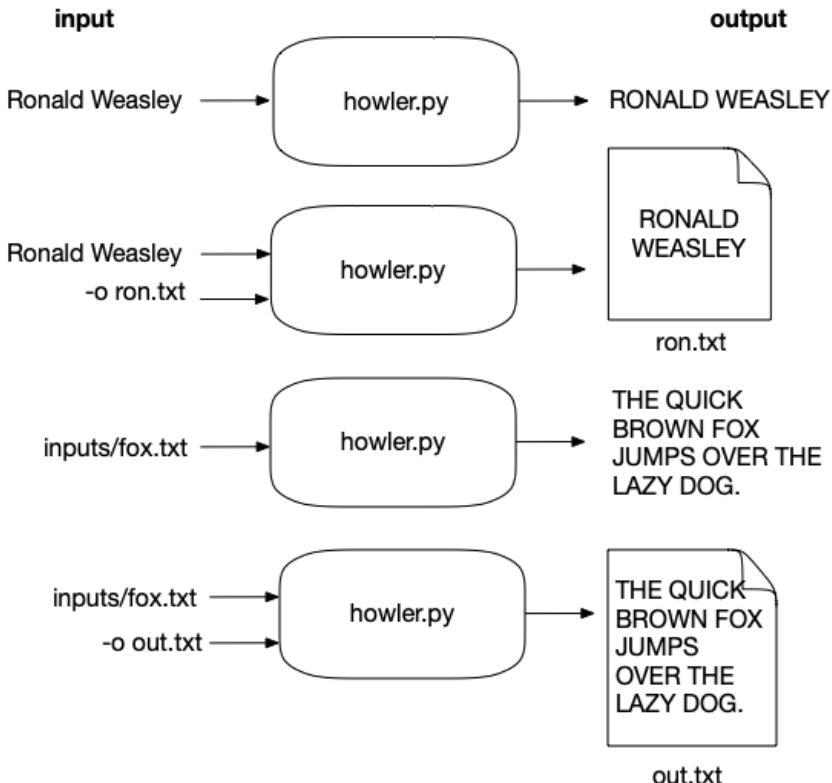
```
>>> open('hagrid.txt').read()
'I am what I am an' I'm not ashamed\n'
```

Just once! Why is that? Remember, each called `open` gives us a new file handle, so calling `open` twice results in new file handles. Both are opened in `write` mode, so that existing data is *overwritten*. It's important to understand how to open files properly or you may end up erasing important data files!

6.3 Writing howler.py

Here is a string diagram showing the overview of the program and some example

inputs and outputs:



When run with no arguments, it should print a short usage:

```
$ ./howler.py
usage: howler.py [-h] [-o str] STR
howler.py: error: the following arguments are required: STR
```

When run with `-h` or `--help`, the program should print a longer usage statement:

```
$ ./howler.py -h
usage: howler.py [-h] [-o str] str

Howler (upper-cases input)

positional arguments:
  str                  Input string or file

optional arguments:
  -h, --help            show this help message and exit
  -o str, --outfile str          Output filename (default: )
```

If the argument is a regular string, it should uppercase that:

```
$ ./howler.py 'How dare you steal that car!'
HOW DARE YOU STEAL THAT CAR!
```

If the argument is the name of a file, it should uppercase the *contents of the file*:

```
$ ./howler.py ../inputs/fox.txt
THE QUICK BROWN FOX JUMPS OVER THE LAZY DOG.
```

If given an --outfile filename, the uppercased text should be written to the indicated file and nothing should be printed to STDOUT:

```
$ ./howler.py -o out.txt ../inputs/fox.txt
$ cat out.txt
THE QUICK BROWN FOX JUMPS OVER THE LAZY DOG.
```

Hints:

- Start with new.py and alter the get_args section until your usage statements match the ones above.
- Run the test suite and try to pass just the first test that handles text on the command line and output to STDOUT.
- The next test is to see if you can write the output to a given file. Figure out how to do that.
- The next test is for reading input from a file. Then work on that. Don't try to do all the things at once!
- There is a special file handle that always exists called "standard out" (often STDOUT). If you print without a file argument, then it defaults to sys.stdout. You will need to import sys in order to use it.

Be sure you really try to write the program and pass all the tests before moving on to read the solution! If you get stuck, maybe whip up a batch of Polyjuice Potion and freak out your friends.

6.4 Solution

```
#!/usr/bin/env python3
"""
Howler"""

import argparse
import os
import sys

# -----
def get_args():
    """
    get command-line arguments
    """

    parser = argparse.ArgumentParser(
        description='Howler (upper-case input)',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('text', metavar='str', help='Input string or file') ①

    parser.add_argument('-o', ②
```

```

        '--outfile',
        help='Output filename',
        metavar='str',
        type=str,
        default='')

args = parser.parse_args()                                (3)

if os.path.isfile(args.text):
    args.text = open(args.text).read().rstrip()          (4)
    (5)

return args                                              (6)

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()                                    (7)
    out_fh = open(args.outfile, 'wt') if args.outfile else sys.stdout
    out_fh.write(args.text.upper() + '\n')                (8)
    out_fh.close()                                       (9)
    (10)

# -----
if __name__ == '__main__':
    main()

```

- ① The text argument is a string that may be the name of a file.
- ② The --outfile option is also a string that names a file.
- ③ Parse the command-line arguments into the variable args so that we can manually check the text argument.
- ④ Check if args.text names an existing file.
- ⑤ If so, overwrite the value of args.text with the results of reading the file.
- ⑥ Now that we've fixed up the args, we can return them to the caller.
- ⑦ Call get_args to get the arguments to the program.
- ⑧ Use an if expression to choose either sys.stdout or a newly opened file handle to write the output.
- ⑨ Use the opened file handle to write the output converted to upper.
- ⑩ Close the file handle.

6.5 Discussion

How did it go for you this time? I hope you avoided Snape's office. You really don't want more Saturday detentions.

6.5.1 Defining the arguments

The get_args function, as always, is the first. Here I define two arguments. The first is a positional text argument. Since it may or may not name a file, all I can know is that it will be a string.

```
parser.add_argument('text', metavar='str', help='Input string or file')
```

The other argument is an option, so I give it a short name of -o and a long name of --

`outfile`. The default type for all arguments is `str`, but I go ahead and state that explicitly. The default value is the empty string. I could just as easily use the special `None` type which is also the default value, but I prefer to use a defined argument like the empty string.

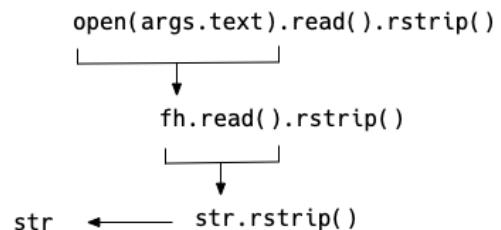
```
parser.add_argument('-o',
                   '--outfile',
                   help='Output filename',
                   metavar='str',
                   type=str,
                   default='')
```

6.5.2 Reading input from a file or the command line

This is a deceptively simple program that demonstrates a couple of very important elements of file input and output. The `text` input might be a plain string or it might be the name of a file. This pattern will come up repeatedly in this book:

```
if os.path.isfile(args.text):
    args.text = open(args.text).read().rstrip()
```

The function `os.path.isfile` will tell us if there is a file with the name in `text`. If that returns `True`, then we can safely `open(file)` to get a *file handle* which has a method called `read` which will return *all* the contents of the file. This is usually safe, but be careful if you write a program that could potentially read gigantic files. For instance, in my day job we regularly deal with files with sizes in the 10s to 100s of gigabytes, so I would need to ensure my system has more memory than the size of the file!



The result of `open(file).read()` is a `str` which itself has a method called `rstrip` that will return a copy of the string *stripped* of the whitespace off the *right* side of the string (hence the name `rstrip`). The longer way to write the above would be:

```
if os.path.isfile(text):
    fh = open(text)
    text = fh.read()
    text = text.rstrip()
```

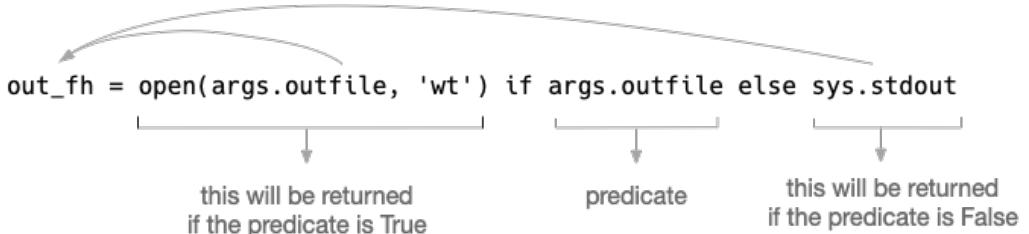
In my version, I choose to handle this inside the `get_args` function. This is the first time I'm showing you that you can intercept and alter the arguments before passing them on to `main`. We'll use this idea quite a bit in more exercises. I like to do all the work to validate the user's arguments inside `get_args`. I could just as easily do this in `main` after the call to `get_args`, so this is entirely a style issue.

6.5.3 Choosing the output file handle

The line decides where to put the output of our program

```
out_fh = open(args.outfile, 'wt') if args.outfile else sys.stdout
```

The `if` expression will open `args.outfile` for writing text (`'wt'`) if the user provided that argument; otherwise, we will use `sys.stdout` which is a file handle to STDOUT. Note that we don't have to call `open` on `sys.stdout` because it is always there and always open for business.



6.5.4 Printing the output

To get uppercase, we can use the `text.upper` method, then we need to find a way to print it to the output file handle. I chose to do:

```
out_fh.write(text.upper())
```

But you could also do:

```
print(text.upper(), file=out_fh)
```

Finally I need to close the file handle with `out_fh.close()`.

6.5.5 A low-memory version

There is a potentially serious problem waiting to bite us in this program. In the `get_args`, we're reading the entire file into memory with this line:

```
if os.path.isfile(args.text):
    args.text = open(args.text).read().rstrip()
```

We could, instead, only open the file:

```
if os.path.isfile(args.text):
    args.text = open(args.text)
```

And later read it line-by-line:

```
for line in args.text:
    out_fh.write(line.upper())
```

The problem, though, is how to handle the times when the `text` argument is actually text and not the name of a file. The `io` (input-output) module in Python has a way to

represent text as *stream*:

```
>>> import io
>>> text = io.StringIO('foo\nbar\nbaz\n')
>>> for line in text:
...     print(line, end='')
...
foo
bar
baz
```

- ① Import the `io` module.
- ② Use the `io.StringIO` function to turn the given `str` value into something we can treat like an open file handle.
- ③ Use a `for` loop to iterate the "lines" of text by separated by newlines.
- ④ Print the line using the `end=''` option to avoid having 2 newlines.

To make this work, we can change how we handle the `args.text` like so:

```
#!/usr/bin/env python3
"""Low-memory Howler"""

import argparse
import os
import io
import sys

# -----
def get_args():
    """get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Howler (upper-cases input)',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('text', metavar='str', help='Input string or file')

    parser.add_argument('-o',
                        '--outfile',
                        help='Output filename',
                        metavar='str',
                        type=str,
                        default='')

    args = parser.parse_args()

    if os.path.isfile(args.text):
        args.text = open(args.text)          ①
    else:
        args.text = io.StringIO(args.text + '\n') ②

    return args

# -----
def main():
    """Make a jazz noise here"""


# -----
```

```

args = get_args()
out_fh = open(args.outfile, 'wt') if args.outfile else sys.stdout
for line in args.text:
    out_fh.write(line.upper())           ④
    out_fh.close()                     ⑤

# -----
if __name__ == '__main__':
    main()

```

- ① Check if `args.text` is a file.
- ② If so, replace `args.text` with the file handle created by opening the file named by it.
- ③ Otherwise, replace `args.text` with an `io.StringIO` value that will act like an open file handle. Note that we need to add a newline to the text so that it will look like the lines of input coming from an actual file.
- ④ Read the input (whether `io.StringIO` or a file handle) line-by-line.
- ⑤ Handle the line.

6.6 Review

- To read or write to files, you must open them.
- The default mode for `open` is for reading a file.
- To write a text file, you must use '`wt`' as the second argument to `open`.
- By default, you write text to a file handle. You must use the '`b`' flag to indicate that you want to write binary data.
- The `os.path` module contains many useful functions such as `isfile` that will tell you if a file exists by the given name.
- `STDOUT` (standard output) is always available via the special `sys.stdout` file handle which is always open.
- The `print` function takes an optional `file` argument of where to put the output. That argument must be an open file handle such as `sys.stdout` (the default) or the result of `open`.



6.7 Going Further

- Add a flag that will lowercase the input instead. Maybe call it `--ee` for the poet e e cummings who liked to write poetry devoid of uppercase letters.
- Alter the program handle multiple input files
- Change `--outfile` to `--outdir` and write each input file to the same file name in the output directory.

Words Count: Reading files/STDIN, iterating lists, formatting strings

"I love to count!" — Count von Count

Counting things is a surprisingly important programming skill. Maybe you're trying to find how many pizzas were sold each quarter or how many times you see certain words in a set of documents. Usually the data we deal with in computing comes to us in files, so we're going to push a little further into reading files and manipulating strings by writing a Python version of the venerable Unix `wc` (word count) program. We're going to write a program called `wc.py` that will count the characters, words, and lines for all the files given as positional arguments.



Given one or more valid files, it should print the number of lines, words, and characters, each in columns 8 characters wide, followed by a space and then the name of the file. Here's what it looks like for one file:

```
$ ./wc.py ../inputs/scarlet.txt
 7035   68061  396320 ../inputs/scarlet.txt
```

When there are many files, print the counts for each file and then print a "total" line summing each column:

```
$ ./wc.py ../inputs/s*.txt
    7035   68061  396320 ../inputs/scarlet.txt
  17 118 661 ../inputs/sonnet-29.txt
    3   7   45 ../inputs/spiders.txt
  7055   68186  397026 total
```

There may also be *no* arguments, in which case we'll read from "standard in" (STDIN) which is the complement to STDOUT we used in "Howler." You can use STDIN to chain programs together where the output of one program becomes the input for the next. To pass text via STDIN, you can use the < redirect from a file:

```
$ ./wc.py < ../inputs/fox.txt
  1   9   45 <stdin>
```

Or pipe (|) the output of one command into another:

```
$ cat ../inputs/fox.txt | ./wc.py
  1   9   45 <stdin>
```

In this exercise, you will:

- Learn how to process zero or more positional arguments
- Validate input files
- Read from files or from "standard in"
- Use multiple levels of `for` loops
- Break files into lines, words, and characters
- Use counter variables
- Format string output

7.1 Writing `wc.py`

Let's get started! Create your program and modify the arguments until it will print the following usage if run with the -h or --help flags:

```
$ ./wc.py -h
usage: wc.py [-h] [FILE [FILE ...]]

Emulate wc (word count)

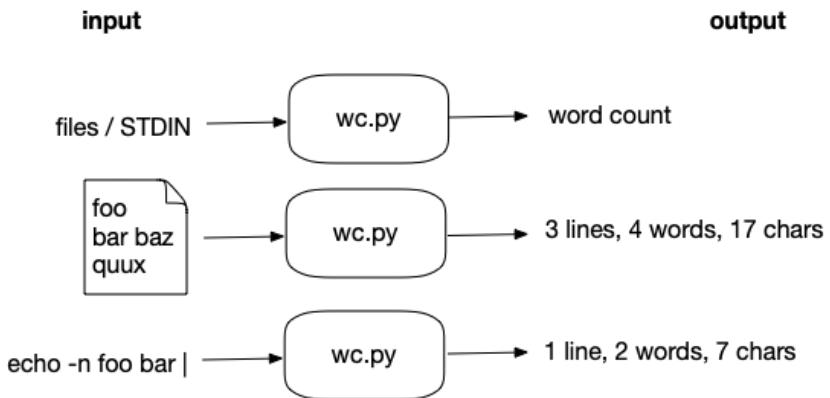
positional arguments:
  FILE      Input file(s) (default: [<io.TextIOWrapper name='<stdin>'  
          mode='r' encoding='UTF-8']])

optional arguments:
  -h, --help  show this help message and exit
```

Given a non-existent file, your program should print an error message and exit with a non-zero exit value:

```
$ ./wc.py foo
usage: wc.py [-h] FILE [FILE ...]
wc.py: error: argument FILE: can't open 'foo': \
[Errno 2] No such file or directory: 'foo'
```

Here is a string diagram to help you think about how the program should work:



7.1.1 Defining file inputs

The first step will be to define your arguments to `argparse`. The program takes *zero or more* positional arguments and nothing else. Remember that you never have to define the `-h` or `--help` arguments as `argparse` handles those automatically.

In "Picnic," we used `nargs='+'` to indicate one or more items for our picnic. Here we want to use `nargs='*'` to indicate *zero or more*. For what it's worth, there's one other value that `nargs` can take and that is `?` for *zero or one*. In all cases, the argument(s) will be returned as a list. Even if there are no arguments, you will still get an empty list (`[]`). For this program, if there are no arguments, we'll read `STDIN`.

Table 7.1. Possible values for `nargs`

Symbol	Meaning
<code>?</code>	<code>zero or one</code>
<code>*</code>	<code>zero or more</code>
<code>+</code>	<code>one or more</code>

In "Howler," we used the "standard out" (`STDOUT`) file handle with `sys.stdout`. To read `STDIN`, we'll use Python's `sys.stdin` file handle which is similarly always open and available to you. Because you are using `nargs='*'`, the values will be a list, and so the default should be a list as well. Can you figure out how to make the default value for your file argument be a list with `sys.stdin`?

Lastly, we should discuss the type of the positional arguments. If they are provided, they should be *readable files*. We saw in "Howler" how to test if the input argument was a file by using `os.path.isfile`. In that program, the input might be either plain text or a file name, so we had to check this ourselves. In this program, however, we will require that any arguments should be files, and so we can define the `type=argparse.FileType('r')`. When you do this, `argparse` takes on all the work to

validate the inputs from the user and produce useful error messages. Be sure to review the "File arguments" section from Chapter 2. If the user provides valid input, then argparse will provide you with a list of *open file handles*. All in all, this saves you quite a bit of time.

7.1.2 Iterating lists

Your program will end up with a list of file handles. In "Jump The Five," we used a for loop to iterate through the characters in the input text. Here we can use a for loop over the file inputs.

```
for fh in args.file:  
    # read each file
```

The `fh` is a "file handle." We saw in "Howler" how to manually open and read a file. Here the `fh` is already open, so we can read the contents from it. There are many ways to read a file, however. The `read` method will give you the *entire contents* of the file in one go. If the file is large — say, if the size of the file exceeds your available memory on your machine — then your program will crash. I would recommend, instead, that you use a for loop on the `fh`. Python will understand this to mean that you wish to read each line of input, one-at-a-time.

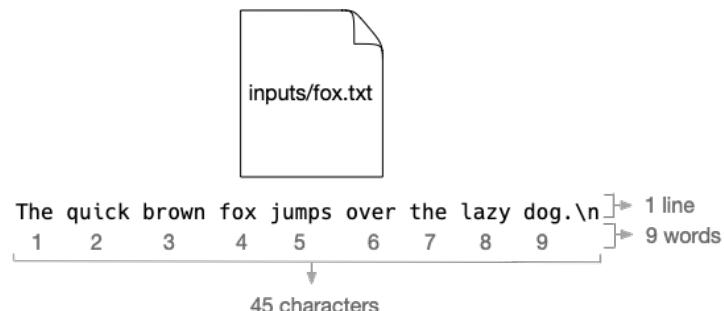
```
for fh in args.file: # ONE LOOP!  
    for line in fh: # TWO LOOPS!  
        # process the line
```

So that's two levels of for loops, one for each file handle and then another for each line in each file handle. TWO LOOPS! I LOVE TO COUNT!

7.1.3 What you're counting

The output for each file will be the number of lines, words, and characters, each printed in a field 8 characters wide followed by the name of the file which will be available to you via `fh.name`. Let's take a look at the output from the standard `wc` program on my system. Notice that when run with just one argument, it produces counts only for that file:

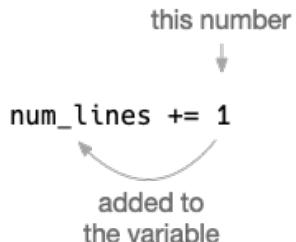
```
$ wc fox.txt  
1 9 45 fox.txt
```



When run with multiple files, it also shows a "total" line:

```
$ wc fox.txt sonnet-29.txt
 1   9  45 fox.txt
17 118 669 sonnet-29.txt
18 127 714 total
```

For each file, you will need to create variables that hold the numbers for lines, words, and characters. For instance, if you use the `for line in fh` loop that I suggest, then you need to have a variable like `num_lines` to increment on each iteration. That is, somewhere in your code you will need to set a variable to `0` and then, inside the `for` loop, make it go up by one. The idiom in Python for this is:



```
num_lines = 0
for line in fh:
    num_lines += 1
```

You also need to count the number of words and characters, so you'll need similar `num_words` and `num_chars` variables. In "Picnic," we discussed how we can convert back and forth between strings and lists. For the purposes of this exercise, we'll use the `str.split` method to break each line on spaces. You can then use the length of the resulting list as the number of words. For the number of characters, you can use the same length function on the line and add that to a `num_chars` variable.

7.1.4 Formatting your results

This is the first exercise where the output needs to be formatted in a particular way. Don't try handle this part manually. That way lies madness. Instead, you need to learn the magic of the `str.format` method. The `help` doesn't have much in the way of documentation, so I'd recommend you read PEP3101 (www.python.org/dev/peps/pep-3101/).

We've seen that the curly braces `({})` inside the `str` part create placeholders that will be replaced by the values passed to the method:

```
>>> import math
>>> 'Pi is {}'.format(math.pi)
'Pi is 3.141592653589793'
```

You can put formatting information inside the curly braces to specify how you want the value displayed. If you are familiar with `printf` from C-type languages, this is the same idea. For instance, I can print just two numbers of pi after the decimal. The `:` introduces the formatting options, and the `0.02f` describes two decimal points of precision:

```
>>> 'Pi is {:.02f}'.format(math.pi)
'Pi is 3.14'
```

The formatting information comes after the colon (:) inside the curly braces. You can also use the f-string method where the variable comes *before* the colon:

```
>>> f'Pi is {math.pi:0.02f}'
'Pi is 3.14'
```

Here you need to use `{:8}` for each of lines, words, and characters so that they all line up in neat columns. The 8 describes the width of the field which is assumed to be a string. The text will be right-justified.

Hints:

- Start with `new.py` and delete all the non-positional arguments.
- Use `nargs='*' to indicate zero or more positional arguments for your file argument.`
- How could you use `sys.stdin` for the `default`? Remember that both `narg='*' and nargs='+' mean that the arguments will be supplied as a list. How can you create a list that contains just sys.stdin for the default value?`
- Remember that you are just trying to pass one test at a time. Create the program, get the help right, then worry about the first test.
- Compare the results of your version to the `wc` installed on your system. Note that not every Unix-like system has the same `wc`, so results may vary.

Time to write this yourself before you read the solution. Fear is the mind-killer. You can do this.

```
<<<
```

7.2 Solution

```
#!/usr/bin/env python3
"""Emulate wc (word count)"""

import argparse
import sys

# -----
def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Emulate wc (word count)',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('file',
                        metavar='FILE',
                        nargs='*',
                        default=[sys.stdin],
                        type=argparse.FileType('r'),
                        help='Input file(s)')①②
```

```

    return parser.parse_args()

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()

    total_lines, total_chars, total_words = 0, 0, 0
    for fh in args.file:
        lines, words, chars = 0, 0, 0
        for line in fh:
            lines += 1
            chars += len(line)
            words += len(line.split())

        total_lines += lines
        total_chars += chars
        total_words += words

        print(f'{lines:8}{words:8}{chars:8} {fh.name}') ⑪

    if len(args.file) > 1:
        print(f'{total_lines:8}{total_words:8}{total_chars:8} total') ⑫ ⑬

# -----
if __name__ == '__main__':
    main()

```

- ① If you set the default to a list with `sys.stdin`, then you have handled the STDIN option.
- ② If the user supplies any arguments, argparse will check if they are valid file inputs. If there is a problem, argparse will halt execution of the program and show the user an error message.
- ③ These are the variables for the "total" line, if we need them.
- ④ Iterate through the list of `arg.file` inputs. I use the variable `fh` to remind me that these are open file handles, even STDIN.
- ⑤ Initialize variables to count *just this file*.
- ⑥ Iterate through each line of `fh`.
- ⑦ For each line, we increment `lines` by 1.
- ⑧ The number of `chars` is incremented by the length of the line.
- ⑨ To get the number of words, we can split the line on spaces (the default). We length of that list is added to the `words`.
- ⑩ We add the numbers for this file to the `total_` variables.
- ⑪ Print the counts for this file using the `{:8}` option to print in a field 8 characters wide.
- ⑫ Check if we had more than 1 input.
- ⑬ Print the "total" line.

7.3 Discussion

7.3.1 Defining the arguments

This program is rather short and seems rather simple, but it's definitely not exactly

easy. One part of the exercise is to really get familiar with argparse and the trouble it can save you. The key is in defining the file positional arguments. If you use nargs='*' to indicate zero or more arguments, then you know argparse is going to give you back a list with zero or more elements. If you use type=argparse.FileType('r'), then any arguments provided must be readable files. The list that argparse returns will be a list of *open file handles*. Lastly, if you use default=[sys.stdin], then you understand that sys.stdin is essentially an open file handle to read from "standard in" (AKA STDIN), and you are letting argparse know that you want the default to be a list containing sys.stdin.

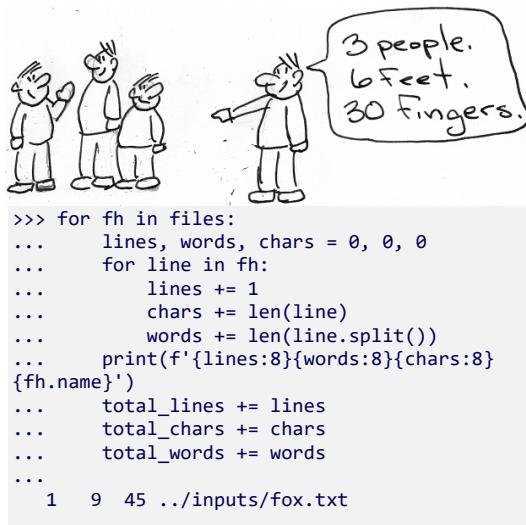
7.3.2 Reading a file using a for loop

I can create a list of open file handles in the REPL to mimic what I'd get from args.file:

```
>>> files = [open('../inputs/fox.txt')]
```

Before I use a for loop to iterate through them, I need to set up three variables to track the *total* number of lines, words, and characters:

```
>>> total_lines, total_chars, total_words = 0, 0, 0
```



Inside the for loop for each file handle, I initialize three more variables to hold the count of lines, characters, and words *for this particular file*. I then use another for loop to iterate over each line in the file handle (fh). For the lines, I can add 1 on each pass through the for loop. For the chars, I can add length of the line (len(line)) to track the number of characters. Lastly for the words, I can use line.split() to break the line on whitespace to create a list of "words." It's not actually a perfect way to count actual words, but it's close enough. I can use the len function on the list to add to the words variable. The for loop ends when the end of the file is reached, and that is when I can print out the counts and the file name using {:8} placeholders in the print template to indicate a text field 8 characters wide.

Lastly, if the number of file arguments is greater than 1, I need to print my totals:

```
if len(args.file) > 1:
    print(f'{total_lines:8}{total_words:8}{total_chars:8} total')
```

7.4 Review

- The `nargs` (number of arguments) option to `argparse` allows you to validate the number of arguments from the user. The star ('*') means zero or more while '+' means one or more.
- If you define an argument using `type=argparse.FileType('r')`, then `argparse` will validate that the user has provided a readable file and will make the value available in your code as an open file handle.
- You can read and write from the Unix standard in/out file handles by using `sys.stdin` and `sys.stdout`.
- You can nest `for` loops to handle multiple levels of processing.
- The `str.split` method will split a string on spaces into words.
- The `len` function can be used on both strings and lists. For the latter, it will tell you the number of elements contained.
- The `str.format` and Python's f-strings both recognize the same printf-style formatting options to allow you to control how a value is displayed.

7.5 Going Further

- By default, `wc` will print all the columns like our program, but it will also accept flags to print `-c` for number of characters, `-l` for number of lines, and `-w` for number of words. When any of these flags are present, only those columns for the given flags are shown, so `-wc` would show just the columns for characters and words. Add both short and long flags for these options to your program so that it behaves exactly like `wc`.
- Implement other system tools like `head`, `tail`, `cat`, and `tac` (the reverse of `cat`).

Gashlycrumb: Looking items up in a dictionary

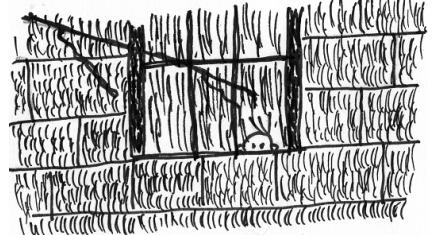
Every time you log into a website, the code behind it has to look up your username and password to compare to the values you put into the login form. Whenever you give your phone number at the hardware store or scan your library card to checkout a book, a computer program uses one piece of information to find other things like how often you buy compost or if you have any overdue books. Probably all these examples would be using a database to find that information. We're going to use a dictionary that we will fill with information from an input file.

In this exercise, we're going to look up a lines of text from an input file that start with the letters provided by the user. The text will come from an input file which will default to Edward Gorey's "The Gashlycrumb Tinies," an abecedarian book that describes various and ghastly ways in which children die. Instead of "A is for artichoke, B is for blackberry," we get:

A is for Amy who fell down the stairs. B is for Basil assaulted by bears.

Our `gashlycrumb.py` program will take one or more letters of the alphabet as positional

Figure 8.1. N is for Neville who died of ennui.



arguments and will look up the lines of text from an *optional* input file that start with that letter. The input file will have each letter on a separate line:

```
$ head -2 gashlycrumb.txt
A is for Amy who fell down the stairs.
B is for Basil assaulted by bears.
```

When our unfortunate user runs our program, here is what they would see. Note that we will consider the letter in a *case-insensitive* fashion:

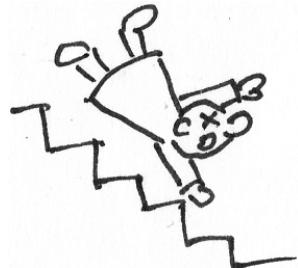
```
$ ./gashlycrumb.py e f
E is for Ernest who choked on a peach.
F is for Fanny sucked dry by a leech.
```

In this exercise, you will:

- Accept one or more positional arguments we'll call `letter`
- Accept an optional `--file` argument which must be a file. The default value will be '`gashlycrumb.txt`' (provided).
- Read the file, find the first letter of each line, and build a data structure that associates that letter to the line of text.
- For each `letter` provided by the user, either print the line of text for the `letter` if present or a message if it isn't.
- Learn how to "pretty print" a data structure.

You can draw from several previous programs:

- From the "Word Count" program, you know how to take a file input and read it line-by-line.
- From the "Article" program, you know how to get the first letter of a bit of text.
- From the "Jump The Five" program, you know how to build a dictionary and lookup a value.



Now you'll put all those skills together to recite morbid poetry!

8.1 Writing `gashlycrumb.py`

Before you begin writing, I would encourage you to run the tests with `make test` or `pytest -xv test.py` in the `gashlycrumb` directory. The first test should fail:

```
test.py::test_exists FAILED
```

This is just a reminder that the first thing to do is to create the file called `gashlycrumb.py`. You can do this however you like, maybe by running `new.py gashlycrumb.py` in the `gashlycrumb` directory or by copying the `template/template.py` file, or by just starting a new file from scratch.

Run your tests again and you should pass the first and possibly the second tests if your program produces a usage statement. Next let's get the arguments straight. Modify

your program's parameters in the `get_args` function so that it will produce the following usage when the program is run with *no arguments* or with the `-h` or `--help` flags:

```
$ ./gashlycrumb.py -h
usage: gashlycrumb.py [-h] [-f str] str [str ...]

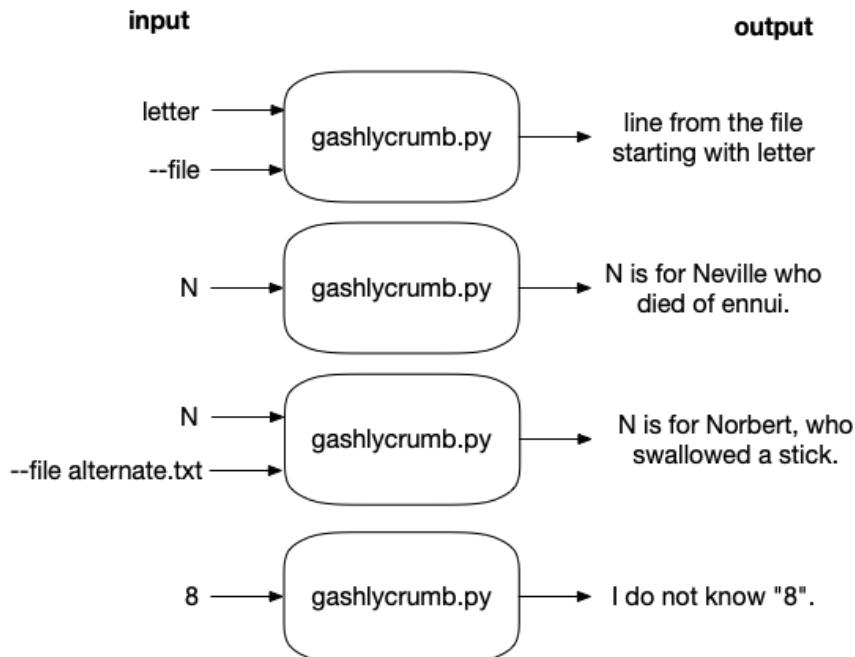
Gashlycrumb

positional arguments:
  str                  Letter(s)          ①

optional arguments:
  -h, --help            show this help message and exit      ②
  -f str, --file str   Input file (default: gashlycrumb.txt) ③
```

- ① The letter is a required positional argument that accepts one or more values (notice `[str ...]` to indicate this).
- ② The `-h` and `--help` arguments are created automatically by argparse.
- ③ The `-f` or `--file` argument is an option with a default value of `gashlycrumb.txt`.

Here is a string diagram showing how the program will work:



Once you have the correct usage, start off by echoing each of the letter arguments:

```
def main():
    args = get_args()
    for letter in args.letter:
        print(letter)
```

Try running it to make sure it works:

```
$ ./gashlycrumb.py a b
a
b
```

Next read the file line-by-line using a for loop. Note that I'm using `end=''` with the print so that it won't print the newline that's already attached to each line of the file:

```
def main():
    args = get_args()
    for letter in args.letter:
        print(letter)

    for line in args.file:
        print(line, end='')
```

Try running it to ensure you can read the input file:

```
$ ./gashlycrumb.py a b | head -4
a
b
A is for Amy who fell down the stairs.
B is for Basil assaulted by bears.
```

Use the "alternate.txt" file, too:

```
$ ./gashlycrumb.py a b --file alternate.txt | head -4
a
b
A is for Alfred, poisoned to death.
B is for Bertrand, consumed by meth.
```

If provided a `--file` argument that does not exist, your program should exit with an error message. Note that, if you use `type=argparse.FileType('r')`, this error should be produced automatically by `argparse`:

```
$ ./gashlycrumb.py -f blaragh b
usage: gashlycrumb.py [-h] [-f str] str
gashlycrumb.py: error: argument -f/--file: can't open 'blaragh': \
[Errno 2] No such file or directory: 'blaragh'
```

Now think about how you can use the first letter of each line to create an entry into a dict. Use the `print` command to look at your dictionary. Figure out how to check if the given letter is in (wink, wink, nudge, nudge) your dictionary. If given a value that does not exist in the list of first characters on the lines from the input file (when searched without regard to case), you should print a message:

```
$ ./gashlycrumb.py 3
I do not know "3".
$ ./gashlycrumb.py CH
I do not know "CH".
```

If the given letter is in the dictionary, print the



value for it:

```
$ ./gashlycrumb.py a
A is for Amy who fell down the stairs.
$ ./gashlycrumb.py z
Z is for Zillah who drank too much gin.
```



Run the test suite to ensure your program meets all the requirements. Read the errors closely and fix your program.

Hints:

- Start with `new.py` and remove everything but the positional and optional `argparse.FileType('r')` parameters.
- Use `nargs='+'` to define the positional argument `letter` so it will require one or more values.
- A dictionary is a natural data structure that you can use to associate some value like the letter "A" to some phrase like "A is for Amy who fell down the stairs." Create a new, empty dict.
- Once you have an open file handle, you can read the file line-by-line with a `for` loop.
- Each line of text is a string. How can you get the first character of a string?
- Using that first character, how can you set the value of a dict to be the key and the line itself to be the value?
- Iterate through each `letter` argument. How can you check that a given value is in the dictionary?

No skipping ahead to the solution until you have written your own version! If you peek, you will die a horrible death stamped by kittens.

8.2 Solution

```
#!/usr/bin/env python3
"""
Lookup tables"""

import argparse

# -----
def get_args():
    """
    get command-line arguments
    """

    parser = argparse.ArgumentParser(
        description='Gashlycrumb',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)
```

```

parser.add_argument('letter',
                    help='Letter(s)',
                    metavar='str',
                    nargs='+',
                    type=str) ①

parser.add_argument('-f',
                    '--file',
                    help='Input file',
                    metavar='str',
                    type=argparse.FileType('r'),
                    default='gashlycrumb.txt') ②

return parser.parse_args()

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()
    lookup = {line[0].upper(): line.rstrip() for line in args.file} ③

    for letter in args.letter:
        if letter.upper() in lookup:
            print(lookup[letter.upper()])
        else:
            print(f'I do not know "{letter}".') ④⑤⑥⑦

# -----
if __name__ == '__main__':
    main()

```

- ① A positional argument we'll call `letter` that uses `nargs='+'` to indicate one or more values are required.
- ② The optional `--file` argument must be a readable file. The `default` value is the `gashlycrumb.txt` file which we know to exist.
- ③ Build the `lookup` using a dictionary comprehension that reads the given file. Use the `upper` function to disregard case.
- ④ Use a `for` loop to iterate over each `letter` in `args.letter`.
- ⑤ See if the `letter` is in the `lookup` dictionary, checking the `upper` value to disregard case.
- ⑥ If so, print the line of text from the `lookup` for the `letter`.
- ⑦ Otherwise, print a message that the `letter` is unknown.

8.3 Discussion

Did the frightful paws of the kittens hurt much? Let's talk about how I solved this problem. Remember, mine is just one of many possible solutions.

8.3.1 Handling the arguments

I prefer to have all the logic for parsing and validating the command-line arguments in the `get_args` function. In particular, `argparse` can do a fine job verifying tedious things such as an argument being an existing, readable file which is why I

use `type=argparse.FileType('r')` for that argument. If the user doesn't supply a valid arguments, then argparse will throw an error, printing a helpful message along with the short usage and exiting with an error code.

By the time I get to the line `args = get_args()`, I know that I have one or more "letter" arguments and a valid, open file handle in the `args.file` slot. In the REPL, I can use `open` to get a file handle which I like to usually like to call `fh`. For copyright purposes, I'll use my alternate text:

```
>>> fh = open('alternate.txt')
```

8.3.2 Reading the input file

I know that I want to use a dictionary where the keys are the first letters of each line and the values are the lines themselves. That means I need to start by creating a new, empty dictionary either by using the `dict()` function or by setting a variable equal to an empty set of curly braces (`{}`). I'll call my variable `lookup`:

```
>>> lookup = {}
```

I will use a `for` loop to read each line of text. From Crow's Nest, I know I can use `line[0].upper()` to get the first letter of `line` and uppercase it. I can use that as the key into `lookup`. Each line of text ends with a newline that I'd like to remove. I can use the `str.rstrip` method to strip whitespace from the right side of the line (`rstrip = right strip`). The result of that will be the value for my `lookup`:

```
for line in fh:
    lookup[line[0].upper()] = line.rstrip()
```

I'd like to look at the resulting `lookup` dictionary. I can `print` it from the program or type `lookup` in the REPL, but it's going to be hard to read. I encourage you to try it. Luckily there is a lovely module called `pprint` to "pretty print" data structures. Here is how I can import the `pprint` function from the `pprint` module with the alias `pp`:

```
from pprint import pprint as pp
```



```
>>> from pprint import pprint as pp
```

Now let's take a peek at `lookup`:

```
>>> pp(lookup)
{'A': 'A is for Alfred, poisoned to death.',
 'B': 'B is for Bertrand, consumed by meth.',
 'C': 'C is for Cornell, who ate some glass.',
 'D': 'D is for Donald, who died from gas.',
 'E': 'E is for Edward, hanged by the neck.',
 'F': 'F is for Freddy, crushed in a wreck.',
 'G': 'G is for Geoffrey, who slit his wrist.',
 'H': "H is for Henry, who's neck got a twist.",
```

```
'I': 'I is for Ingrid, who tripped down a stair.',
'J': 'J is for Jered, who fell off a chair,',
'K': 'K is for Kevin, bit by a snake.',
'L': 'L is for Lauryl, impaled on a stake.',
'M': 'M is for Moira, hit by a brick.',
'N': 'N is for Norbert, who swallowed a stick.',
'O': 'O is for Orville, who fell in a canyon,',
'P': 'P is for Paul, strangled by his banyan,',
'Q': 'Q is for Quintanna, flayed in the night,',
'R': 'R is for Robert, who died of spite,',
'S': 'S is for Susan, stung by a jelly,',
'T': 'T is for Terrange, kicked in the belly,',
'U': 'U is for Uma, who's life was vanquished,",
'V': 'V is for Victor, consumed by anguish,',
'W': "W is for Walter, who's socks were too long,",
'X': 'X is for Xavier, stuck through with a prong,',
'Y': 'Y is for Yoeman, too fat by a piece,',
'Z': 'Z is for Zora, smothered by a fleece.'}
```

Hey, that looks like a handy data structure. Hooray for us!

8.3.3 Looping with `for` versus a list comprehensions

We used three lines of code to build our lookup dictionary. We can actually accomplish that in *one* line by using a "dictionary comprehension." We've created list comprehensions by sticking a `for` inside brackets `[]`. A dictionary comprehension is the same but the `for` loop is inside curly braces `{}`.

```
lookup = []
for line in fh:
    lookup[line[0].upper()] = line.rstrip()

lookup = {line[0].upper(): line.rstrip() for line in fh}
```

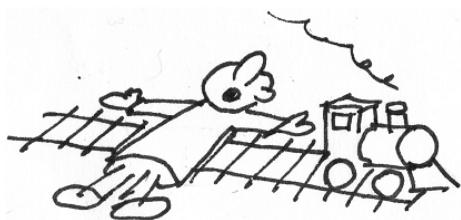
If you are following along by pasting code into the REPL, note that we have exhausted the file handle `fh` just above by reading it. (Refer back to the "Howler" to read about file handles.) I need to open it again for this next bit:

```
>>> fh = open('gashlycrumb.txt')
>>> lookup = {line[0].upper(): line.rstrip() for line in fh}
```

If you print it again, you should see the same output as above. It may seem like showing off to write one line of code instead of three, but it really does make a good deal of sense to write compact, idiomatic code. More code always means more chances for bugs, so I usually try to write code that is as simple as possible (but no simpler).

8.3.4 Dictionary lookups

Now that I have a lookup, I can ask if some value is in the keys. Note that I know the letters are in uppercase and I assume the user could give me lower, so I use `letter.upper()` to only compare that case:



```
>>> letter = 'a'
>>> letter.upper() in lookup
True
>>> lookup[letter.upper()]
'A is for Amy who fell down the stairs.'
```

If the letter is found, I can print the line of text for that letter; otherwise, I can print the message that I don't know that letter:

```
>>> letter = '4'
>>> if letter.upper() in lookup:
...     print(lookup[letter.upper()])
... else:
...     print('I do not know "{}".format(letter))
...
I do not know "4".
```

An even shorter way to write that would use the `dict.get` method:

```
def main():
    args = get_args()
    lookup = {line[0].upper(): line.rstrip() for line in args.file}

    for letter in args.letter:
        print(lookup.get(letter.upper(), f'I do not know "{letter}"')) ①
```

① If `lookup.get` will return the value for `letter.upper()` or the phrase "I don't know..."

8.4 Review

- A dictionary comprehension is a way to build a dictionary in a one-line `for` loop.
- Defining file input arguments using `argparse.FileType` saves you time and code.
- Python's `pprint` module is used for "pretty printing" complex data structures.

8.5 Going Further

- Write a phonebook that reads a file and creates a dictionary from the names of your friends to their email or phone numbers.
- Create a program that uses a dictionary to count the number of times you see each word in a document.
- Write an interactive version that takes input directly from the user. Use `while True` to set up an infinite loop and keep using `input` to get the user's next letter:

```
$ ./gashlycrumb_interactive.py
Please provide a letter [! to quit]: t
T is for Titus who flew into bits.
Please provide a letter [! to quit]: 7
I do not know "7".
Please provide a letter [! to quit]: !
Bye
```